

## 13. RANDOMIZED ALGORITHMS

---

- ▶ *contention resolution*
- ▶ *global min cut*
- ▶ *linearity of expectation*
- ▶ *max 3-satisfiability*
- ▶ *universal hashing*
- ▶ *Chernoff bounds*
- ▶ *load balancing*

Lecture slides by Kevin Wayne

Copyright © 2005 Pearson–Addison Wesley

<http://www.cs.princeton.edu/~wayne/kleinberg-tardos>

# Randomization

---

## Algorithmic design patterns.

- Greedy.
- Divide-and-conquer.
- Dynamic programming.
- Network flow.
- **Randomization.**

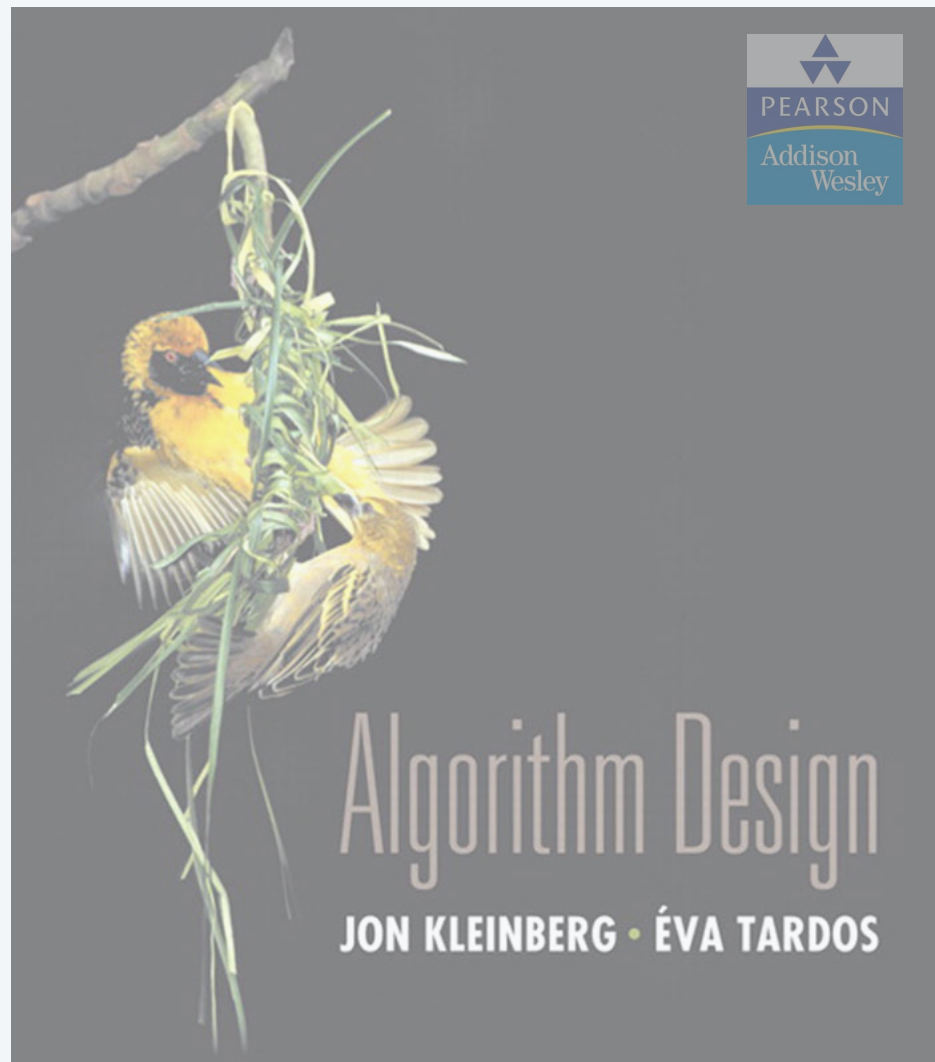
in practice, access to a pseudo-random number generator



**Randomization.** Allow fair coin flip in unit time.

**Why randomize?** Can lead to simplest, fastest, or only known algorithm for a particular problem.

**Ex.** Symmetry-breaking protocols, graph algorithms, quicksort, hashing, load balancing, Monte Carlo integration, cryptography.



## 13. RANDOMIZED ALGORITHMS

---

- ▶ *contention resolution*
- ▶ *global min cut*
- ▶ *linearity of expectation*
- ▶ *max 3-satisfiability*
- ▶ *universal hashing*
- ▶ *Chernoff bounds*
- ▶ *load balancing*


# Contention resolution in a distributed system

---

**Contention resolution.** Given  $n$  processes  $P_1, \dots, P_n$ , each competing for access to a shared database. If two or more processes access the database simultaneously, all processes are locked out. Devise protocol to ensure all processes get through on a regular basis.

**Restriction.** Processes can't communicate.

**Challenge.** Need symmetry-breaking paradigm.



# Contention resolution: randomized protocol

**Protocol.** Each process requests access to the database at time  $t$  with probability  $p = 1/n$ . *# processes.*

**Claim.** Let  $S[i, t]$  = event that process  $i$  succeeds in accessing the database at time  $t$ . Then  $1 / (e \cdot n) \leq \Pr [S(i, t)] \leq 1/(2n)$ .

**Pf.** By independence,  $\Pr [S(i, t)] = p (1 - p)^{n-1}$ .

*process i requests access*      *Process i successful at time t*  
*none of remaining n-1 processes request access*

- Setting  $p = 1/n$ , we have  $\Pr [S(i, t)] = \underbrace{1/n (1 - 1/n)^{n-1}}_{\text{between } 1/e \text{ and } 1/2}$ . ■
- value that maximizes  $\Pr[S(i, t)]$*

**Useful facts from calculus.** As  $n$  increases from 2, the function:

- $(1 - 1/n)^n$  converges monotonically from  $1/4$  up to  $1/e$ .
- $(1 - 1/n)^{n-1}$  converges monotonically from  $1/2$  down to  $1/e$ .

# Contention resolution: randomized protocol

---

**Claim.** The probability that process  $i$  fails to access the database in  $en$  rounds is at most  $1/e$ . After  $e \cdot n (c \ln n)$  rounds, the probability  $\leq n^{-c}$ .

**Pf.** Let  $F[i, t]$  = event that process  $i$  fails to access database in rounds 1 through  $t$ . By independence and previous claim, we have

$$\Pr[F[i, t]] \leq (1 - 1/(en))^t.$$

- Choose  $t = \lceil e \cdot n \rceil$ :

$$\Pr[F(i, t)] \leq \left(1 - \frac{1}{en}\right)^{\lceil en \rceil} \leq \left(1 - \frac{1}{en}\right)^{en} \leq \frac{1}{e}$$

- Choose  $t = \lceil e \cdot n \rceil \lceil c \ln n \rceil$ :

$$\Pr[F(i, t)] \leq \left(\frac{1}{e}\right)^{c \ln n} = n^{-c}$$

constant  
prob. of failure.

Even at 8:59

# Contention resolution: randomized protocol

**Claim.** The probability that all processes succeed within  $2e \cdot n \ln n$  rounds is  $\geq 1 - 1/n$ .

**Pf.** Let  $F[t]$  = event that at least one of the  $n$  processes fails to access database in any of the rounds 1 through  $t$ .

$$\Pr[F[t]] = \Pr\left[\bigcup_{i=1}^n F[i, t]\right] \leq \sum_{i=1}^n \Pr[F[i, t]] \leq n \left(1 - \frac{1}{en}\right)^t$$

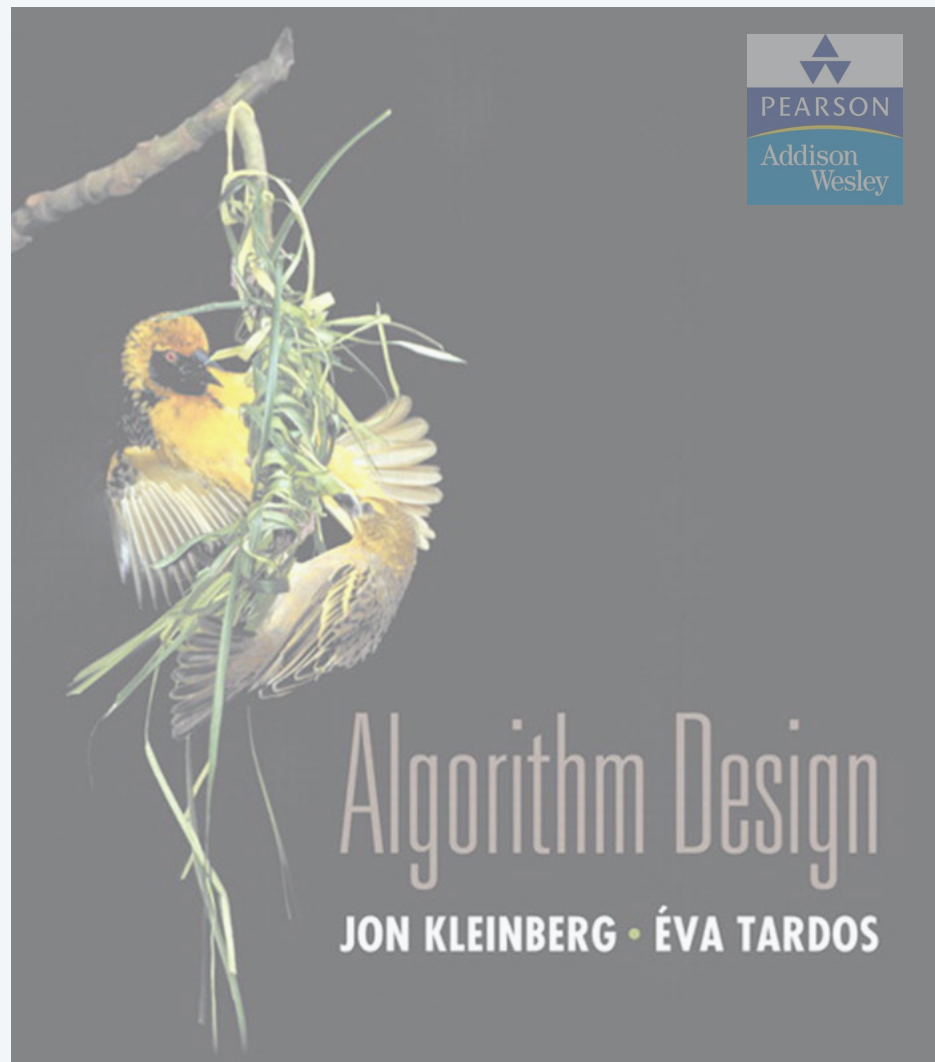
union bound                      previous slide

- Choosing  $t = 2 \lceil en \rceil \lceil c \ln n \rceil$  yields  $\Pr[F[t]] \leq n \cdot n^{-2} = 1/n$ . ■

**Union bound.**

Given events  $E_1, \dots, E_n$ ,

$$\Pr\left[\bigcup_{i=1}^n E_i\right] \leq \sum_{i=1}^n \Pr[E_i]$$



## 13. RANDOMIZED ALGORITHMS

---

- ▶ *contention resolution*
- ▶ *global min cut*
- ▶ *linearity of expectation*
- ▶ *max 3-satisfiability*
- ▶ *universal hashing*
- ▶ *Chernoff bounds*
- ▶ *load balancing*



# Global minimum cut

---

**Global min cut.** Given a connected, undirected graph  $G = (V, E)$ , find a cut  $(A, B)$  of minimum cardinality.

**Applications.** Partitioning items in a database, identify clusters of related documents, network reliability, network design, circuit design, TSP solvers.

**Network flow solution.**

- Replace every edge  $(u, v)$  with two antiparallel edges  $(u, v)$  and  $(v, u)$ .
- Pick some vertex  $s$  and compute min  $s-v$  cut separating  $s$  from each other vertex  $v \in V$ . *Repeat for every vertex.*


**False intuition.** Global min-cut is harder than min  $s-t$  cut.

# Contraction algorithm


---

Contraction algorithm. [Karger 1995]

- Pick an edge  $e = (u, v)$  uniformly at random.
- Contract edge  $e$ .
  - replace  $u$  and  $v$  by single new super-node  $w$
  - preserve edges, updating endpoints of  $u$  and  $v$  to  $w$
  - keep parallel edges, but delete self-loops
- Repeat until graph has just two nodes  $u_1$  and  $v_1$ .
- Return the cut (all nodes that were contracted to form  $v_1$ ).



$\Rightarrow$   
contract  $u-v$




# Contraction algorithm

## Contraction algorithm. [Karger 1995]

- Pick an edge  $e = (u, v)$  uniformly at random.
- **Contract** edge  $e$ .
  - replace  $u$  and  $v$  by single new super-node  $w$
  - preserve edges, updating endpoints of  $u$  and  $v$  to  $w$
  - keep parallel edges, but delete self-loops
- Repeat until graph has just two nodes  $u_1$  and  $v_1$ .
- Return the cut (all nodes that were contracted to form  $v_1$ ).

No constraints on choosing edges (only randomly)



*We obtain  
2 because there  
are finite.*


# Contraction algorithm

**Claim.** The contraction algorithm returns a min cut with prob  $\geq 2/n^2$ .

**Pf.** Consider a global min-cut  $(A^*, B^*)$  of  $G$ .

- Let  $F^*$  be edges with one endpoint in  $A^*$  and the other in  $B^*$ .
- Let  $k = |F^*|$  = size of min cut.
- In **first step**, algorithm contracts an edge in  $F^*$  with probability  $k/|E|$ .
- Every node has  $\text{degree} \geq k$  since otherwise  $(A^*, B^*)$  would not be a min-cut  $\Rightarrow |E| \geq \frac{1}{2} k n \Leftrightarrow k/|E| \leq 2/n \rightarrow \mathcal{P}(F)$
- Thus, algorithm contracts an edge in  $F^*$  with probability  $\leq 2/n$ .

Contr.  
of edges.



# Contraction algorithm

**Claim.** The contraction algorithm returns a min cut with prob  $\geq 2 / n^2$ .

**Pf.** Consider a global min-cut  $(A^*, B^*)$  of  $G$ .

- Let  $F^*$  be edges with one endpoint in  $A^*$  and the other in  $B^*$ .
- Let  $k = |F^*|$  = size of min cut.
- Let  $G'$  be graph after  $j$  iterations. There are  $n' = n - j$  supernodes.
- Suppose no edge in  $F^*$  has been contracted. The min-cut in  $G'$  is still  $k$ .
- Since value of min-cut is  $k$ ,  $|E'| \geq \frac{1}{2} k n' \Leftrightarrow k / |E'| \leq 2 / n'$ .
- Thus, algorithm contracts an edge in  $F^*$  with probability  $\leq 2 / n'$ .
- Let  $E_j$  = event that an edge in  $F^*$  is not contracted in iteration  $j$ .

$$\begin{aligned} \Pr[E_1 \cap E_2 \cdots \cap E_{n-2}] &= \Pr[E_1] \times \Pr[E_2 \mid E_1] \times \cdots \times \Pr[E_{n-2} \mid E_1 \cap E_2 \cdots \cap E_{n-3}] \\ &\geq \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \cdots \left(1 - \frac{2}{4}\right) \left(1 - \frac{2}{3}\right) \\ &= \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \cdots \left(\frac{2}{4}\right) \left(\frac{1}{3}\right) \\ &= \frac{2}{n(n-1)} \\ &\geq \frac{2}{n^2} \end{aligned}$$

# Contraction algorithm

---

**Amplification.** To amplify the probability of success, run the contraction algorithm many times.

with independent random choices,

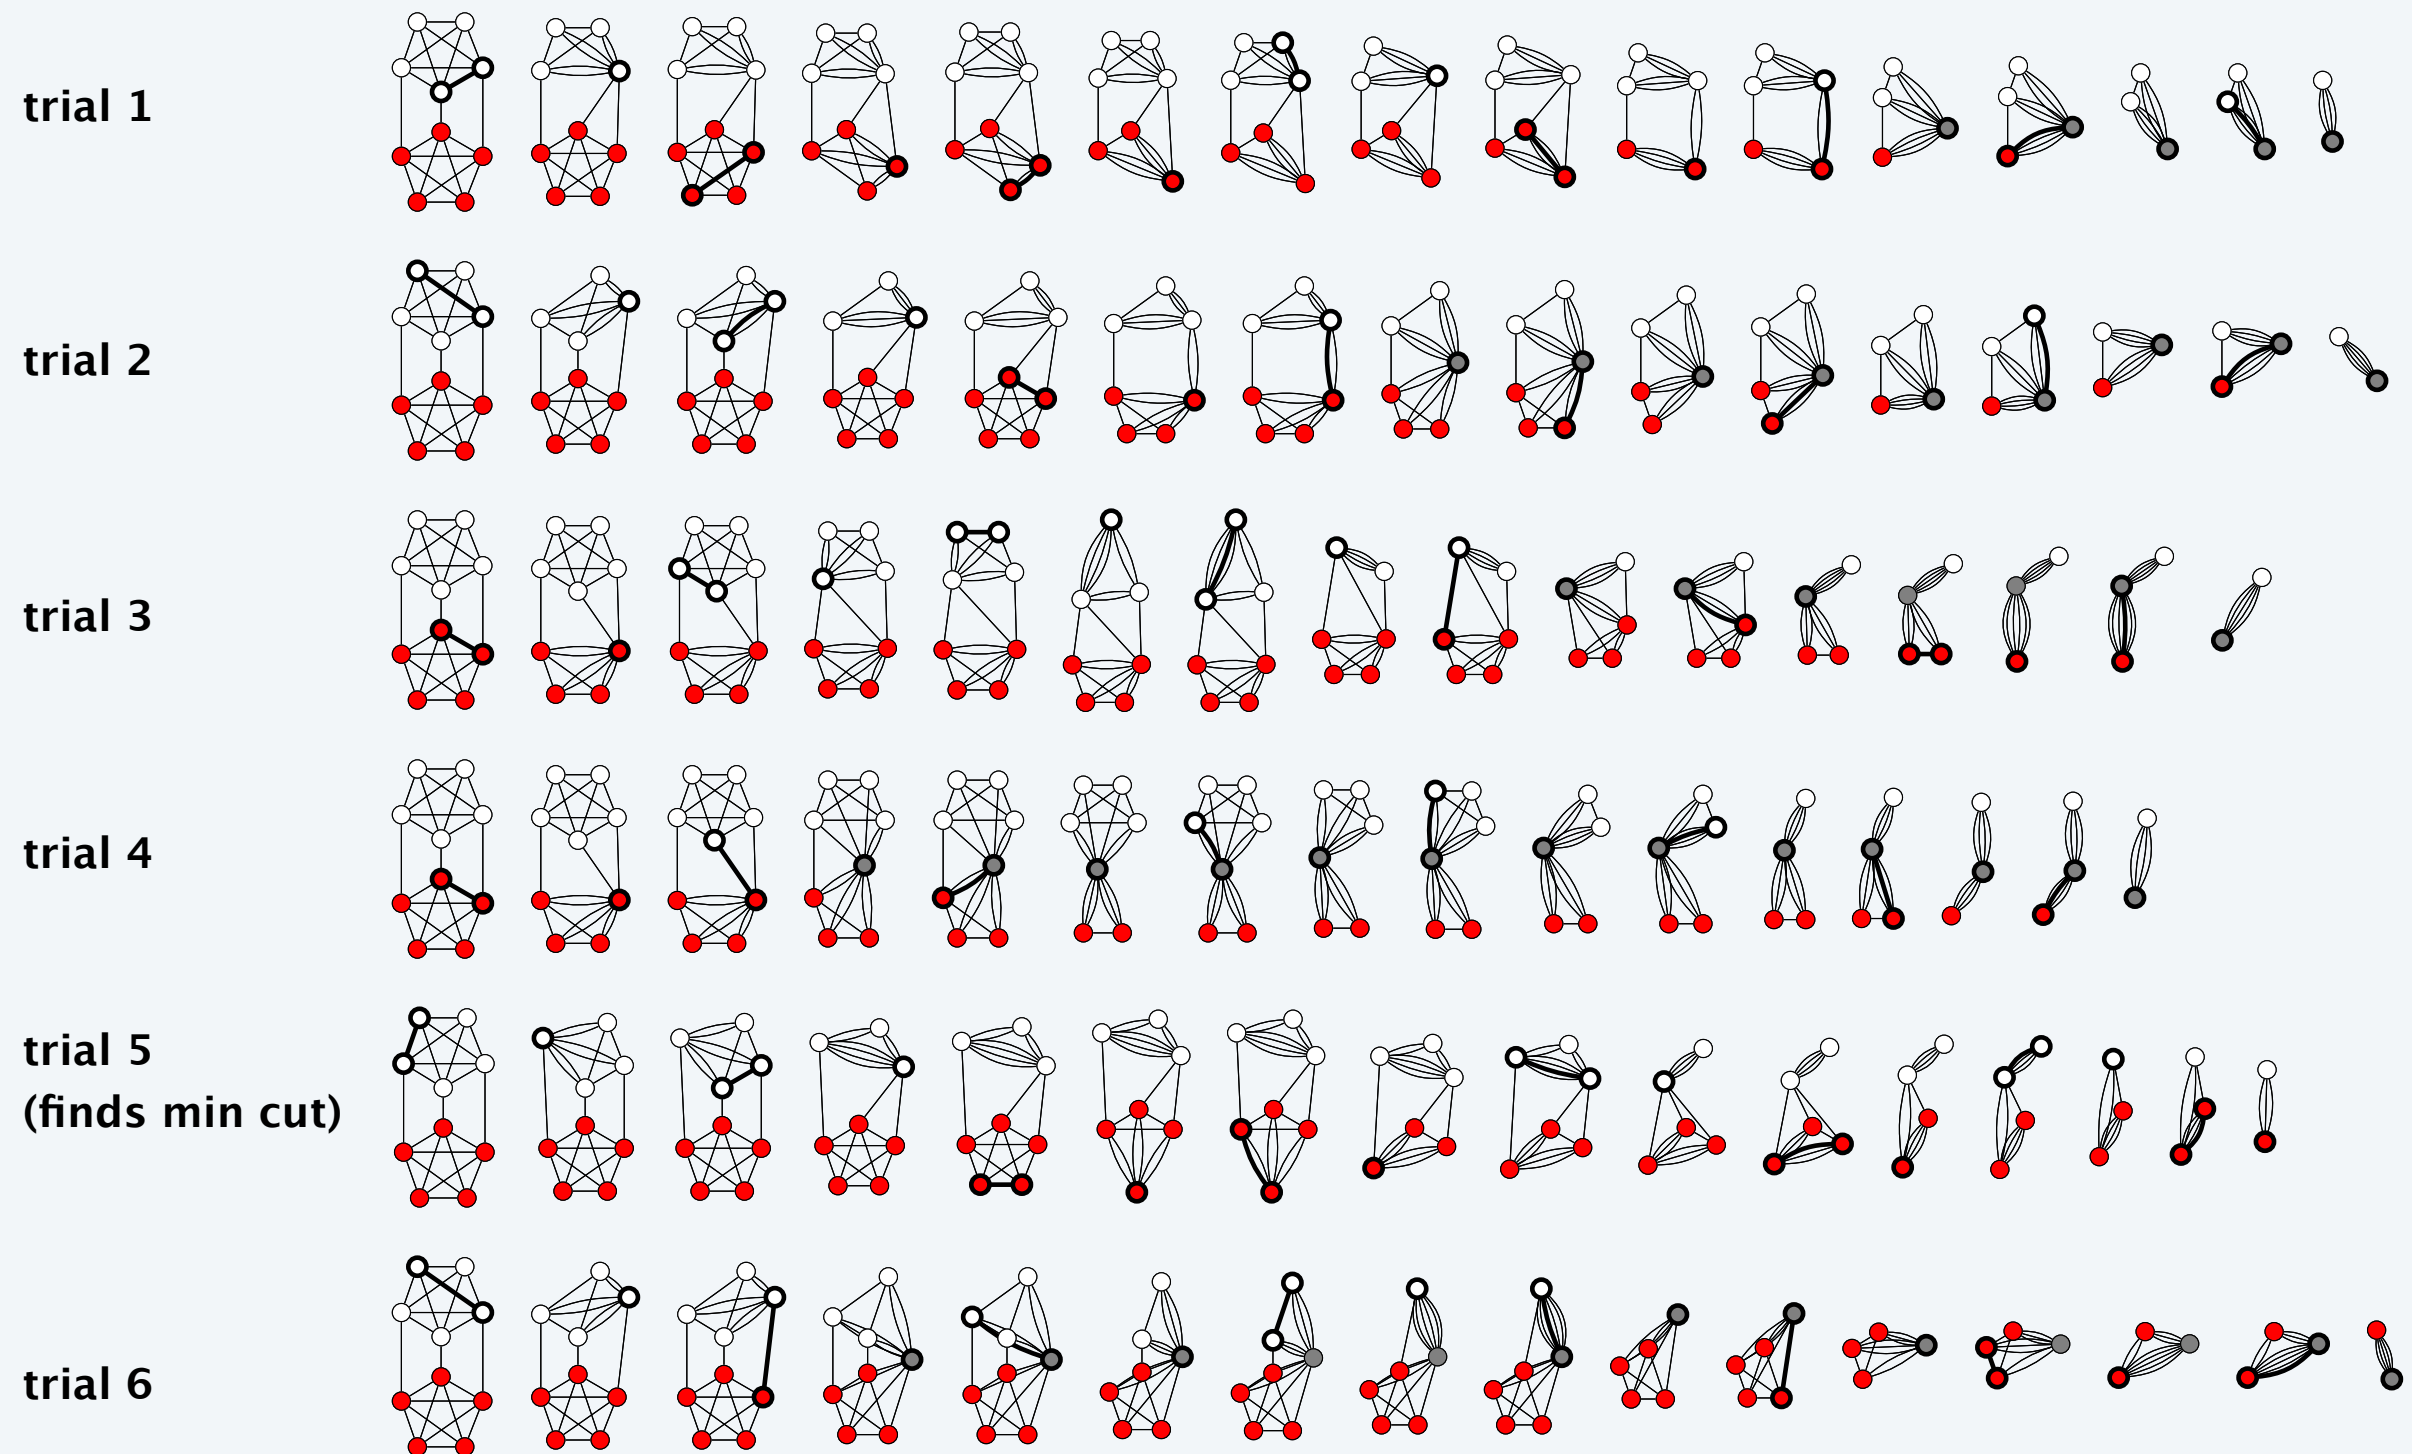
**Claim.** If we repeat the contraction algorithm  $n^2 \ln n$  times, then the probability of failing to find the global min-cut is  $\leq 1 / n^2$ .

**Pf.** By independence, the probability of failure is at most

$$\left(1 - \frac{2}{n^2}\right)^{n^2 \ln n} = \left[\left(1 - \frac{2}{n^2}\right)^{\frac{1}{2}n^2}\right]^{2 \ln n} \leq \left(e^{-1}\right)^{2 \ln n} = \frac{1}{n^2}$$

$(1 - 1/x)^x \leq 1/e$

# Contraction algorithm: example execution



...

Reference: Thore Husfeldt

# Global min cut: context

---

**Remark.** Overall running time is slow since we perform  $\Theta(n^2 \log n)$  iterations and each takes  $\Omega(m)$  time.

**Improvement.** [Karger–Stein 1996]  $O(n^2 \log^3 n)$ .

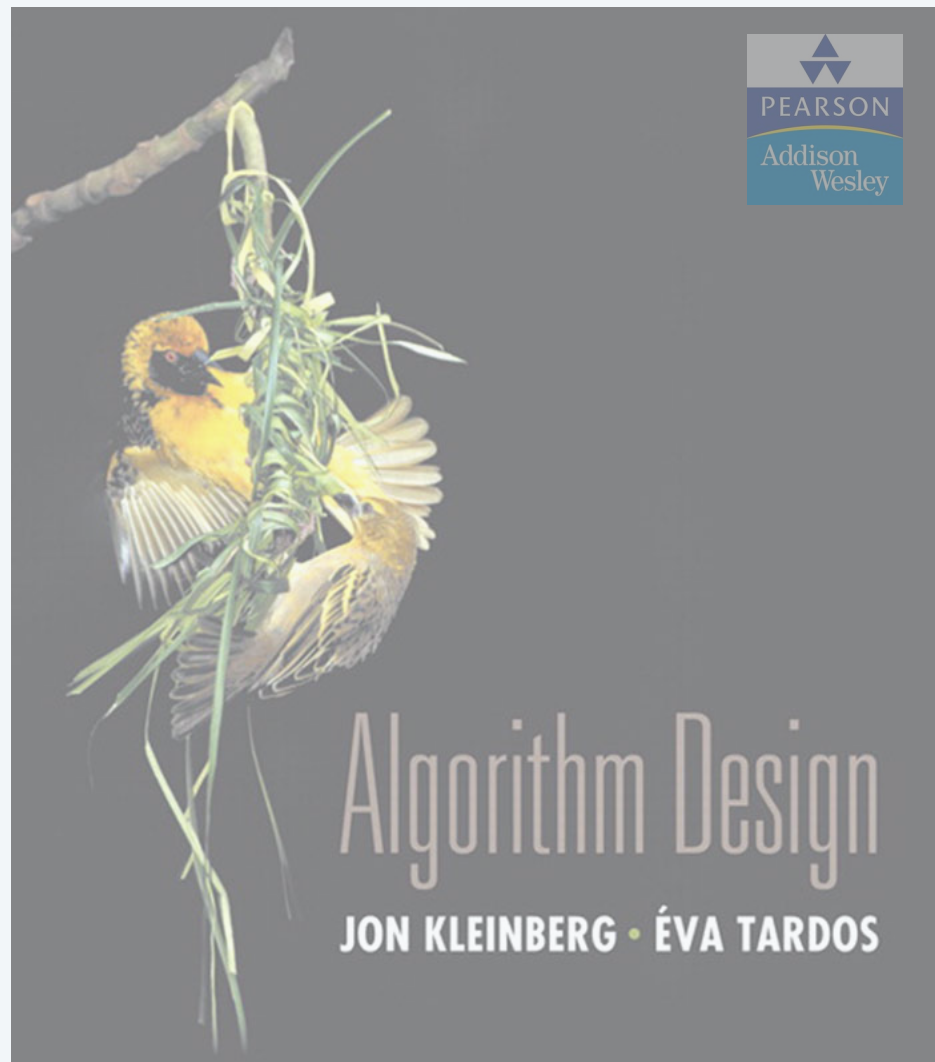
- Early iterations are less risky than later ones: probability of contracting an edge in min cut hits 50% when  $n / \sqrt{2}$  nodes remain.
- Run contraction algorithm until  $n / \sqrt{2}$  nodes remain.
- Run contraction algorithm **twice** on resulting graph and return **best** of two cuts.

**Extensions.** Naturally generalizes to handle positive weights.

**Best known.** [Karger 2000]  $O(m \log^3 n)$ .

faster than best known max flow algorithm or  
deterministic global min cut algorithm





## 13. RANDOMIZED ALGORITHMS

---

- ▶ *contention resolution*
- ▶ *global min cut*
- ▶ ***linearity of expectation***
- ▶ *max 3-satisfiability*
- ▶ *universal hashing*
- ▶ *Chernoff bounds*
- ▶ *load balancing*

# Expectation

---

**Expectation.** Given a discrete random variable  $X$ , its expectation  $E[X]$  is defined by:

$$E[X] = \sum_{j=0}^{\infty} j \Pr[X = j]$$

**Waiting for a first success.** Coin is heads with probability  $p$  and tails with probability  $1-p$ . How many independent flips  $X$  until first heads?

$$E[X] = \sum_{j=0}^{\infty} j \cdot \Pr[X = j] = \sum_{j=0}^{\infty} j \underset{\substack{\uparrow \\ j-1 \text{ tails}}}{(1-p)^{j-1}} \underset{\substack{\uparrow \\ 1 \text{ head}}}{p} = \frac{p}{1-p} \sum_{j=0}^{\infty} j (1-p)^j = \frac{p}{1-p} \cdot \frac{1-p}{p^2} = \frac{1}{p}$$

$$\frac{1}{\left(\frac{1}{2}\right)} = 2 \text{ flips.}$$

## Expectation: two properties

---

**Useful property.** If  $X$  is a 0/1 random variable,  $E[X] = \Pr[X = 1]$ .

**Pf.** 
$$E[X] = \sum_{j=0}^{\infty} j \cdot \Pr[X = j] = \sum_{j=0}^1 j \cdot \Pr[X = j] = \Pr[X = 1]$$

not necessarily independent

**Linearity of expectation.** Given two random variables  $X$  and  $Y$  defined over the same probability space,  $E[X + Y] = E[X] + E[Y]$ .

**Benefit.** Decouples a complex calculation into simpler pieces.

# Guessing cards

---

**Game.** Shuffle a deck of  $n$  cards; turn them over one at a time; try to guess each card.

---

**Memoryless guessing.** No psychic abilities; can't even remember what's been turned over already. Guess a card from full deck uniformly at random.

**Claim.** The expected number of correct guesses is 1.

**Pf.** [surprisingly effortless using linearity of expectation]

- Let  $X_i = 1$  if  $i^{th}$  prediction is correct and 0 otherwise.
- Let  $X = \text{number of correct guesses} = X_1 + \dots + X_n$ .

- $E[X_i] = \Pr[X_i = 1] = 1/n$ .

- $E[X] = E[X_1] + \dots + E[X_n] = 1/n + \dots + 1/n = 1.$  ■

$\uparrow \quad \Sigma$   
linearity of expectation

# Guessing cards

---

**Game.** Shuffle a deck of  $n$  cards; turn them over one at a time; try to guess each card.

**Guessing with memory.** Guess a card uniformly at random from cards not yet seen.

**Claim.** The expected number of correct guesses is  $\Theta(\log n)$ .

**Pf.**

- Let  $X_i = 1$  if  $i$ th prediction is correct and 0 otherwise.
- Let  $X =$  number of correct guesses  $= X_1 + \dots + X_n$ .
- $E[X_i] = \Pr[X_i = 1] = 1 / (n - (i - 1))$ .
- $E[X] = E[X_1] + \dots + E[X_n] = 1/n + \dots + 1/2 + 1/1 = H(n)$ . ■

↑  
linearity of expectation

↑  
 $\ln(n+1) < H(n) < 1 + \ln n$

# Coupon collector

**Coupon collector.** Each box of cereal contains a coupon. There are  $n$  different types of coupons. Assuming all boxes are equally likely to contain each coupon, how many boxes before you have  $\geq 1$  coupon of each type?


**Claim.** The expected number of steps is  $\Theta(n \log n)$ .

**Pf.**

- Phase  $j$  = time between  $j$  and  $j + 1$  distinct coupons.
- Let  $X_j$  = number of steps you spend in phase  $j$ .
- Let  $X$  = number of steps in total =  $X_0 + X_1 + \dots + X_{n-1}$ .

$$E[X] = \sum_{j=0}^{n-1} E[X_j] = \sum_{j=0}^{n-1} \frac{n}{n-j} = n \sum_{i=1}^n \frac{1}{i} = n H(n)$$

prob of success =  $(n - j) / n$   
 $\Rightarrow$  expected waiting time =  $n / (n - j)$



## 13. RANDOMIZED ALGORITHMS

---

- ▶ *contention resolution*
- ▶ *global min cut*
- ▶ *linearity of expectation*
- ▶ ***max 3-satisfiability***
- ▶ *universal hashing*
- ▶ *Chernoff bounds*
- ▶ *load balancing*

# Maximum 3-satisfiability

---

exactly 3 distinct literals per clause

**Maximum 3-satisfiability.** Given a 3-SAT formula, find a truth assignment that satisfies as many clauses as possible.

$$\begin{aligned}C_1 &= x_2 \vee \overline{x_3} \vee \overline{x_4} \\C_2 &= x_2 \vee x_3 \vee \overline{x_4} \\C_3 &= \overline{x_1} \vee x_2 \vee x_4 \\C_4 &= \overline{x_1} \vee \overline{x_2} \vee x_3 \\C_5 &= x_1 \vee \overline{x_2} \vee \overline{x_4}\end{aligned}$$

**Remark.** **NP-hard** search problem.

**Simple idea.** Flip a coin, and set each variable true with probability  $\frac{1}{2}$ , independently for each variable.



# Maximum 3-satisfiability: analysis

---

**Claim.** Given a 3-SAT formula with  $k$  clauses, the expected number of clauses satisfied by a random assignment is  $7k / 8$ .

**Pf.** Consider random variable  $Z_j = \begin{cases} 1 & \text{if clause } C_j \text{ is satisfied} \\ 0 & \text{otherwise.} \end{cases}$

- Let  $Z$  = number of clauses satisfied by random assignment.

$$\begin{aligned} E[Z] &= \sum_{j=1}^k E[Z_j] \\ &= \sum_{j=1}^k \Pr[\text{clause } C_j \text{ is satisfied}] \\ &= \frac{7}{8} k \end{aligned}$$

linearity of expectation

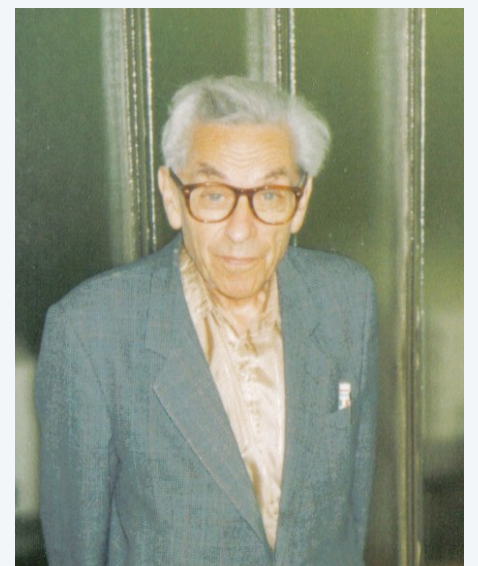
# The probabilistic method

---

**Corollary.** For any instance of 3-SAT, there exists a truth assignment that satisfies at least a  $7/8$  fraction of all clauses.

**Pf.** Random variable is at least its expectation some of the time. ■

**Probabilistic method.** [Paul Erdős] Prove the existence of a non-obvious property by showing that a random construction produces it with positive probability!



# Maximum 3-satisfiability: analysis

---

Q. Can we turn this idea into a  $7/8$ -approximation algorithm?

A. Yes (but a random variable can almost always be below its mean).

**Lemma.** The probability that a random assignment satisfies  $\geq 7k/8$  clauses is at least  $1/(8k)$ .

**Pf.** Let  $p_j$  be probability that exactly  $j$  clauses are satisfied;  
let  $p$  be probability that  $\geq 7k/8$  clauses are satisfied.

$$\begin{aligned}\frac{7}{8}k &= E[Z] = \sum_{j \geq 0} j p_j \\ &= \sum_{j < 7k/8} j p_j + \sum_{j \geq 7k/8} j p_j \\ &\leq \left(\frac{7k}{8} - \frac{1}{8}\right) \sum_{j < 7k/8} p_j + k \sum_{j \geq 7k/8} p_j \\ &\leq \left(\frac{7}{8}k - \frac{1}{8}\right) \cdot 1 + k p\end{aligned}$$

Rearranging terms yields  $p \geq 1/(8k)$ . ■

## Maximum 3-satisfiability: analysis

---

**Johnson's algorithm.** Repeatedly generate random truth assignments until one of them satisfies  $\geq 7k / 8$  clauses.

**Theorem.** Johnson's algorithm is a  $7/8$ -approximation algorithm.

**Pf.** By previous lemma, each iteration succeeds with probability  $\geq 1 / (8k)$ . By the waiting-time bound, the expected number of trials to find the satisfying assignment is at most  $8k$ . ■

# Maximum satisfiability

---

## Extensions.

- Allow one, two, or more literals per clause.
- Find max **weighted** set of satisfied clauses.

**Theorem.** [Asano–Williamson 2000] There exists a 0.784-approximation algorithm for MAX-SAT.

**Theorem.** [Karloff–Zwick 1997, Zwick+computer 2002] There exists a  $7/8$ -approximation algorithm for version of MAX-3-SAT in which each clause has **at most** 3 literals.

**Theorem.** [Håstad 1997] Unless  $\mathbf{P} = \mathbf{NP}$ , no  $\rho$ -approximation algorithm for MAX-3-SAT (and hence MAX-SAT) for any  $\rho > 7/8$ .



very unlikely to improve over simple randomized  
algorithm for MAX-3-SAT

# Monte Carlo vs. Las Vegas algorithms

---

**Monte Carlo.** Guaranteed to run in poly-time, likely to find correct answer.

**Ex:** Contraction algorithm for global min cut.

**Las Vegas.** Guaranteed to find correct answer, likely to run in poly-time.

**Ex:** Randomized quicksort, Johnson's MAX-3-SAT algorithm.

stop algorithm  
after a certain point

**Remark.** Can always convert a Las Vegas algorithm into Monte Carlo,  
but no known method (in general) to convert the other way.

# RP and ZPP


---

**RP.** [Monte Carlo] Decision problems solvable with **one-sided error** in poly-time.

## One-sided error.


- If the correct answer is *no*, always return *no*.
- If the correct answer is *yes*, return *yes* with probability  $\geq \frac{1}{2}$ .

can decrease probability of false negative  
to  $2^{-100}$  by 100 independent repetitions



**ZPP.** [Las Vegas] Decision problems solvable in **expected** poly-time.


running time can be unbounded,  
but fast on average



**Theorem.**  $\mathbf{P} \subseteq \mathbf{ZPP} \subseteq \mathbf{RP} \subseteq \mathbf{NP}$ .

**Fundamental open questions.** To what extent does randomization help?

Does  $\mathbf{P} = \mathbf{ZPP}$ ? Does  $\mathbf{ZPP} = \mathbf{RP}$ ? Does  $\mathbf{RP} = \mathbf{NP}$ ?



## 13. RANDOMIZED ALGORITHMS

---

- ▶ *contention resolution*
- ▶ *global min cut*
- ▶ *linearity of expectation*
- ▶ *max 3-satisfiability*
- ▶ ***universal hashing***
- ▶ *Chernoff bounds*
- ▶ *load balancing*



# Dictionary data type

---

**Dictionary.** Given a universe  $U$  of possible elements, maintain a subset  $S \subseteq U$  so that **inserting**, **deleting**, and **searching** in  $S$  is efficient.

## Dictionary interface.

- *create()*: initialize a dictionary with  $S = \emptyset$ .
- *insert( $u$ )*: add element  $u \in U$  to  $S$ .
- *delete( $u$ )*: delete  $u$  from  $S$  (if  $u$  is currently in  $S$ ).
- *lookup( $u$ )*: is  $u$  in  $S$ ?

**Challenge.** Universe  $U$  can be extremely large so defining an array of size  $|U|$  is infeasible.

**Applications.** File systems, databases, Google, compilers, checksums P2P networks, associative arrays, cryptography, web caching, etc.

# Hashing

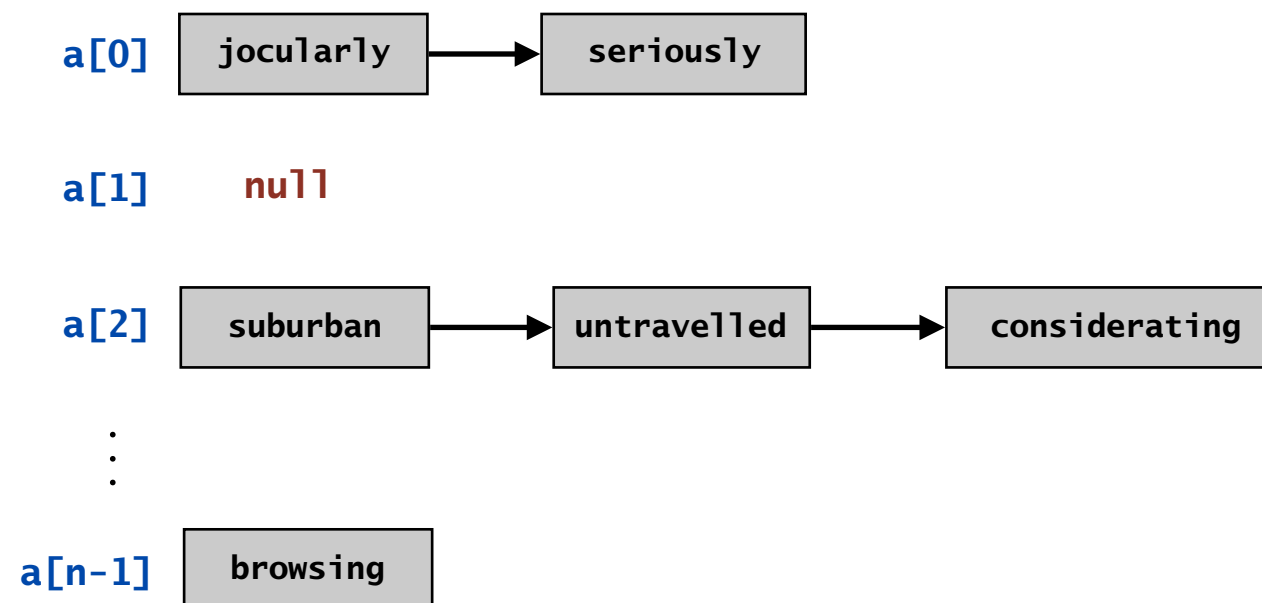
**Hash function.**  $h : U \rightarrow \{0, 1, \dots, n-1\}$ .

**Hashing.** Create an array  $a$  of length  $n$ . When processing element  $u$ , access array element  $a[h(u)]$ .

**Collision.** When  $h(u) = h(v)$  but  $u \neq v$ .

birthday paradox

- A collision is expected after  $\Theta(\sqrt{n})$  random insertions.
- Separate chaining:  $a[i]$  stores linked list of elements  $u$  with  $h(u) = i$ .



# Ad-hoc hash function

---

## Ad-hoc hash function.

```
int hash(String s, int n) {  
    int hash = 0;  
    for (int i = 0; i < s.length(); i++)  
        hash = (31 * hash) + s[i];  
    return hash % n;  
}
```

hash function à la Java string library

**Deterministic hashing.** If  $|U| \geq n^2$ , then for any fixed hash function  $h$ , there is a subset  $S \subseteq U$  of  $n$  elements that all hash to same slot. Thus,  $\Theta(n)$  time per lookup in worst-case.


**Q.** But isn't ad-hoc hash function good enough in practice?

# Algorithmic complexity attacks

---

## When can't we live with ad-hoc hash function?

- Obvious situations: aircraft control, nuclear reactor, pace maker, ....
- Surprising situations: denial-of-service attacks.



malicious adversary learns your ad-hoc hash function  
(e.g., by reading Java API) and causes a big pile-up  
in a single slot that grinds performance to a halt

## Real world exploits. [Crosby–Wallach 2003]

- Linux 2.4.20 kernel: save files with carefully chosen names.
- Perl 5.8.0: insert carefully chosen strings into associative array.
- Bro server: send carefully chosen packets to DOS the server,  
using less bandwidth than a dial-up modem.

# Hashing performance

---

**Ideal hash function.** Maps  $m$  elements uniformly at random to  $n$  hash slots.

- Running time depends on length of chains.
- Average length of chain  $= \alpha = m / n$ .
- Choose  $n \approx m \Rightarrow$  expect  $O(1)$  per insert, lookup, or delete.

**Challenge.** Explicit hash function  $h$  that achieves  $O(1)$  per operation.

**Approach.** Use randomization for the choice of  $h$ .

adversary knows the randomized algorithm you're using,  
but doesn't know random choice that the algorithm makes

# Universal hashing (Carter–Wegman 1980s)

A **universal family of hash functions** is a set of hash functions  $H$  mapping a universe  $U$  to the set  $\{0, 1, \dots, n-1\}$  such that

- For any pair of elements  $u \neq v$ :  $\Pr_{h \in H} [h(u) = h(v)] \leq 1/n$
- Can select random  $h$  efficiently.
- Can compute  $h(u)$  efficiently.

chosen uniformly at random

**Ex.**  $U = \{a, b, c, d, e, f\}$ ,  $n = 2$ .

	a	b	c	d	e	f
$h_1(x)$	0	1	0	1	0	1
$h_2(x)$	0	0	0	1	1	1

	a	b	c	d	e	f
$h_1(x)$	0	1	0	1	0	1
$h_2(x)$	0	0	0	1	1	1
$h_3(x)$	0	0	1	0	1	1
$h_4(x)$	1	0	0	1	1	0

$$H = \{h_1, h_2\}$$

$$\Pr_{h \in H} [h(a) = h(b)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(c)] = 1$$

$$\Pr_{h \in H} [h(a) = h(d)] = 0$$

...

$$H = \{h_1, h_2, h_3, h_4\}$$

$$\Pr_{h \in H} [h(a) = h(b)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(c)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(d)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(e)] = 1/2$$

$$\Pr_{h \in H} [h(a) = h(f)] = 0$$

...

**not universal**

**universal**

# Universal hashing: analysis

---

**Proposition.** Let  $H$  be a universal family of hash functions mapping a universe  $U$  to the set  $\{0, 1, \dots, n-1\}$ ; let  $h \in H$  be chosen uniformly at random from  $H$ ; let  $S \subseteq U$  be a subset of size at most  $n$ ; and let  $u \notin S$ . Then, the expected number of items in  $S$  that collide with  $u$  is at most 1.

**Pf.** For any  $s \in S$ , define random variable  $X_s = 1$  if  $h(s) = h(u)$ , and 0 otherwise. Let  $X$  be a random variable counting the total number of collisions with  $u$ .

$$E_{h \in H}[X] = E[\sum_{s \in S} X_s] \underset{\text{linearity of expectation}}{=} \sum_{s \in S} E[X_s] \underset{X_s \text{ is a 0-1 random variable}}{=} \sum_{s \in S} \Pr[X_s = 1] \underset{\text{universal}}{\leq} \sum_{s \in S} \frac{1}{n} = |S| \frac{1}{n} \leq 1$$

**Q.** OK, but how do we design a universal class of hash functions?

# Designing a universal family of hash functions

---

**Modulus.** We will use a prime number  $p$  for the size of the hash table.

**Integer encoding.** Identify each element  $u \in U$  with a base- $p$  integer of  $r$  digits:  $x = (x_1, x_2, \dots, x_r)$ .

**Hash function.** Let  $A$  = set of all  $r$ -digit, base- $p$  integers. For each  $a = (a_1, a_2, \dots, a_r)$  where  $0 \leq a_i < p$ , define

$$h_a(x) = \left( \sum_{i=1}^r a_i x_i \right) \bmod p \quad \leftarrow \text{maps universe } U \text{ to set } \{0, 1, \dots, p-1\}$$

**Hash function family.**  $H = \{ h_a : a \in A \}$ .



# Designing a universal family of hash functions

---


**Theorem.**  $H = \{ h_a : a \in A \}$  is a universal family of hash functions.

**Pf.** Let  $x = (x_1, x_2, \dots, x_r)$  and  $y = (y_1, y_2, \dots, y_r)$  be two distinct elements of  $U$ .

We need to show that  $\Pr[h_a(x) = h_a(y)] \leq 1/p$ .

- Since  $x \neq y$ , there exists an integer  $j$  such that  $x_j \neq y_j$ .
- We have  $h_a(x) = h_a(y)$  iff

$$a_j \underbrace{(y_j - x_j)}_z \equiv \underbrace{\sum_{i \neq j} a_i (x_i - y_i)}_m \pmod{p}$$

- Can assume  $a$  was chosen uniformly at random by first selecting all coordinates  $a_i$  where  $i \neq j$ , then selecting  $a_j$  at random. Thus, we can assume  $a_i$  is fixed for all coordinates  $i \neq j$ .
- Since  $p$  is prime,  $a_j z \equiv m \pmod{p}$  has at most one solution among  $p$  possibilities.  see lemma on next slide
- Thus  $\Pr[h_a(x) = h_a(y)] \leq 1/p$ . ■

# Number theory fact

---

**Fact.** Let  $p$  be prime, and let  $z \not\equiv 0 \pmod{p}$ . Then  $\alpha z \equiv m \pmod{p}$  has at most one solution  $0 \leq \alpha < p$ .

**Pf.**

- Suppose  $0 \leq \alpha_1 < p$  and  $0 \leq \alpha_2 < p$  are two different solutions.
- Then  $(\alpha_1 - \alpha_2) z \equiv 0 \pmod{p}$ ; hence  $(\alpha_1 - \alpha_2) z$  is divisible by  $p$ .
- Since  $z \not\equiv 0 \pmod{p}$ , we know that  $z$  is not divisible by  $p$ .
- It follows that  $(\alpha_1 - \alpha_2)$  is divisible by  $p$ .
- This implies  $\alpha_1 = \alpha_2$ . ■

here's where we  
use that  $p$  is prime

**Bonus fact.** Can replace “at most one” with “exactly one” in above fact.

**Pf idea.** Euclid's algorithm.


# Universal hashing: summary

---

**Goal.** Given a universe  $U$ , maintain a subset  $S \subseteq U$  so that insert, delete, and lookup are efficient.


**Universal hash function family.**  $H = \{ h_a : a \in A \}$ .

$$h_a(x) = \left( \sum_{i=1}^r a_i x_i \right) \bmod p$$

- Choose  $p$  prime so that  $m \leq p \leq 2m$ , where  $m = |S|$ .
- Fact: there exists a prime between  $m$  and  $2m$ .  can find such a prime using another randomized algorithm (!)

**Consequence.**

- Space used =  $\Theta(m)$ .
- Expected number of collisions per operation is  $\leq 1$   
 $\Rightarrow O(1)$  time per insert, delete, or lookup.



## 13. RANDOMIZED ALGORITHMS

---

- ▶ *contention resolution*
- ▶ *global min cut*
- ▶ *linearity of expectation*
- ▶ *max 3-satisfiability*
- ▶ *universal hashing*
- ▶ **Chernoff bounds**
- ▶ *load balancing*

# Chernoff Bounds (above mean)

**Theorem.** Suppose  $X_1, \dots, X_n$  are independent 0-1 random variables. Let  $X = X_1 + \dots + X_n$ . Then for any  $\mu \geq E[X]$  and for any  $\delta > 0$ , we have

$$\Pr[X > (1 + \delta)\mu] < \left[ \frac{e^\delta}{(1 + \delta)^{1+\delta}} \right]^\mu$$

↑  
sum of independent 0-1 random variables  
is tightly centered on the mean

**Pf.** We apply a number of simple transformations.

- For any  $t > 0$ ,

$$\Pr[X > (1 + \delta)\mu] = \Pr[e^{tX} > e^{t(1+\delta)\mu}] \leq e^{-t(1+\delta)\mu} \cdot E[e^{tX}]$$

↑  
 $f(x) = e^{tx}$  is monotone in  $x$

↑  
Markov's inequality:  $\Pr[X > a] \leq E[X] / a$

- Now  $E[e^{tX}] = E[e^{t \sum_i X_i}] = \prod_i E[e^{tX_i}]$

↑  
definition of  $X$

↑  
independence

# Chernoff Bounds (above mean)

---

Pf. [ continued ]

- Let  $p_i = \Pr [X_i = 1]$ . Then,

$$E[e^{tX_i}] = p_i e^t + (1 - p_i) e^0 = 1 + p_i(e^t - 1) \leq e^{p_i(e^t - 1)}$$

↑  
for any  $\alpha \geq 0$ ,  $1 + \alpha \leq e^\alpha$

- Combining everything:

$$\Pr[X > (1 + \delta)\mu] \leq e^{-t(1+\delta)\mu} \prod_i E[e^{tX_i}] \leq e^{-t(1+\delta)\mu} \prod_i e^{p_i(e^t - 1)} \leq e^{-t(1+\delta)\mu} e^{\mu(e^t - 1)}$$

↑  
previous slide

↑  
inequality above

↑  
 $\sum_i p_i = E[X] \leq \mu$

- Finally, choose  $t = \ln(1 + \delta)$ . ■

## Chernoff Bounds (below mean)

---

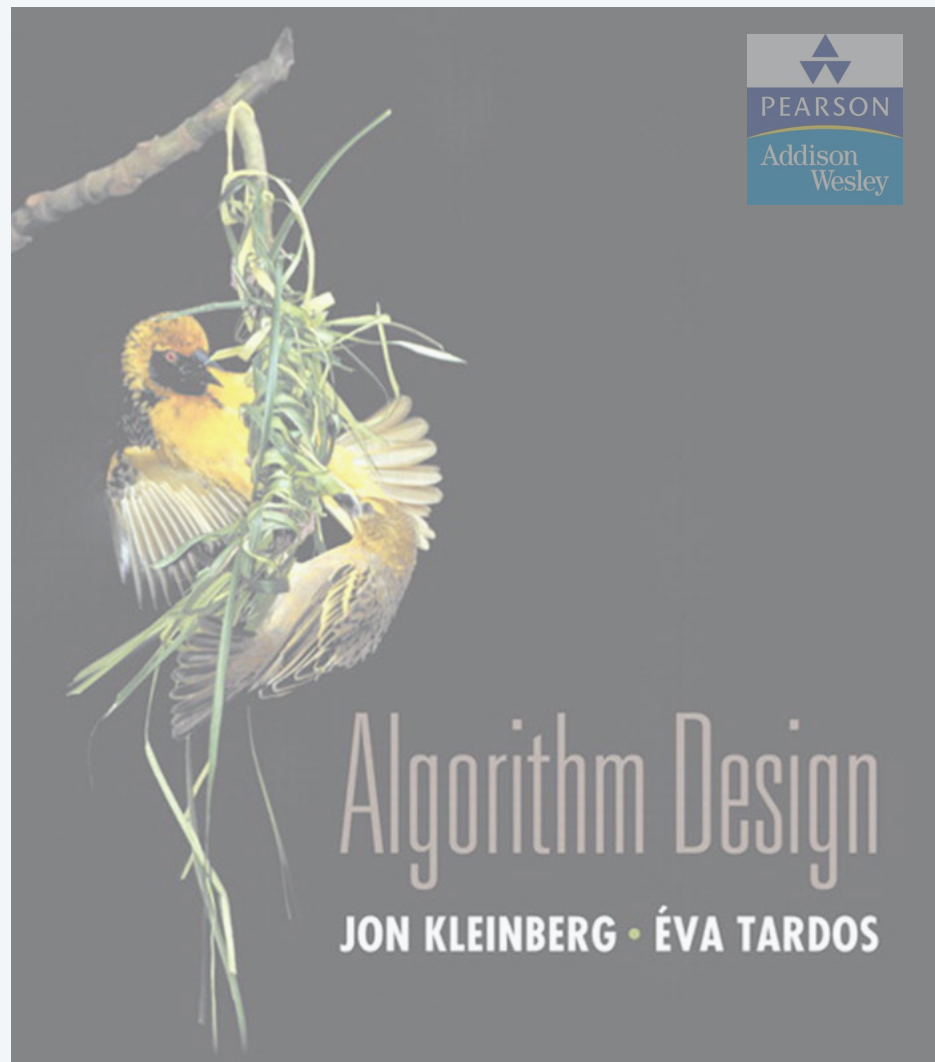
**Theorem.** Suppose  $X_1, \dots, X_n$  are independent 0-1 random variables.

Let  $X = X_1 + \dots + X_n$ . Then for any  $\mu \leq E[X]$  and for any  $0 < \delta < 1$ , we have

$$\Pr[X < (1 - \delta)\mu] < e^{-\delta^2 \mu / 2}$$

**Pf idea.** Similar.

**Remark.** Not quite symmetric since only makes sense to consider  $\delta < 1$ .



## 13. RANDOMIZED ALGORITHMS

---

- ▶ *contention resolution*
- ▶ *global min cut*
- ▶ *linearity of expectation*
- ▶ *max 3-satisfiability*
- ▶ *universal hashing*
- ▶ *Chernoff bounds*
- ▶ ***load balancing***



# Load balancing

---

**Load balancing.** System in which  $m$  jobs arrive in a stream and need to be processed immediately on  $m$  identical processors. Find an assignment that balances the workload across processors.

**Centralized controller.** Assign jobs in round-robin manner. Each processor receives at most  $\lceil m / n \rceil$  jobs.

**Decentralized controller.** Assign jobs to processors uniformly at random. How likely is it that some processor is assigned “too many” jobs?

# Load balancing

---

## Analysis.

- Let  $X_i$  = number of jobs assigned to processor  $i$ .
- Let  $Y_{ij} = 1$  if job  $j$  assigned to processor  $i$ , and 0 otherwise.
- We have  $E[Y_{ij}] = 1/n$ .
- Thus,  $X_i = \sum_j Y_{ij}$ , and  $\mu = E[X_i] = 1$ .
- Applying Chernoff bounds with  $\delta = c - 1$  yields  $\Pr[X_i > c] < \frac{e^{c-1}}{c^c}$ .
- Let  $\gamma(n)$  be number  $x$  such that  $x^x = n$ , and choose  $c = e \gamma(n)$ .

$$\Pr[X_i > c] < \frac{e^{c-1}}{c^c} < \left(\frac{e}{c}\right)^c = \left(\frac{1}{\gamma(n)}\right)^{e\gamma(n)} < \left(\frac{1}{\gamma(n)}\right)^{2\gamma(n)} = \frac{1}{n^2}$$

- Union bound  $\Rightarrow$  with probability  $\geq 1 - 1/n$  no processor receives more than  $e \gamma(n) = \Theta(\log n / \log \log n)$  jobs.

↑  
Bonus fact: with high probability,  
some processor receives  $\Theta(\log n / \log \log n)$  jobs

# Load balancing: many jobs

---

**Theorem.** Suppose the number of jobs  $m = 16 n \ln n$ . Then on average, each of the  $n$  processors handles  $\mu = 16 \ln n$  jobs. With high probability, every processor will have between half and twice the average load.

**Pf.**

- Let  $X_i, Y_{ij}$  be as before.
- Applying Chernoff bounds with  $\delta = 1$  yields

$$\Pr[X_i > 2\mu] < \left(\frac{e}{4}\right)^{16n \ln n} < \left(\frac{1}{e}\right)^{\ln n} = \frac{1}{n^2}$$

$$\Pr[X_i < \frac{1}{2}\mu] < e^{-\frac{1}{2} (\frac{1}{2})^2 16n \ln n} = \frac{1}{n^2}$$

- Union bound  $\Rightarrow$  every processor has load between half and twice the average with probability  $\geq 1 - 2/n$ . ■