

Disk Arrays

Single Disk-Array Disk

Single disk

Nomenclature

- n_r : number of stripe units reads by a read request
- n_w : number of stripe units modified by a write request
- λ_{array} : arrival rate of read requests to the disk array
- λ_{array}^w : arrival rate of write requests to the disk array
- λ_{disk} : arrival rate of read requests to any of the disk in the array
- λ_{disk}^w : arrival rate of write requests to any of the disk in the array
- N : number of physical disks in the array
- S_{array} : average service time at a disk array for read
- S_{array}^w : average service time at a disk array for write

21

- Fraction of read requests to the array that goes to any of the disks (uniform distribution to all disks):

$$\frac{n_r}{N}$$

- Arrival rate of disk requests at a disk:

$$\lambda_{disk} = \frac{n_r}{N} \lambda_{array} + \frac{n_w}{N} \lambda_{array}^w$$

where $\langle n_r, n_w \rangle$ is the number of read request to a disk as a result of \sim write requests

- Arrival rate of write requests to any disk in the array:

$$\lambda_{disk}^w = \frac{n_w}{N} \lambda_{array}^w$$

⇒ Average service time of read request at disk array:

$$S_{array} = \max_{i=1, \dots, N} [S_{disk,i}]$$

22

where $R_{disk,i}$ is the average response time of read requests at disk i :

⇒ Average service time of write request at disk array:

$$S_{array}^w = \max_{i=1, \dots, N} [R_{disk,i}^w] + \max_{i=1, \dots, N} [R_{disk,i}]$$

where $R_{disk,i}^w$ is the average response time of write requests at disk i .

Nomenclature

Sd: service time
SeekTime
Seek_{rand}
DiskSpeed
DiskRevolutionTime
RotationalLatency
BlockSize
TransferRate
TransferTime
ControllerTime
P_{miss}

Computing Disk Service Time

$$s_d = ContrTime +$$

$$P_{miss} (Seek + Latency + TransferT)$$

$$TransferT = \frac{BlockSize}{TransferRate}$$

6

Types of workloads

Random Workload:

10, 201, 15, 1023, 45, 39, 782

Sequential Workload:

4, 102, 103, 104, 105, 106, 25, 88, 32, 33, 34, 35, 36, 37, 38, 29, 15
 run length= 5 run length= 7

Random Workload:

$$P_{miss} = 1$$

$$RunLength = 1$$

$$SeekTime = S_{rand}$$

$$Latency = 1 / 2 \times RevolutionTime$$

Sequential Workload:

$$P_{miss} = 1 / RunLength$$

$$SeekTime = Seek_{rand} / RunLength$$

$$Latency = \frac{1 / 2 + (RunLength - 1) [(1 + U_d) / 2]}{RunLength} \times RevolutionTime$$

$$TransferTime = \frac{BlockSize}{10^6 \times TransferRate} \quad (3.3.19)$$

Random workloads:

$$\bar{S}_d = ControllerTime + Seek_{rand} + \frac{DiskRevolutionTime}{2} + TransferTime \quad (3.3.20)$$

Sequential workloads:

$$\bar{S}_d = ControllerTime + \frac{Seek_{rand}}{RunLength} + \frac{[1/2 + (RunLength - 1) (1 + U_d) / 2] \times DiskRevolutionTime}{RunLength} + TransferTime / RunLength \quad (3.3.21)$$

$$RunLength = NumberRequests / NumRuns$$

$$\begin{aligned} SeekTime &= \frac{NumRuns \times Seek_{rand}}{NumberRequests} \\ &= \frac{Seek_{rand}}{NumberRequests / NumRuns} \\ &= Seek_{rand} / RunLength. \end{aligned}$$