

Distributed Systems – 6 CFU module

A.A. 2020/2021

Instructor: Silvia Bonomi

Tentative Program

Topic	References
Introduction to Distributed Systems	[T1] - Chapter 1 [T2] - Chapter 1 and Chapter 2 [S]
Basic Abstractions <ul style="list-style-type: none">• Distributed Computations• Abstracting Processes• Abstracting Communications• Timing Assumptions• Abstracting Time	[T1] - Chapter 2 (except Section 2.3, Sections 2.4.5-2.4.7, Sections 2.6.6, Section 2.7) [S]
Clock Synchronization <ul style="list-style-type: none">• Internal and External synchronization• Cristian's algorithm• Berkley's algorithm• NTP• Happened-before relation• Application of scalar logical clocks to the mutual exclusion<ul style="list-style-type: none">◦ Lamport's algorithm◦ Ricarta-Agrawala's algorithm	[T2] - Chapter 14 (until Section 14.4 included) [S] [R1]
Broadcast <ul style="list-style-type: none">• Best Effort Broadcast• Reliable Broadcast• Uniform Reliable Broadcast• Probabilistic Broadcast	[T1] - Chapter 3 - from Section 3.1 to Section 3.4 (included) [T1] - Chapter 3 – Section 3.8 except Section 3.8.5 [S]
Consensus <ul style="list-style-type: none">• Regular Consensus• FLP Impossibility Result• Uniform Consensus• Paxos Algorithm	[T1] - Chapter 5, Sections 5.1.1, 5.1.2, 5.2.1, 5.2.2 [R2] [S]
Ordered Communication Primitives <ul style="list-style-type: none">• FIFO Broadcast	[T1] - Chapter 3 - from Section 3.9 (except 3.9.6)

<ul style="list-style-type: none"> • Causal Order Broadcast • Total Order Broadcast <ul style="list-style-type: none"> ○ TO Hierarchy 	[T1] - Chapter 6 – Section 6.1 [R3] [S]
Registers <ul style="list-style-type: none"> • Regular Register • Atomic Register • Message Passing Implementations • Transformation from (1, N) Regular to (1, N) Atomic 	[T1] - Chapter 4 - until Section 4.3 [S]
Software Replication <ul style="list-style-type: none"> • Linearizability • Primary-backup • Active Replication 	[R4] [S]
CAP Theorem	[R5] - [R6] [S]
Byzantine Fault Tolerance <ul style="list-style-type: none"> • Authenticated point-to-point links • Byzantine Broadcast • Byzantine Tolerant Registers • The Byzantine General Problem • State Machine Replication - PBFT 	[T1] - Chapter 2 – Section 2.4.6 [T1] - Chapter 3 – Section 3.10 (except 3.10.4), Section 3.11 [T1] - Chapter 4 – Sections 4.6 and 4.7 [R10] [S]
Broadcasting Information in Multi-hop Networks <ul style="list-style-type: none"> • Static networks • Dynamic networks 	[S] and references listed at the end of the slides
Blockchain and Distributed Ledgers	[S]

Main Text Book

[T1] - C. Cachin, R. Guerraoui and L. Rodrigues. Introduction to Reliable and Secure Distributed Programming, Springer, 2011

[S] – Slides from Lectures

Suggested Readings

[T2] - George Coulouris, Jean Dollimore and Tim Kindberg, Gordon Blair "Distributed Systems: Concepts and Design (5th Edition)". Addison - Wesley, 2012.

[R1] - Roberto Baldoni, Michel Raynal, "*Fundamentals of Distributed Computing: A Practical Tour of Vector Clock Systems*", IEEE Distributed Systems Online 3(2) (2002) <https://www.computer.org/csdl/mags/ds/2002/02/o2001.pdf>

[R2] - L. Lamport "Paxos Made Simple", <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/12/paxos-simple-Copy.pdf>

[R3] - Stefano Cimmino, Carlo Marchetti, Roberto Baldoni "A Guided Tour on Total Order Specifications" WORDS Fall 2003: 187-194

[R4] - Rachid Guerraoui and André Schiper: "Fault-Tolerance by Replication in Distributed Systems". In Proceedings of the 1996 Ada-Europe International Conference on Reliable Software Technologies (Ada-Europe '96).

[R5] - Brewer "CAP twelve years later: How the "rules" have changed" <http://ieeexplore.ieee.org/document/6133253/> (see NOTE above)

[R6] - Abadi "Consistency Tradeoffs in Modern Distributed Database System Design: CAP is Only Part of the Story" <http://ieeexplore.ieee.org/document/6127847/> (see NOTE above)

[R10] - Leslie Lamport, Robert Shostak, and Marshall Pease "The Byzantine Generals Problem " in ACM TOPLAS 1982 Available at <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/12/The-Byzantine-Generals-Problem.pdf>

NOTE: Use the Sapienza proxy to access these papers. Instruction on how to do it can be found here <https://web.uniroma1.it/sbs/easybixy/easybixy>