

# Distributed Systems

## Master of Science in Engineering in Computer Science

AA 2020/2021

---

LECTURE 12: TOTAL ORDER BROADCAST

# System model

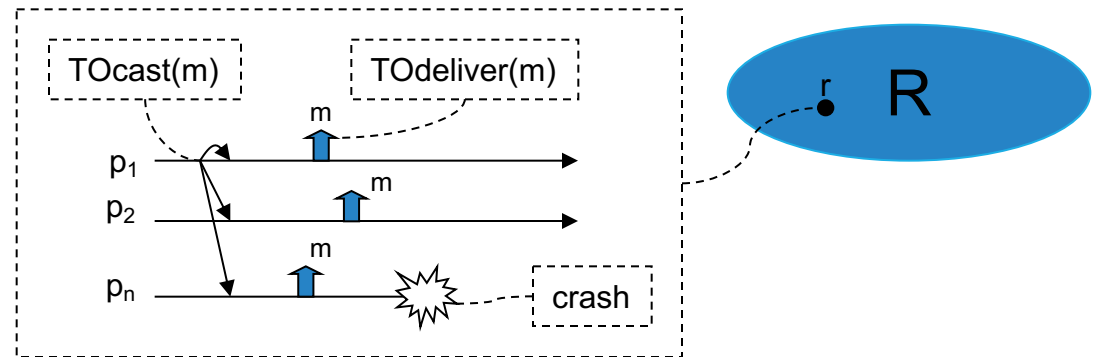
Static set of processes  $\Pi = \{p_1 \dots p_n\}$

Message passing over perfect channels (message exchanging between correct processes is reliable)

Asynchronous

Crash fault model for processes

We characterize the system in terms of its possible runs  $R$




# A few notation

---



Property  $P$ : predicate on the system, identifying a set of runs  $R_P \subseteq R$

- $P \Rightarrow P'$  iff  $R_P \subseteq R_{P'}$



Specification  $S(P_1, \dots, P_m)$ : logical and of  $m$  properties, identifying a set of runs  
 $R_S = R_{P_1} \cap \dots \cap R_{P_m} \subseteq R$

- $S \rightarrow S'$  iff  $R_S \subseteq R_{S'}$



# TO specifications

---

Total order specifications are usually composed by four properties:

- A Validity property guarantees that messages sent by correct processes will eventually be delivered at least by correct processes;
- An Integrity property guarantees that no spurious or duplicate messages are delivered;
- An Agreement property ensures that (at least correct) processes deliver the same set of messages;
- An Order property constrains (at least correct) processes delivering the same messages to deliver them in the same order.

# TO specifications

---

Total Order Broadcast =  $S(V, I, A, O)$

- **NUV**  
**UI**

- ~~A~~ = Agreement
- ~~O~~ = Order

→  $TO(A, O)$

Distinct specifications arise from distinct formulations of each property

- uniform vs non-uniform
- A uniform property imposes restrictions on the behavior of (at least) correct processes on the basis of events occurred in some process

# TO Specifications

---

Crash failure + Perfect channels  $\Rightarrow$

- **NUV**: if a correct process TOCAST a message  $m$  then some *correct* process will eventually deliver  $m$
- **UI**: For any message  $m$ , every process  $p$  delivers  $m$  at most once and only if  $m$  was previously TOCAST by some (correct or not) process.

# The Agreement property

---

## UNIFORM AGREEMENT (UA)

If a process (correct or not) TDelivers a message  $m$ , then all correct processes will eventually TDeliver  $m$

## NON-UNIFORM AGREEMENT (NUA)

If a correct process TDelivers a message  $m$ , then all correct processes will eventually TDeliver  $m$

## CONSTRAINS THE SET OF DELIVERED MESSAGES

**Correct processes always deliver the same set of messages  $M$**

**Each faulty process  $p$  delivers a set  $M_p$**


UA:  $M_p \subseteq M$

NUA:  $M_p$  *can be* s.t.  $M_p - M \neq \emptyset$

# The Agreement property


## UNIFORM AGREEMENT (UA)

If a process (correct or not) TODelivers a message  $m$ , then all correct processes will eventually TODeliver  $m$



## NON-UNIFORM AGREEMENT (NUA)

If a correct process TODelivers a message  $m$ , then all correct processes will eventually TODeliver  $m$





# The Ordering Property

---

## STRONG UNIFORM TOTAL ORDER (SUTO)

If some process  $TODelivers$  some message  $m$  before message  $m'$ , then a process  $TODelivers$   $m'$  only after it has  $TODelivered$   $m$ .



- same order
- same prefix of the set of delivered messages
- after an omission, disjoint sets of delivered messages

## WEAK UNIFORM TOTAL ORDER (WUTO)

If process  $p$  and process  $q$  both  $TODeliver$  messages  $m$  and  $m'$ , then  $p$   $TODelivers$   $m$  before  $m'$  if and only if  $q$   $TODelivers$   $m$  before  $m'$ .




- no restrictions on the set of delivered messages

# The Order Property

## STRONG UNIFORM TOTAL ORDER (SUTO)

If some process TODelivers some message  $m$  before message  $m'$ , then a process TODelivers  $m'$  only after it has TODelivered  $m$ .


**SUTO**



## WEAK UNIFORM TOTAL ORDER (WUTO)

If process  $p$  and process  $q$  both TODeliver messages  $m$  and  $m'$ , then  $p$  TODelivers  $m$  before  $m'$  if and only if  $q$  TODelivers  $m$  before  $m'$ .

**WUTO**



# The Order Property

---

SUTO and WUTO are uniform but they both have a non-uniform counterpart

## STRONG NON-UNIFORM TOTAL ORDER (SNUTO)

If some correct process TODElivers some message  $m$  before message  $m'$ , then a correct process TODElivers  $m'$  only after it has TODElivered  $m$ .



## WEAK NON-UNIFORM TOTAL ORDER (WNUTO)

If correct processes  $p$  and  $q$  both TODEliver messages  $m$  and  $m'$ , then  $p$  TODElivers  $m$  before  $m'$  if and only if  $q$  TODElivers  $m$  before  $m'$ .


# The Order property (2)

SUTO  $\Rightarrow$  WUTO

SNUTO  $\Rightarrow$  WNUTO





# TO specifications



TO(UA,SUTO)



- The strongest TO spec.

TO(NUA,SUTO)




# TO specifications (2)

TO(UA,WUTO)




TO(NUA,WUTO)

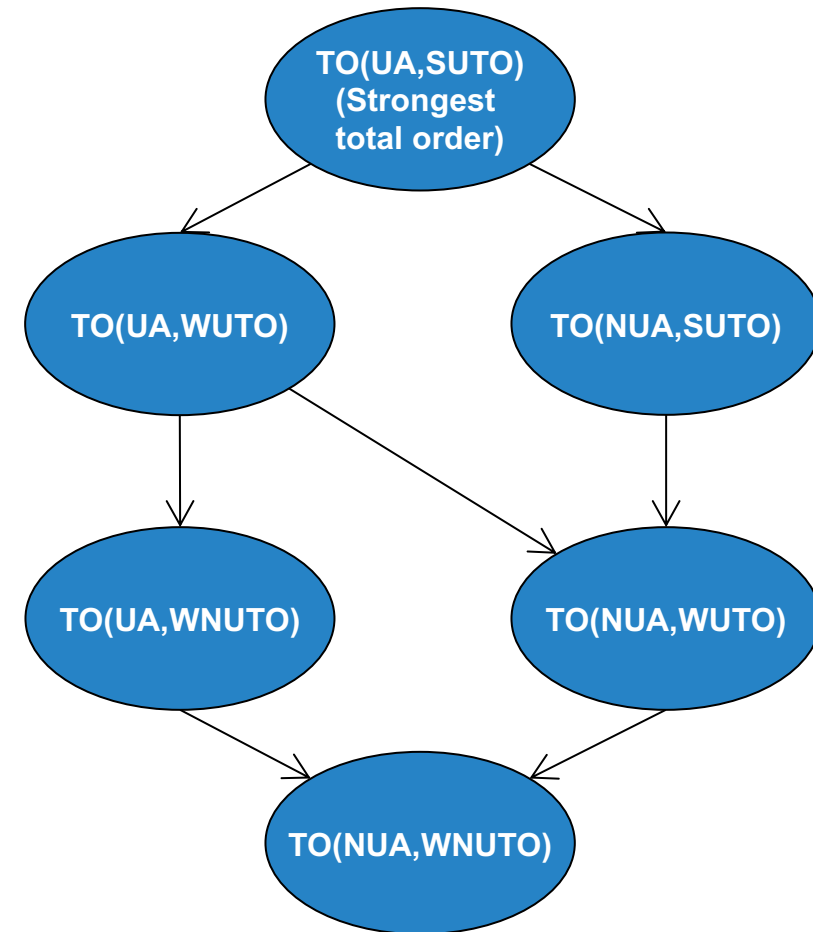



# TO specifications (3)

TO(UA,WNUTO)




TO(NUA,WNUTO)



# Total Order Implementation

---





# Total Order Algorithm

---

---

**Algorithm 6.1:** Consensus-Based Total-Order Broadcast

---

**Implements:**

TotalOrderBroadcast, **instance** *tob*.

**Uses:**

ReliableBroadcast, **instance** *rb*;  
Consensus (multiple instances).

**upon event**  $\langle tob, Init \rangle$  **do**

*unordered* :=  $\emptyset$ ;  
*delivered* :=  $\emptyset$ ;  
*round* := 1;  
*wait* := FALSE;

**upon event**  $\langle tob, Broadcast \mid m \rangle$  **do**

**trigger**  $\langle rb, Broadcast \mid m \rangle$ ;

**upon event**  $\langle rb, Deliver \mid p, m \rangle$  **do**

**if**  $m \notin delivered$  **then**  
*unordered* := *unordered*  $\cup \{(p, m)\}$ ;


**upon** *unordered*  $\neq \emptyset \wedge wait = \text{FALSE}$  **do**

*wait* := TRUE;  
Initialize a new instance *c.round* of consensus;  
**trigger**  $\langle c.round, Propose \mid unordered \rangle$ ;

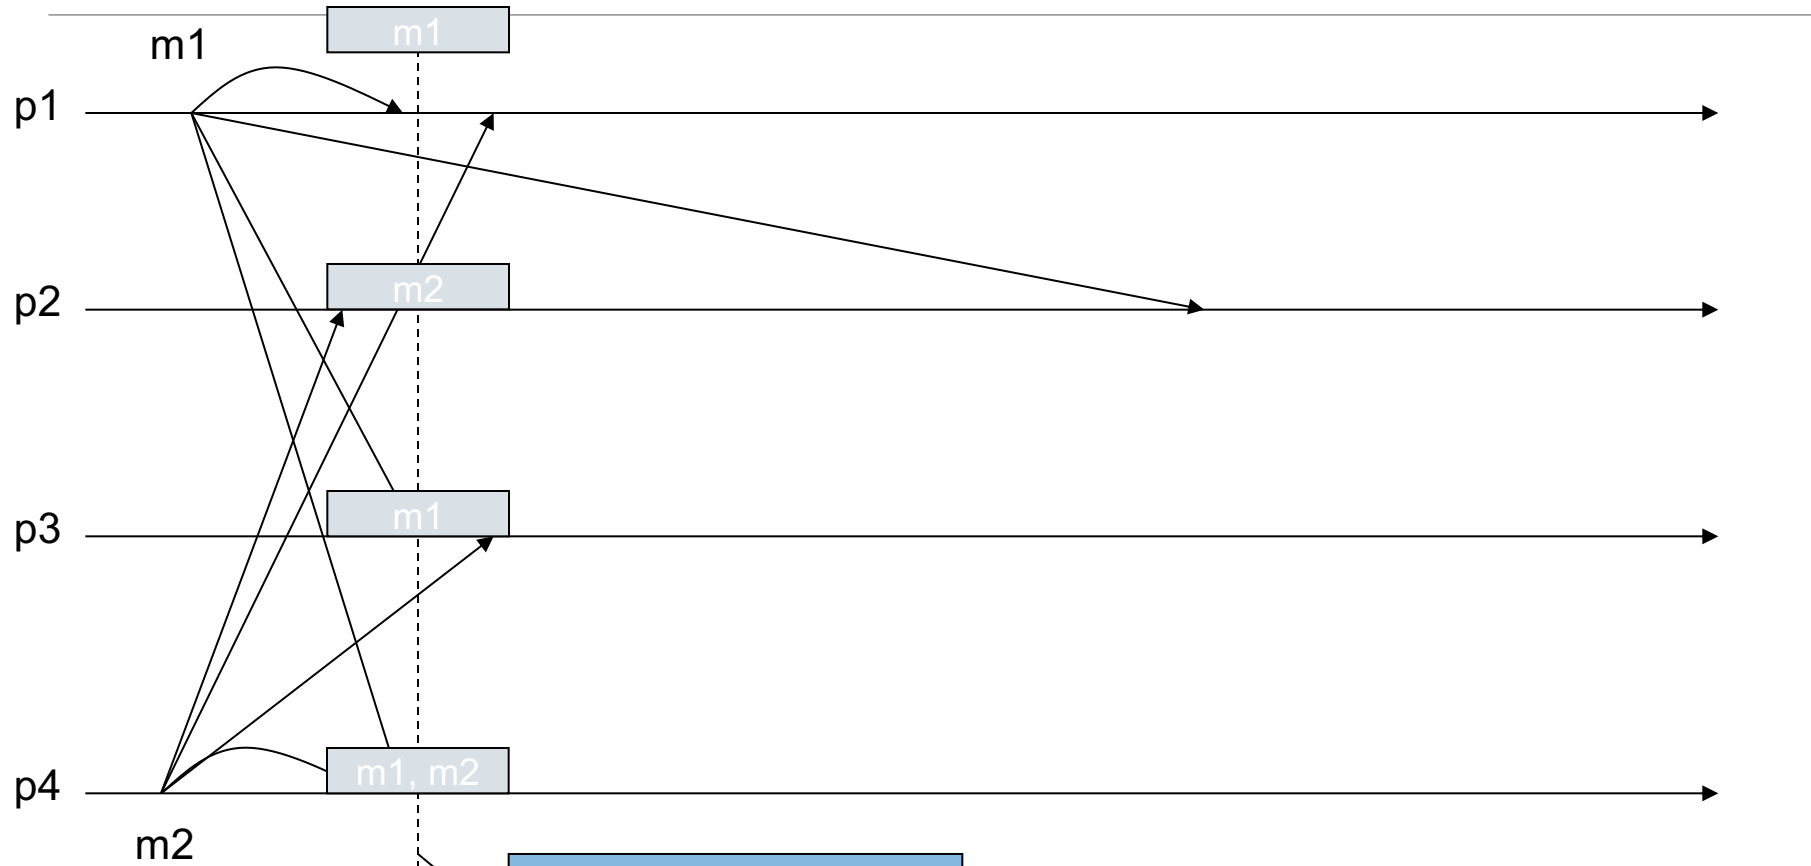
**upon event**  $\langle c.r, Decide \mid decided \rangle$  **such that**  $r = round$  **do**

**forall**  $(s, m) \in \text{sort}(decided)$  **do** // by the order in the resulting sorted list  
**trigger**  $\langle tob, Deliver \mid s, m \rangle$ ;  
*delivered* := *delivered*  $\cup decided$ ;  
*unordered* := *unordered*  $\setminus decided$ ;  
*round* := *round* + 1;  
*wait* := FALSE;

# Example




# Example




Each process  
proposes its  
*unordered* buffer

# Example




# Example



An instance of  
consensus is  
triggered

Each process  
proposes its  
*unordered* buffer

# Example



Consensus return m2, m3  
and all processes deliver  
m2 and then m3

# Exercise


---

Which TO specification is satisfied by this algorithm?

It depends from the assumptions about Reliable Broadcast and Consensus


Consensus Reliable Broadcast	Uniform	Non Uniform
Uniform		
Non Uniform		

# Example 1 (UC and URB)





# Example 2 (UC and URB)



# Uniform Consensus (UC) and Uniform Reliable Broadcast (URB)

---


Assuming both Consensus and Reliable Broadcast uniform we have

TO (UA, SUTO)

*Proof.*

- Due to URB all the processes (even the faults) deliver the same set of messages
- The unordered buffer contains the same set of messages for each process
  - All the processes will deliver the same set of messages (UA)
- Due to UC, all processes (even the faults) decide for the same list of messages
- Messages are sorted by a deterministic rule
  - All processes will deliver the messages in the same order

# Example (NUC and URB)



# Non Uniform Consensus (NUC) and Uniform Reliable Broadcast (URB)

---


Assuming both Consensus and Reliable Broadcast uniform we have

TO (UA, WNUTO)

*Proof.*

- Due to URB all the processes (even the faults) deliver the same set of messages
- The unordered buffer contains the same set of messages for each process
  - All the processes will deliver the same set of messages (UA)
- Due to NUC, all correct processes decide for the same list of messages
- Faulty processes can decide differently
  - All correct processes will deliver the messages in the same order
  - Faulty processes will deliver, just before a crash, a different sequence of messages

# Example (NUC and NURB)



# Non Uniform Consensus (NUC) and Non Uniform Reliable Broadcast (NURB)

---

Assuming both Consensus and Reliable Broadcast uniform we have

TO (NUA, WNUTO)

*Proof.*

- Due to NURB correct processes deliver the same set of messages
- Faulty processes can deliver other messages
  - Only correct processes will deliver the same set of messages (NUA)
- Due to NUC, all correct processes decide for the same list of messages
- Faulty processes can decide differently
  - All correct processes will deliver the messages in the same order
  - Faulty processes will deliver, just before a crash, a different sequence of messages


Consensus Reliable Broadcast	Uniform	Non Uniform
	Uniform	Non Uniform
Uniform	UA SUTO	UA WNUTO
Non Uniform		NUA WNUTO

## Exercise

Which specification is satisfied assuming UC and NURB?

# Exercise

Consider the run depicted in the figure:



1.  $TO(NUA, SUTO)$
2. Adding  $m_4$  to  $UA$   
Add  $m_4$  after  $m_3$  (WUTO)
3. Put  $m_3$  before  $m_1$

1. Which type of total ordering is satisfied by the run? Specify both the agreement and the ordering properties.
2. Modify the run in order to satisfy  $TO(UA, WUTO)$  but not  $TO(UA, SUTO)$
3. Modify the run in order to satisfy  $TO(NUA, WNUTO)$  but not  $TO(NUA, WUTO)$



# References

---

C. Cachin, R. Guerraoui and L. Rodrigues. Introduction to Reliable and Secure Distributed Programming, Springer, 2011

- Chapter 6 – Section 6.1

Stefano Cimmino, Carlo Marchetti, Roberto Baldoni "A Guided Tour on Total Order Specifications" WORDS Fall 2003: 187-194