

2 Beyond Relational Database Systems

- Column Stores
- NoSQL Database Systems
- Graph-DB

2.1 Column Stores¹

Table:

A	B	C
A1	B1	C1
A2	B2	C2
A3	B3	C3

row-wise storage:

A1	B1	C1	A2	B2	C2	A3	B3	C3
----	----	----	----	----	----	----	----	----

¹see for more information:

http://nms.csail.mit.edu/~stavros/pubs/tutorial2009-column_stores.pdf

2.1 Column Stores¹

Table:

A	B	C
A1	B1	C1
A2	B2	C2
A3	B3	C3

row-wise storage:

A1	B1	C1	A2	B2	C2	A3	B3	C3
----	----	----	----	----	----	----	----	----

column-wise storage:

A1	A2	A3	B1	B2	B3	C1	C2	C3
----	----	----	----	----	----	----	----	----

¹see for more information:

http://nms.csail.mit.edu/~stavros/pubs/tutorial2009-column_stores.pdf

2.1 Column Stores¹

Table:

A	B	C
A1	B1	C1
A2	B2	C2
A3	B3	C3

row-wise storage:

A1	B1	C1	A2	B2	C2	A3	B3	C3
----	----	----	----	----	----	----	----	----

column-wise storage:

A1	A2	A3	B1	B2	B3	C1	C2	C3
----	----	----	----	----	----	----	----	----

¹see for more information:

http://nms.csail.mit.edu/~stavros/pubs/tutorial2009-column_stores.pdf

Why Column-stores?

Advantages:

- Buffer management, only columns needed are read,
- data locality, column block iteration supports value aggregations,
- column-specific compression, values are from the same type,
- column insertion,
- late tuple materialization.

Disadvantage:

- Tuple insertion, and
- tuple reconstruction.

Compression and Suppression


Column compression

- Run-length encoding,
- Bit-vector encoding,
- Dictionary encoding,
- Reference encoding,
- Differential encoding.

Null suppression

- Position list encoding.

Run-length encoding



Bit-vector encoding

<i>someAttribute</i>		<i>someAttribute</i> 200	<i>someAttribute</i> 300	<i>someAttribute</i> 400
200		1	0	0
200		1	0	0
200		1	0	0
300		0	1	0
300	⇒	0	1	0
300		0	1	0
300		0	1	0
400		0	0	1
300		0	1	0
300		0	1	0

Dictionary encoding

someAttribute

https : //www.informatik.uni – freiburg.de/
 https : //wiki.dbpedia.org/
 https : //de.wikipedia.org/wiki/World_Wide_Web
 https : //www.informatik.uni – freiburg.de/
 https : //wiki.dbpedia.org/
 https : //de.wikipedia.org/wiki/World_Wide_Web
 https : //www.uni – freiburg.de/

⇒

someAttribute_encoded

1
 2
 3
 1
 2
 3
 4

Dictionary

1 → https : //www.informatik.uni – freiburg.de/
 2 → https : //wiki.dbpedia.org/
 3 → https : //de.wikipedia.org/wiki/World_Wide_Web
 4 → https : //www.uni – freiburg.de/

Reference encoding

<i>someAttribute</i>		<i>someAttribute_compressed</i> Reference : 200
200		0
205		5
208		8
199		-1
195	⇒	-5
300		#300
300		#300
200		0
210		10
205		5

Exeptions are indicated by #; necessary when difference to reference-value is larger then given range.

Differential encoding

<u><i>someAttribute</i></u>		<u><i>someAttribute_compressed</i></u> <i>Reference : 200</i>
200		0
205		5
208		3
199		-9
195	⇒	-4
300		#300
300		#300
200		5
210		10
205		-5

Exeptions are indicated by #; necessary when difference to previous value is larger then given range.

Null suppression: Position list encoding

someAttribute

NULL

https : //wiki.dbpedia.org/

NULL

NULL

https : //wiki.dbpedia.org/

NULL

https : //www.uni – freiburg.de/

NULL

NULL


 \Rightarrow *someAttribute_encoded*2 \rightarrow https : //wiki.dbpedia.org/5 \rightarrow https : //wiki.dbpedia.org/7 \rightarrow https : //www.uni – freiburg.de/*Metadata*

total number of rows : 9


number of nonnull value : 3

Column striping

How to store nested records in a column-store?



record-oriented



column-oriented

Lit.: Dremel: Interactive Analysis of Web-Scale Datasets By Sergey Melnik et al., Communications of the ACM, Vol. 54 No. 6, Pages 114-123

Dremel made simple with Parquet,

https://blog.twitter.com/engineering/en_us/a/2013/dremel-made-simple-with-parquet.html

Data model

Strongly-typed nested records with abstract syntax:

$$\tau = \text{dom} \mid \langle A_1 : \tau[* \mid ?], \dots, A_n : \tau[* \mid ?] \rangle$$

where τ is an atomic type or a record type.

- Atomic types in dom: integers, floating-point numbers, strings, etc.
- Records consist of one or multiple fields. Field i in a record has a name A_i and an optional multiplicity label.
- Repeated fields (*) may occur 0 or more times in a record.
- Optional fields (?) may be missing from the record.
- Otherwise, a field is required, i.e., must appear exactly once.

Data model

Strongly-typed nested records with abstract syntax:

$$\tau = dom \mid \langle A_1 : \tau[* \mid ?], \dots, A_n : \tau[* \mid ?] \rangle$$

where τ is an atomic type or a record type.

- Atomic types in *dom*: integers, floating-point numbers, strings, etc.
- Records consist of one or multiple fields. Field i in a record has a name A_i and an optional multiplicity label.
- Repeated fields $(*)$ may occur 0 or more times in a record.
- Optional fields $(?)$ may be missing from the record.
- Otherwise, a field is required, i.e., must appear exactly once.

Sample nested tree-shaped records and their schema.

```

DocId: 10      r1
Links
  Forward: 20
  Forward: 40
  Forward: 60
Name
  Language
    Code: 'en-us'
    Country: 'us'
  Language
    Code: 'en'
  Url: 'http://A'
Name
  Url: 'http://B'
Name
  Language
    Code: 'en-gb'
    Country: 'gb'

```

```

message Document {
  required int64 DocId;
  optional group Links {
    repeated int64 Backward;
    repeated int64 Forward; }
  repeated group Name {
    repeated group Language {
      required string Code;
      optional string Country; }
    optional string Url; }}

```

```

DocId: 20      r2
Links
  Backward: 10
  Backward: 30
  Forward: 80
Name
  Url: 'http://C'

```


Column-stripped representation: one column for each path from the root to a leaf.

DocId	Name.Url	Links.Forward	Links.Backward
value r d	value r d	value r d	value r d
10 0 0	http://A 0 2	20 0 2	NULL 0 1
20 0 0	http://B 1 2	40 1 2	10 0 2
	NULL 1 1	60 1 2	30 1 2
	http://C 0 2	80 0 2	

Name.Language.Code	Name.Language.Country
value r d	value r d
en-us 0 2	us 0 3
en 2 2	NULL 2 2
NULL 1 1	NULL 1 1
en-gb 1 2	gb 1 3
NULL 0 1	NULL 0 1

... Repetition level r and definition level d are used to be able to reconstruct the original nested record!

Column striping

DocId: 10 **r₁**
 Links
 Forward: 20
 Forward: 40
 Forward: 60
 Name
 Language
 Code: 'en-us'
 Country: 'us'
 Language
 Code: 'en'
 Url: 'http://A'
 Name
 Url: 'http://B'
 Name
 Language
 Code: 'en-gb'
 Country: 'gb'

```
message Document {
  required int64 DocId;
  optional group Links {
    repeated int64 Backward;
    repeated int64 Forward; }
  repeated group Name {
    repeated group Language {
      required string Code;
      optional string Country; }
    optional string Url; }}
```

DocId: 20 **r₂**
 Links
 Backward: 10
 Backward: 30
 Forward: 80
 Name
 Url: 'http://C'

DocId	Name.Url	Links.Forward	Links.Backward
value	value	value	value
10	http://A	20	NULL
20	http://B	40	10
	NULL	60	30
	http://C	80	

Column striping

```

DocId: 10      r1
Links
  Forward: 20
  Forward: 40
  Forward: 60
Name
  Language
    Code: 'en-us'
    Country: 'us'
  Language
    Code: 'en'
  Url: 'http://A'
Name
  Url: 'http://B'
Name
  Language
    Code: 'en-gb'
    Country: 'gb'

```

```

message Document {
  required int64 DocId;
  optional group Links {
    repeated int64 Backward;
    repeated int64 Forward; }
  repeated group Name {
    repeated group Language {
      required string Code;
      optional string Country; }
    optional string Url; }}

```

```

DocId: 20      r2
Links
  Backward: 10
  Backward: 30
  Forward: 80
Name
  Url: 'http://C'

```

Name.Language.Code

value	r	d
en-us	0	2
en	2	2
NULL	1	1
en-gb	1	2
NULL	0	1

Name.Language.Country

value	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
NULL	0	1

Record structure is captured by *repetition level* and *definition level*.

Repetition level:

The *repetition level* for a path denotes at which level in the path the last repetition occurred. Only repeated fields in the path are counted.

Allows to derive when to start a new list and at which level.

Definition level:

The *definition level* specifies how many fields in a path p that could be undefined (because they are optional or repeated) are present.
In other words, stores the level for which the field is null: from 0 at the root of the schema up to the maximum level for this column as can be derived from the schema.

- A field is defined: all its parents are defined too,
- a field is null: record the level at which it started being null.
- null values must be inserted in case no value for the entire path but only for a prefix of the path exists.

A required field is always defined and does not get a definition level.

Record structure is captured by *repetition level* and *definition level*.

Repetition level:

The *repetition level* for a path denotes at which level in the path the last repetition occurred. Only repeated fields in the path are counted.

Allows to derive when to start a new list and at which level.

Definition level:

The *definition level* specifies how many fields in a path p that could be undefined (because they are optional or repeated) are present.
In other words, stores the level for which the field is null: from 0 at the root of the schema up to the maximum level for this column as can be derived from the schema.

- A field is defined: all its parents are defined too,
- a field is null: record the level at which it started being null.
- null values must be inserted in case no value for the entire path but only for a prefix of the path exists.

A required field is always defined and does not get a definition level.

The *definition level* specifies how many fields in a path p that could be undefined (because they are optional or repeated) are present.

```
message ExampleDefinitionLevel {  
    optional group a {  
        optional group b {  
            optional string c;  
        }  
    }  
}
```

Value	Definition Level
a: null	0
a: { b: null }	1
a: { b: { c: null } }	2
a: { b: { c: "foo" } }	3 (actually defined)

The *definition level* specifies how many fields in a path p that could be undefined (because they are optional or repeated) are present.

```
message ExampleDefinitionLevel {  
  optional group a {  
    required group b {  
      optional string c;  
    }  
  }  
}
```

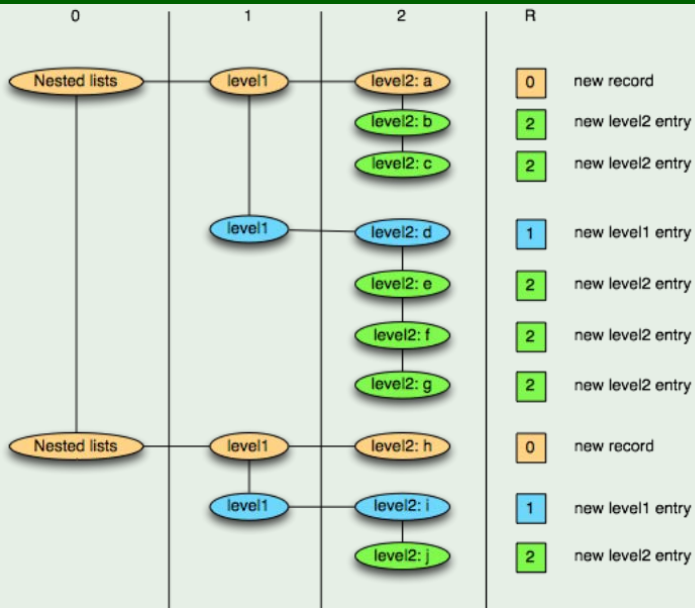
Value	Definition Level
a: null	0
a: { b: null }	Impossible, as b is required
a: { b: { c: null } }	1
a: { b: { c: "foo" } }	2 (actually defined)

The *repetition level* for a path denotes at which level in the path the last repetition occurred. Only repeated fields in the path are counted.

Schema:	Data: [[a,b,c],[d,e,f,g]],[[h],[i,j]]
<pre> message nestedLists { repeated group level1 { repeated string level2; } } </pre>	<pre> { level1: { level2: a level2: b level2: c }, level1: { level2: d level2: e level2: f level2: g } } { level1: { level2: h }, level1: { level2: i level2: j } } </pre>

- 0: a new record,
new level lists,
- 1: new level 1 list,
new level 2 list
- 2: new element
in level 2 list

Repetition level	Value
0	a
2	b
2	c
1	d
2	e
2	f
2	g
0	h
1	i
2	j



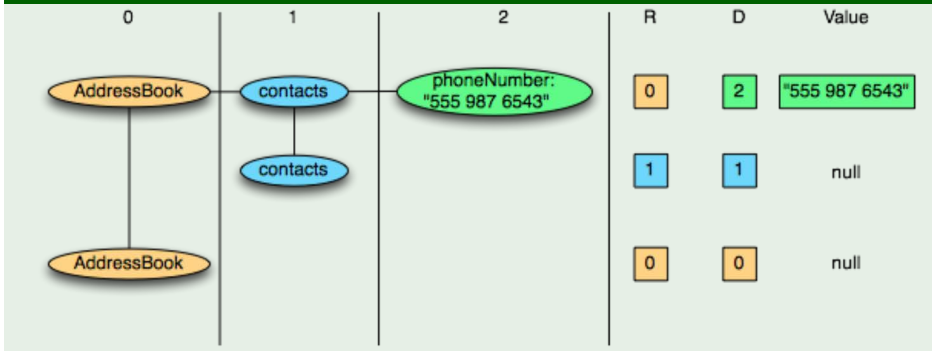
Striping and assembly

```
message AddressBook {  
  required string owner;  
  repeated string ownerPhoneNumbers;  
  repeated group contacts {  
    required string name;  
    optional string phoneNumber;  
  }  
}
```

considering only `contacts.phoneNumber`:

```
AddressBook {  
  contacts: {  
    phoneNumber: "555 987 6543"  
  }  
  contacts: {  
  }  
}  
AddressBook {  
}
```

Striping



Striped representation

contacts.phoneNumber		
R	D	Value
0	2	" 555 987 6543"
1	1	NULL
0	0	NULL

Striped representation

contacts.phoneNumber		
R	D	Value
0	2	"555 987 6543"
1	1	NULL
0	0	NULL

Assembly

R=0, D=2, Value = "555 987 6543":

R = 0 means a new record.

We recreate the nested records from the root
until the definition level (here 2)

D = 2 which is the maximum.

The value is defined and is inserted.

R=1, D=1:

R = 1 means a new entry in the contacts list at level 1.

D = 1 means contacts is defined but not phoneNumber,
so we just create an empty contacts.

R=0, D=0:

R = 0 means a new record. we create the nested records
from the root until the definition level

D = 0 => contacts is actually null, so we only
have an empty AddressBook


MonetDB²

In the past decades, advances in speed of commodity CPUs have far outpaced advances in RAM latency. Main-memory access has become a performance bottleneck for many computer applications: "memory wall."

- Redesign of the query execution model to better exploit pipelined CPU architectures and CPU instruction caches;
- use of columnar rather than row-wise data storage to better exploit CPU data caches;
- design of new cache-conscious query processing algorithms;
- design and automatic calibration of memory cost models to choose and tune these cache-conscious algorithms in the query optimizer.

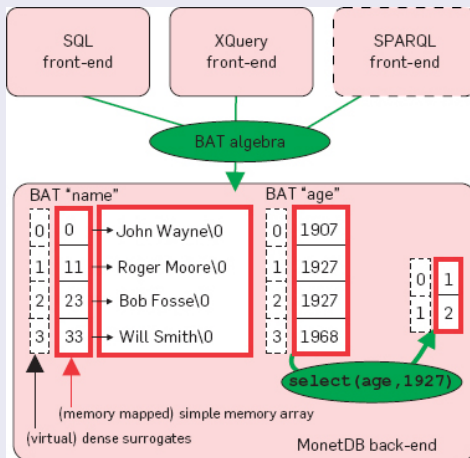
²Breaking the Memory Wall in MonetDB. Peter A. Boncz, Martin L. Kersten, Stefan Manegold. Communications of the ACM, Vol. 51 No. 12

CPU-Architecture




A translation lookaside buffer (TLB) is a memory cache that is used to reduce the time taken to access a user memory location. The TLB stores the recent translations of virtual memory to physical memory and can be called an address-translation cache.

MonetDB Architecture



BAT: Binary Association Table

Cash-conscious join: Partitioned hash-join



- Divide a relation U into H clusters using P sequential passes.
- Radix-clustering is on the lower B bits of the integer hash-value of a column. Each pass clusters tuples on B_p bits, starting with the leftmost bits, where $\sum_{p=1}^P B_p = B$.
- The number of clusters created by the Radix-Cluster is $H = \prod_{p=1}^P H_p$, where each pass subdivides each cluster into $H_p = 2^{B_p}$ new ones.
- When the algorithm starts, the entire relation is considered one single cluster, and is subdivided into $H_1 = 2^{B_1}$ clusters. The next pass takes these clusters and subdivides each into $H_1 = 2^{B_1}$ new ones, yielding $H_1 * H_2$ clusters in total, etc.

- Divide a relation U into H clusters using P sequential passes.
- Radix-clustering is on the lower B bits of the integer hash-value of a column. Each pass clusters tuples on B_p bits, starting with the leftmost bits, where $\sum_{p=1}^P B_p = B$.
- The number of clusters created by the Radix-Cluster is $H = \prod_{p=1}^P H_p$, where each pass subdivides each cluster into $H_p = 2^{B_p}$ new ones.
- When the algorithm starts, the entire relation is considered one single cluster, and is subdivided into $H_1 = 2^{B_1}$ clusters. The next pass takes these clusters and subdivides each into $H_1 = 2^{B_1}$ new ones, yielding $H_1 * H_2$ clusters in total, etc.

- Keep $H_x = 2^{B_x}$ smaller than the number of cache lines and the number of TLB entries, this will completely avoid both TLB and cache thrashing.
- After Radix-Clustering a column on B bits, all tuples that have the same B lowest bits in its column hash-value, appear consecutively in the relation, typically forming clusters of $|U|/2^B$ tuples.

- Divide a relation U into H clusters using P sequential passes.
- Radix-clustering is on the lower B bits of the integer hash-value of a column. Each pass clusters tuples on B_p bits, starting with the leftmost bits, where $\sum_{p=1}^P B_p = B$.
- The number of clusters created by the Radix-Cluster is $H = \prod_{p=1}^P H_p$, where each pass subdivides each cluster into $H_p = 2^{B_p}$ new ones.
- When the algorithm starts, the entire relation is considered one single cluster, and is subdivided into $H_1 = 2^{B_1}$ clusters. The next pass takes these clusters and subdivides each into $H_1 = 2^{B_1}$ new ones, yielding $H_1 * H_2$ clusters in total, etc.

- Keep $H_x = 2^{B_x}$ smaller than the number of cache lines and the number of TLB entries, this will completely avoid both TLB and cache thrashing.
- After Radix-Clustering a column on B bits, all tuples that have the same B lowest bits in its column hash-value, appear consecutively in the relation, typically forming clusters of $|U|/2^B$ tuples.

2.2 Not only SQL (NoSQL) Database System

Not "No SQL", but "Not only SQL"

Non-relational databases are the answer to new data management challenges:

- *Complex structures*: e.g. graphs,
- *Schema independence*: e.g. to support flexible integration of data from different sources into the same store,
- *Sparseness*: e.g. do not store NULL-values,
- *Self-descriptiveness*: e.g. attach meta-data to individual values,
- *Variability*: e.g. constantly changing data and schemata,
- *Scalability*: data is processed in a distributed fashion,
- *Volume*: large volumes have to be processed.

A NoSQL database system may be non-relational, support non-standard query languages, support schema evolution and independence, support data distribution with possibly a weaker form of consistency.

2.2 Not only SQL (NoSQL) Database System

Not "No SQL", but "Not only SQL"

Non-relational databases are the answer to new data management challenges:

- *Complex structures*: e.g. graphs,
- *Schema independence*: e.g. to support flexible integration of data from different sources into the same store,
- *Sparseness*: e.g. do not store NULL-values,
- *Self-descriptiveness*: e.g. attach meta-data to individual values,
- *Variability*: e.g. constantly changing data and schemata,
- *Scalability*: data is processed in a distributed fashion,
- *Volume*: large volumes have to be processed.

A NoSQL database system may be non-relational, support non-standard query languages, support schema evolution and independence, support data distribution with possibly a weaker form of consistency.

Key-Value Store

Key-value pair: $\langle \text{key}, \text{value} \rangle$.

- *key*: string, unique identifier,
- *value*: string.

Prototype of a schema-less database: arbitrary key-value pairs can be stored and retrieved; *simple, but quick*.

- Disadvantage: values cannot be searched, no query language;
- Advantage: data can be easily distributed.

Key-Value Store

Key-value pair: $\langle \text{key}, \text{value} \rangle$.

- *key*: string, unique identifier,
- *value*: string.

Prototype of a schema-less database: arbitrary key-value pairs can be stored and retrieved; *simple, but quick*.

- Disadvantage: values cannot be searched, no query language;
- Advantage: data can be easily distributed.

Values may have a type other than *string*, e.g. collection types *list* or *array*, even nested.

⇒ *Document databases*; data formats XML or JSON: a key-value pair may be the value of another key-value pair.

Document database systems, e.g. MongoDB, CouchDB provide additional processing functionality.

Key-Value Store

Key-value pair: $\langle \text{key}, \text{value} \rangle$.

- *key*: string, unique identifier,
- *value*: string.

Prototype of a schema-less database: arbitrary key-value pairs can be stored and retrieved; *simple, but quick*.

- Disadvantage: values cannot be searched, no query language;
- Advantage: data can be easily distributed.

Values may have a type other than *string*, e.g. collection types *list* or *array*, even nested.

⇒ *Document databases*; data formats XML or JSON: a key-value pair may be the value of another key-value pair.

Document database systems, e.g. MongoDB, CouchDB provide additional processing functionality.

Wide-column Store

Example Google BigTable/Apache HBase: extensible record stores

Data is stored in tables, where records have the ability to hold very large numbers of dynamic columns. Tables are typically not normalized!

- Column management implements the concept of *column families*.
- Akin to key-value stores, *multidimensional* keys are mapped to a value.
- The column names as well as the record keys are not fixed; a record can have billions (?) of columns.
- Characteristic of being schema-free is shared with document stores, however the implementation is very different.
- Wide column stores must not be confused with the column oriented storage in some relational systems.

Wide-column Store

Example Google BigTable/Apache HBase: extensible record stores

Data is stored in tables, where records have the ability to hold very large numbers of dynamic columns. Tables are typically not normalized!


- Column management implements the concept of *column families*.
- Akin to key-value stores, *multidimensional* keys are mapped to a value.
- The column names as well as the record keys are not fixed; a record can have billions (?) of columns.
- Characteristic of being schema-free is shared with document stores, however the implementation is very different.
- Wide column stores must not be confused with the column oriented storage in some relational systems.

2.3 Graph-DB

(A) Edge-labelled Graphs: RDF

To draw an RDF graph, for each Subject-Predicate-Object-triple introduce $s \xrightarrow{P} o$.

Subject	Predicate	Object
A	likes	I_1
A	likes	I_2
C	likes	I_2
A	follows	B
B	follows	C
B	follows	D
C	follows	D



RDF-storage by Relational Databases.

Different approaches:

- Giant Tripel-Table
- Vertical Partitioning
- Property-Table

Other options: Native RDF-stores, Graph-Databases.

Triple Table

Triple-Table

Subj.	Prop.	Obj.
ID1	type	BookType
ID1	title	"XYZ"
ID1	author	"Fox, Joe"
ID1	copyright	"2001"
ID2	type	CDType
ID2	title	"ABC"
ID2	artist	"Orr, Tim"
ID2	copyright	"1985"
ID2	language	"French"
ID3	type	BookType
ID3	title	"MNO"
ID3	language	"English"
ID4	type	DVDType
ID4	title	"DEF"
ID5	type	CDType
ID5	title	"GHI"
ID5	copyright	"1995"
ID6	type	BookType
ID6	copyright	"2004"

Property Table

Property Table Approach

Property Table

Subj.	Type	Title	copyright
ID1	BookType	"XYZ"	"2001"
ID2	CDType	"ABC"	"1985"
ID3	BookType	"MNP"	NULL
ID4	DVDType	"DEF"	NULL
ID5	CDType	"GHI"	"1995"
ID6	BookType	NULL	"2004"

Left-Over Triples

Subj.	Prop.	Obj.
ID1	author	"Fox, Joe"
ID2	artist	"Orr, Tim"
ID2	language	"French"
ID3	language	"English"

Clustered property table

Class: BookType

Subj.	Title	Author	copyright
ID1	"XYZ"	"Fox, Joe"	"2001"
ID3	"MNP"	NULL	NULL
ID6	NULL	NULL	"2004"

Class: CDType

Subj.	Title	Artist	copyright
ID2	"ABC"	"Orr, Tim"	"1985"
ID5	"GHI"	NULL	"1995"

Left-Over Triples

Subj.	Prop.	Obj.
ID2	language	"French"
ID3	language	"English"
ID4	type	DVDType
ID4	title	"DEF"

Property-class table

Discussion: Triple vs. Property Tables

- Triple table:

Advantage: Simplicity.

Problems: Multi-way self-joins may result, which, when not selective, are very expensive.

- Property table:

Advantage: Subject-Subject self-joins are reduced; attribute-typing possible.

Problems:

- Queries with unspecified property values are problematic.
- NULL-values may yield sparse tables.
- What to do with multi-valued attributes?

Discussion: Triple vs. Property Tables

- Triple table:

Advantage: Simplicity.

Problems: Multi-way self-joins may result, which, when not selective, are very expensive.

- Property table:

Advantage: Subject-Subject self-joins are reduced; attribute-typing possible.

Problems:

- Queries with unspecified property values are problematic.
- NULL-values may yield sparse tables.
- What to do with multi-valued attributes?

Vertically partitioned approach

Table per property

Type

ID1	BookType
ID2	CDType
ID3	BookType
ID4	DVDType
ID5	CDType
ID6	BookType

Author

ID1	"Fox, Joe"
-----	------------

Title

ID1	"XYZ"
ID2	"ABC"
ID3	"MNO"
ID4	"DEF"
ID5	"GHI"

Artist

ID2	"Orr, Tim"
-----	------------

Copyright

ID1	"2001"
ID2	"1985"
ID5	"1995"
ID6	"2004"

Language

ID2	"French"
ID3	"English"

Author

ID1	"Fox, Joe"
ID1	"Green, John"

Multi-valued attributes.

Vertically partitioned approach

Each table is sorted by subject. Object column can optionally be indexed.

- Advantages:

- Support for multi-valued attributes.
- Support for heterogenous records.
- Reduced I/O-costs: only that what indeed is needed will be accessed.
- No property-clustering required.

- Problems:

- Increased number of joins.
- Queries with unspecified properties.
- Inserts.

Vertically partitioned approach

Each table is sorted by subject. Object column can optionally be indexed.

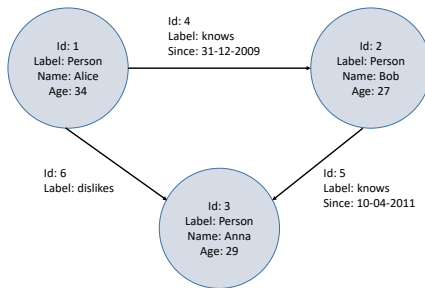
- Advantages:

- Support for multi-valued attributes.
- Support for heterogenous records.
- Reduced I/O-costs: only that what indeed is needed will be accessed.
- No property-clustering required.

- Problems:

- Increased number of joins.
- Queries with unspecified properties.
- Inserts.

(B) Property Graphs



Storing in relational tables

Nodes	NodeID	NodeLabel
	1	Person
	2	Person
	3	Person

PersonAttributes	NodeID	Name	Age
	1	Alice	34
	2	Bob	27
	3	Charlene	29

Edges	EdgeID	EdgeLabel	Source	Target
	4	knows	1	2
	5	knows	2	3
	6	dislikes	1	3

KnowsAttributes	EdgeID	Since
	4	31-12-2009
	5	10-04-2011

Native Graph Storage

- Property stores: key-value pairs;
- Graph structure: by node and relationship stores, sophisticated highly efficient linked list structures.
Akin adjacency lists.