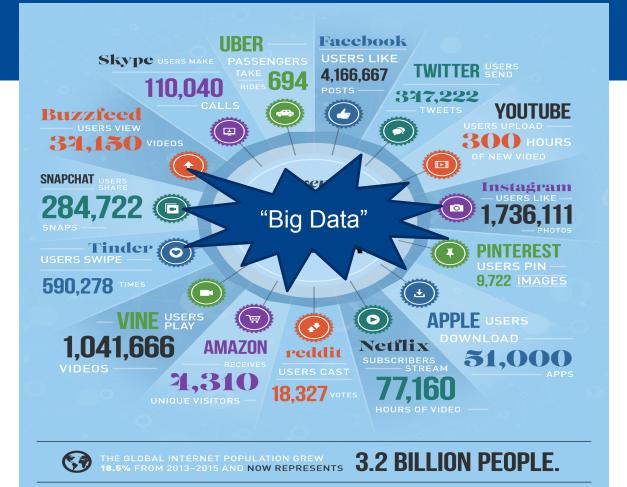
Distributed Processing with Computer Clusters I

Advanced Databases and Information Systems SS18

Albert-Ludwigs-Universität Freiburg



Prof. Dr. Georg Lausen Patrick Philipp





Examples of Big Data

- Google processes 3.5 billion requests per day, storing 10 Exabytes of data.
- Amazon hosts the most servers of any company (est. 1,400,000 servers), having Google and Microsoft close behind. Amazon Web Services (AWS) are used by 60,000 companies and field more than 650,000 requests every second. It is estimated that 1/3 of all Internet users use a website hosted on AWS daily, and that 1% of all Internet traffic goes through Amazon.
- **Facebook** collects 500 terabytes of data daily, including 2.5 billion pieces of content, 2.7 billion likes and 300 million photos.
- 90% of all the data in the world was produced in the last 2 years.
- It is estimated that 40 Zettabytes (40,000 Exabytes) of data will be created by 2020.

See http://cloudtweaks.com/wp-content/uploads/2015/03/bigdata-facts.png

Publicly Available Big Data Sets

- Common crawl https://commoncrawl.org/
- Wikipedia edit history
- AWS Public Datasets
 https://aws.amazon.com/de/datasets/
- Google BigQuery Public Datasets https://cloud.google.com/bigquery/public-data/
- · ...

What is Big Data?

THE PROPERTY OF THE PROPERTY O

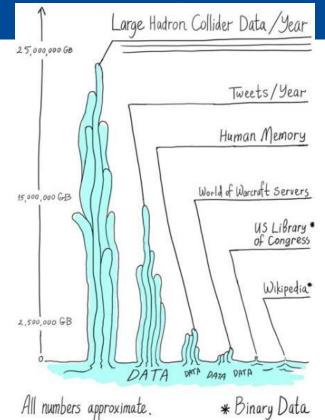
- 'Big data' is similar to 'Small data', but bigger
 - and often less structured
 - often characterized by the 3 (or 4) V's:
 - Volume
 - Variety
 - Velocity
 - (Veracity)

What is Big Data? (2)

- ... but it requires different approaches:
 - techniques, tools, architectures
- ... with the aim to solve new problems
 - or old problems in a better way.

What is Big Data and What not?





from: Symmetry Magazine

Big Data related tools





http://www.bigdata-startups.com/open-source-tools/

Distributed processing



- Distributed processing of Big Data requires non-standard programming models
 - beyond single machines or traditional parallel programming models (like Message Passing Interface, MPI)
 - the aim is to simplify complex programming tasks
- One of the most popular frameworks is MapReduce
 - suitable for commodity hardware to reduce costs

MapReduce: Why?



- On 1 node
 - Scanning @ 50 MB/s ≠ **23 days**
- On 1000 node cluster
 - Scanning @ 50 MB/s = 33 min

Current development

- Companies often can't cope with logged user behavior and throw away data after some time → lost opportunities
- Growing cloud-computing capacities
- Price/performance advantage of low-end servers increases to about a factor of twelve

Big Ideas behind MapReduce

STATE OF THE PARTY OF THE PARTY

- Scale "out", not "up"
 - Large number of commodity servers
- Assume failures are common
 - In a cluster of 10000 servers, expect 10 failures a day.
- Move processing to the data
 - Take advantage of data locality and avoid to transfer large datasets through the network
- Process data sequentially and avoid random access
 - Random disk access causes seek times
- Hide system-level details from the application developer
 - Developers can focus on their problems instead of dealing with distributed programming issues
- Seamless scalability
 - Scaling "out" improves the performance of an algorithm without any modifications

DFS & MapReduce

- Distributed File System (Google: GFS, Hadoop: HDFS)
 - Data is split into equally sized blocks and stored distributed
 - Fault tolerance by replication
 - Very large files / write-once, read-many pattern

MapReduce

- Algorithms expressed as a sequence of Map() and Reduce() functions
- Automatically parallelized by the framework

Advantages

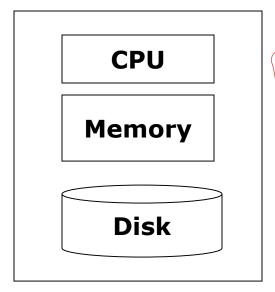
- Partitioning + distribution of data
- Parallelization and assigning of task
- Scalability, fault-tolerance, scheduling,...

That all is done automatically



Single-node architecture





Machine Learning, Statistics

"Classical" Data Mining

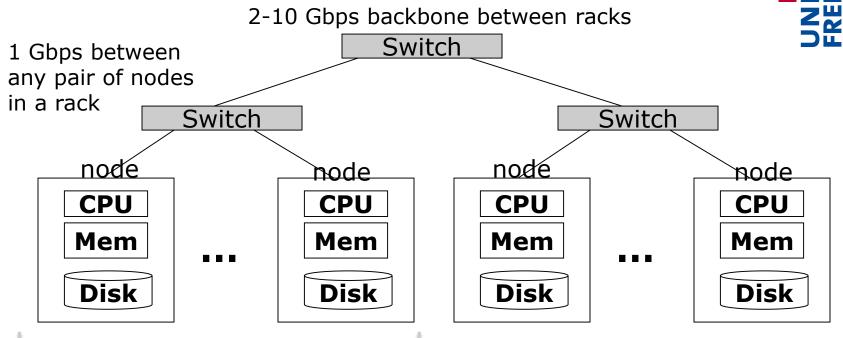
Commodity Clusters

THE CONTROLLED TO THE PARTY OF THE PARTY OF

- Data mining on a single server not enough
- Standard architecture emerging:
 - Cluster of commodity Linux nodes
 - Gigabit ethernet interconnect
- How to organize computations on this architecture?
 - Mask issues such as hardware failure

Cluster Architecture





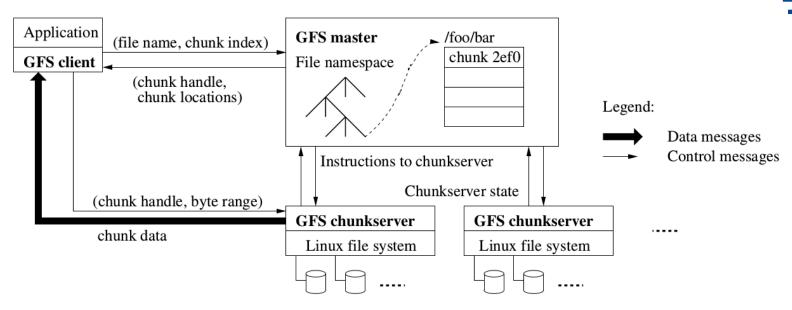
rack
Each rack contains 16-64 nodes

Google File System (GFS)

- very first proposal of a DFS
- Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. "The Google file system." SOSP'2003.
- Available online: http://blough.ece.gatech.edu/6102/presentations/gfssosp2003.pdf
- designed for Web search / Matrix-vector multiplication

GFS Architecture

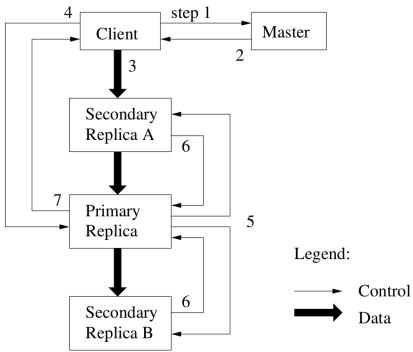




from: Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. "The Google file system." SOSP'2003.

13.06.2018

Writing Workflow in GFS



from: Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. "The Google file system." SOSP'2003.

13.06.2018

DFS Usage Pattern



- Typical usage pattern for a DFS (and MapReduce)
 - Huge files (100s of GB to TB)
 - Data is rarely updated in place
 - Reads and appends are common

Basic Components of a DFS



Chunk Servers

- File is split into contiguous chunks
- Typically each chunk is nowadays between 64 and 256MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

2. Master node

- a.k.a. Name Nodes in HDFS
- Stores metadata (about files)
- Might be replicated

Basic Components of a DFS



- 3. Client library for file access
 - Talks to master to find chunk servers
 - Connects directly to chunk servers to access data

13.06.2018 21

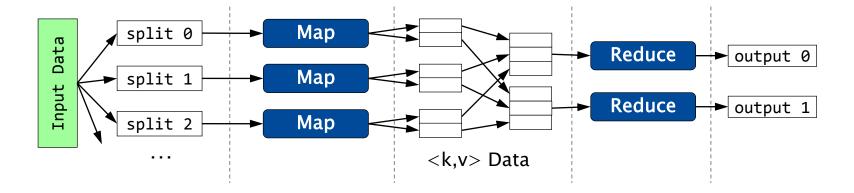
MapReduce



- The key idea of the MapReduce approach:
 - A target problem needs to be parallelizable
 - First, the problem gets split into a set of smaller problems (Map step)
 - 2. Next, smaller problems are solved in a parallel way
 - 3. Finally, a set of solutions to the smaller problems get synthesized into a solution of the original problem (Reduce step)

MapReduce overview





UNI FREIBURG

MapReduce example:

Counting words in documents

Google Maps charts new territory into businesses

Google selling new tools for businesses to build their own maps

Google promises consumer experience for businesses with Maps Engine Pro

Google is trying to get its Maps service used by more businesses



Google	4
Maps	4
Businesses	4
Engine	1
Charts	1
Territory	1
Tools	1

Map task

Google Maps charts new territory into businesses

Google selling new tools for businesses to build their own maps_____

Google promises consumer experience for businesses with Maps Engine Pro

Google is trying to get its Maps service used by more businesses

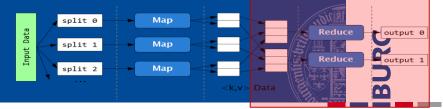


Google	2	Y
Charts	1	
Maps	2	
Territory	1	



Google	2
Businesses	2
Maps	2
Service	1

Reduce task



Google	2
Google	2
Maps	2
Maps	2



Google	4	30
Maps	4	

Businesses	2
Businesses	2
Charts	1
Territory	1



Businesses	4
Charts	1
Territory	1

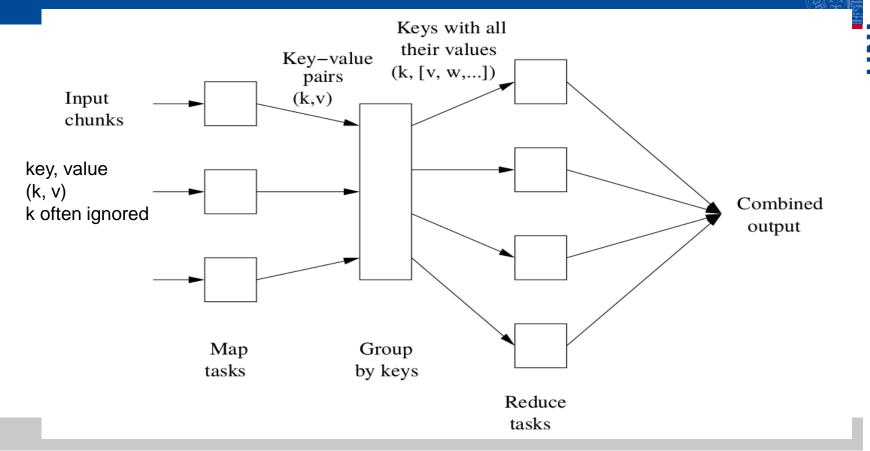
MapReduce: General idea

THE PROPERTY OF THE PROPERTY O

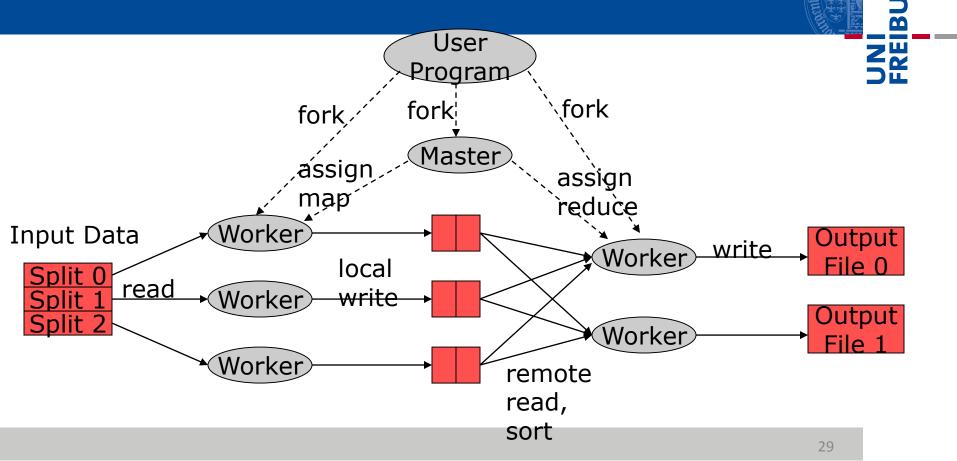
- Two functions: map() and reduce()
- Input/Output: always key-value-pairs
- => allow composition of MapReduce processes

28

Distributed Execution Overview



Distributed Execution in more Detail



Data Flow



- Input/final output are stored on a distributed file system
 - Scheduler tries to schedule map tasks "close" to physical storage location of input data
- Intermediate results are stored on local file system of map- and reduce-workers
- Output is often input to another MapReduce task

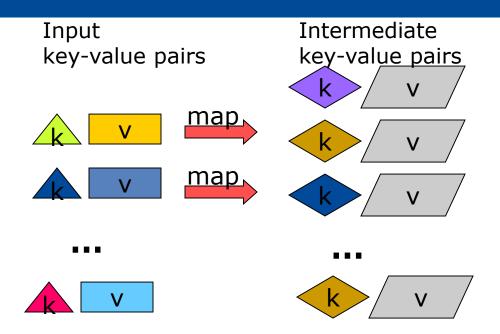
MapReduce: General procedure

- Input: a set of key/value pairs
- User supplies two functions:

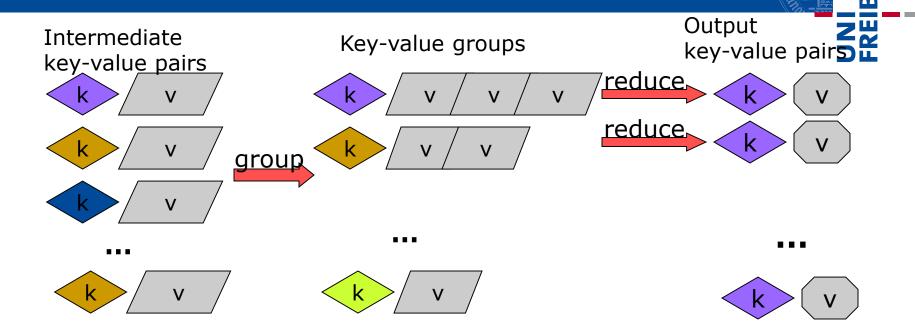
- (k1,v1) is an intermediate key/value pair
- Output is the set of (k1,v2) pairs

MapReduce: The Map Step





MapReduce: The Reduce Step



Example: Word Count



- We have a large file of words, one word to a line
- Count the number of times each distinct word appears in the file
- Sample application:
 - analyze web server logs to find popular URLs
 - analyze Web page contents
- Great: This is naturally parallelizable

Word Count using MapReduce

```
map(key, value):
// key: document name; value: text of document
   for each word w in value:
        emit(w, 1)
 reduce(key, values):
 // key: a word; value: an iterator over counts
          result = 0
          for each count v in values:
                  result += v
         emit(result)
```

Note: Hadoop versions



- Two versions of the APIs
 - The **old API** is in the org.apache.hadoop.**mapred** package
 - The **new API** is in theorg.apache.hadoop.**mapreduce** package
- Two versions of the runtime
 - MRv1 (MapReduce)
 - MRv2 (MapReduce 2, YARN)
- Never mix the old and new API

```
import org.apache.hadoop.*;
public class WordCountDriver {
  public static void main(String[] args) {
    JobClient client = new JobClient();
    JobConf conf = new JobConf(WordCountDriver.class);
    // specify output types
    conf.setOutputKeyClass(Text.class);
    conf.setOutputValueClass(IntWritable.class);
    // specify input and output dirs
    FileInputFormat.addInputPath(conf, new Path("input"));
    FileOutputFormat.setOutputPath(conf, new Path("output"));
```

13.06.2018

```
// specify a mapper
conf.setMapperClass(WordCountMapper.class);
// specify a reducer
conf.setReducerClass(WordCountReducer.class);
conf.setCombinerClass(WordCountReducer.class);
client.setConf(conf);
try {
  JobClient.runJob(conf);
} catch (Exception e) {
  e.printStackTrace();
```

13.06.2018

```
import org.apache.hadoop.*;
public class WordCountMapper extends MapReduceBase implements
Mapper < Long Writable, Text, Text, Int Writable >
        private final IntWritable one = new IntWritable(1);
        private Text word = new Text();
        @Override
        public void map(LongWritable key, Text value, OutputCollector<Text,</pre>
IntWritable> output, Reporter reporter) throws IOException
                 String line = value.toString();
                 StringTokenizer itr = new StringTokenizer(line.toLowerCase());
                 while(itr.hasMoreTokens()) {
                   word.set(itr.nextToken());
                   output.collect(word, one);
```

13.06.2018

```
import org.apache.hadoop.mapred.*;
public class WordCountReducer extends MapReduceBase
    implements Reducer<Text, IntWritable, Text, IntWritable> {
  public void reduce (Text key, Iterator values,
      OutputCollector output, Reporter reporter) throws IOException {
    int sum = 0:
    while (values.hasNext()) {
      IntWritable value = (IntWritable) values.next();
      sum += value.get(); // process value
    output.collect(key, new IntWritable(sum));
```

13.06.2018 40

Coordination

THE REPORT OF THE PARTY OF THE

- Master as "central unit" for coordination
 - Task status: (idle, in-progress, completed)
 - Idle tasks get scheduled as workers become available
 - When a map task completes, it sends the master the location and sizes of its R intermediate files, one for each reducer
 - Master pushes this information to reducers
- Master checks periodically if workers have failures

Failures



- Different kinds of failure possible:
- Map worker failure
 - Map tasks completed or in-progress at worker are reset to idle
 - Reduce workers are notified when task is rescheduled on another worker
- Reduce worker failure
 - Only in-progress tasks are reset to idle
- Master failure
 - MapReduce task is aborted and client is notified

How many Map and Reduce jobs?



- M map tasks, R reduce tasks
- Rule of thumb:
 - Make M and R much larger than the number of nodes in cluster
 - One DFS chunk per map is common
 - Improves dynamic load balancing and speeds recovery from worker failure
- Usually R is smaller than M, because output is spread across R files

- Often a map task will produce many pairs of the form (k,v1), (k,v2), ... for the same key k
 - E.g., popular words in Word Count
- Can save network time by pre-aggregating at mapper
 - combine(k1, list(v1)) \rightarrow v2
 - Usually same as reduce function
 - Can make a huge difference! Optimization
- Works only if reduce function is commutative and associative
- still grouping and aggregation in reduce() necessary

- Inputs to map tasks are created by contiguous splits of input file
- For reduce, we need to ensure that records with the same intermediate key end up at the same worker
- System uses a default partition function e.g., hash(key) mod R
- Sometimes useful to override
 - E.g., hash(hostname(URL)) mod R ensures URLs from a host end up in the same output file

Implementations of MapReduce



- Google
 - Not available outside Google
- Apache Hadoop
 - Well-known Open-Source implementation of Google's MapReduce & Google File System (GFS)
 - Enriched by many subprojects
 - Cloudera's Distribution: VMWare images, tutorials, patches
 - Used by Yahoo, Facebook, Amazon, IBM, Last.fm, EBay ...
- Hadoop and other MapReduce implementations are available on EC2

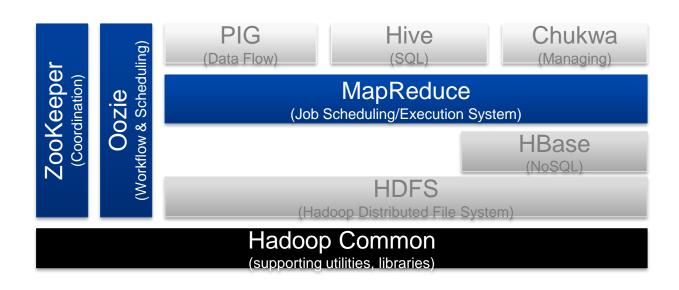
Implementations (2)



- Aster Data
 - Cluster-optimized SQL Database that also implements MapReduce
- MongoDB
- CouchDB
 - having an architecture that "completely embraces the Web"
 - ACID Semantics
 - Built for Offline
 - Distributed Architecture with Replication
 - Document Storage
 - Eventual Consistency
 - Map/Reduce Views and Indexes
 - HTTP API

Hadoop Ecosystem 2011





Hadoop Ecosystem 2015



and many more ...

Further Hadoop Components



- a. Pig
- b. Hive
- c. HBase

13.06.2018 50

Pig (Latin): Why?



- Many parallel algorithms can be expressed by a series of MapReduce jobs
- But MapReduce is fairly low-level:
 - must think about keys, values, partitioning, etc
- Can we capture common "job building blocks"?

Pig (Latin)

- William Charles and the control of t
- Apache Pig: a high-level platform for programs that run on Apache Hadoop
- Pig Latin: Language for this platform
- Started at Yahoo! Research in 2006, since 2007 Apache project
- Runs about 30% of Yahoo!'s jobs*
- Features:
 - Expresses sequences of MapReduce jobs
 - Extract-transform-load (ETL) data pipelines
 - Data model: nested "bags" of items
 - Provides relational (SQL) operators (JOIN, GROUP BY, etc.)
 - Easy to plug in Java functions
 - Pig Pen development environment for Eclipse



* http://blog.cloudera.com/wp-content/uploads/2010/01/IntroToPig.pdf

Example Task

- Suppose you have user data in one file, page view data in another
- and you need to find the top 5 most visited pages by users aged 18 - 25



In MapReduce

```
impost 1ava.util.Itesatos
 impost org.apache.hadoop.fa.Path,
 import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
 import org.apache.hadoop.io.Writable;
import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.mapred.FileInputFormat;
 import org.apache.hadoop.mapred.FileOutputFormat,
import org.apache.hadoop.mapred.JobConf,
import org.apache.hadoop.mapred.KeyValueTextInputFormat,
 import org.a pache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
 import org.apache.hadoop.mapred.RecordReader;
import org. apache. hadoop, mapred. RedoordReader;
import org. apache. hadoop, mapred. Redoordre,
import org. apache. hadoop, mapred. Reporter;
import org. apache. hadoop, mapred. SequenceFileInputFormat;
import org. apache. hadoop, mapred. SequenceFileOutputFormat;
import org. apache. hadoop, mapred. TextImputFormat;
 import org.apache.hadoop.mapred.jobcontrol.Job/
import org.apache.hadoop.mapred.jobcontrol.JobC
import org.apache.hadoop.mapred.lib.IdentityMapper/
public class MRExample {
   public static class LoadPages extends MapReduceBase
                            implements Mapper<LongWeitable, Text, Text, Text> (
                            public void map(LongWritable k, Text val,
                                                        OutputCollector<Text, Text> oc,
Reporter reporter) throws IOException {
                                         // Mail the faportery introd licensplient (
)/ Paul the faportery introduced in the faportery interest of the faportery interest of the faportery interest faportery 
                                             // it came from.
                                         // it came from.
Text outVal - new Text("1 " + value);
oc.collect(outKey, outVal);
              public static class LoadAndFilterUsers extends MapReduceBase
                             implements Mapper LongWritable, Text. Text. Text.
                           public void map(LongWeitable N. Tart val.)
    OutputCollestoreflant, Teath or,
    Reporter resporter) throws IOEnception {
    Fall the kay out
    int in kay out
    int firstComma - line.indexOf(',');
    String value - line.substring() firstComma + 1);
                                        straing value - line.substring( firstComint age - Integer, parseInt(value);
if (age < 18 || age > 25) return;
firing key - line.substring(0, firstComma);
Text outKey - new Text(key);
// Beneard = inter-
                                         // Feepend an index to the value so w

// it came from.

Text outVal - new Text("2" + value);

oc.collect(outKey, outVal);
                                                                                                                                                                                         e know which file
              public static class Join extends MapReduceBase implements Reducer<Text, Text, Text, Text>
                             public woid reduce (Text key,
                                                        Iterator<Text> iter,
                                        OutputCollector(Text, Text> oc, Reporter reporter) throws IOException (
// For each value, figure out which file it's from and
                                          // accordingly.
List<String> first - new ArrayList<String>();
                                          List<String> second - new ArrayList<String>();
                                            while (iter.hasNext()) {
                                                       Text t - iter.next();
String value - t.to Stri:
if (value.chasAt(0) -- 'l')
 first.add(value.substring(1))
                                                         else second.add(value.substring(1));
```

```
// Do the cross product and collect the values
                     for (String al : first) {
  for (String a2 : second) {
    String outval - key + "," + al + ","
    oc.collect(null, new Text(outval));
}
                                   reporter.setStatus("OK");
       public static class LoadJoined extends MapReduceBase
              implements Mapper<Text, Text, Text, LongWritable> {
                            Text k,
Text val,
OutputColle
                                                  ctor<Text, LongWritable> oc,
                             Reporter reporter) throws ICException (
                    Reporter seporter) throws IOExcaption {
// Find the wel
// Find the wel
String line - val.toString();
int firstComma - line.indexOf(',');
int secondComma - line.indexOf(',', first Comma);
String Kay - line.subtring(firstComma, secondComma);
                     // deop the rest of the record, I don't need it anymore,
// just pass a 1 for the combiner/reducer to sum instead.
Text outKey - new Text (key);
                     oc.collect(outKey, new LongWeitable(lL));
public static class ReduceUrls extends MapReduceBase implements Reducer<Text, LongWritzble, WritzbleComparable, Writzbler (
                            Test ke y,
Iterator<LongWritable> iter,
                             Trefator-Longwertzblev leer,
OutputCollector-WeitzbleComparable, Weitzble> oc,
Reporter reporter) throws IOException {
                     // Add up all the values we see
                     wh ile (iter.hasNext()) {
   sum +- iter.hext().get(),
   reporter.setStatus("OK"),
                     oc.collect(key, new LongWritable(sum));
       public static class LoadClicks extends MapReduceBase
              i mplementa Mapper-WeitableComparable, Weitable, LongWeitable,
Text>
            public void map (

**Textshacomparable key, 
**Westable val, 
**CutputCollector=CongWestable, Text> oc, 
**Reporter seporter) throws ToException (
oc.ollect(CongWestable)val, (Text)key) (
       public static class LimitClicks extends MapReduceBase
              implements Reducer<LongWritable, Text, LongWritable, Text>
              int count - 0;
public void meduce(
                     LongWritable key,
                     Iterator<Text> iter.
                     OutputCollector<LongWritable, Text> oc,
Reporter reporter) throws IOException (
                     // Only output the first 100 seconds
while (count < 100 ss iter.hasNext()) {
  oc.collect(key, iter.hasNext()) </pre>
                            count++,
      public static void main(String[] args) throws IOException {
    JobConf lp = new JobConf (MREmample.class);
    lp.se tJobname("Load Pages")
              ip.setInputFormat(TextInputFormat.class);
```

```
ip.setMappe=Class(LoadPages.class);
               FileInputFormat.addInputPath(lp, new
user/qates/pages"));
                FileOutputFormat.setOutputPath(lp,
                new Path("/user/gates/tmp/indexed_pages"));
lp.setNumReduceTasks(0);
                Job loadPages - new Job(lp);
               JobConf ifu = new JobConf (MREwample.class);

ifu.s etJobName("Load and Filter Users");

ifu.setInputFormat(TextInputFormat.class);

ifu.setOutputKayClass(Text.class);
                lfu.setOutputValueClass(Text.class);
                Ifu.setMapperClass(LandFilterUsers.class);
FileInputFormat.add InputPath(Lfu, new
 Path("/user/cates/users"));
                FileOutputFormat.setOutputFath(lfu,
new Path("/user/gates/tmp/filtered_users"));
                lfu.setNumReduceTasks(0);
                Job loadUsers - new Job (lfu);
                 JobConf join - new JobConf(
                                                                           MREwample.class);
                 join.setJobName("Join Users and Pages"),
join.setInputFormat(KeyValueTextInputFormat.class),
                join.setOutputKeyClass(Text.class);
join.setOutputValueClass(Text.class);
join.setMappe=Class(IdentityMap p
                 ioin setReduce=Class(Join class).
join.setkeduderCiass(Join.class))
FileInputFormat.addInputFath(join, new
Path("/user/gates/tmp/indexed_pages"));
Path ("/usas/gatea/tmp/induxed_pagea"));

Sth ("usasputPemat.add.mputPath(join, new

Sth ("usasputPemat.as tOutputPath (join, new

Path ("join.astTumReducaTakka (50);

Joh join.astTumReducaTakka (50);

Job join.ob - new Job (join);
                 inin.Tob.addDemending.Tob(loadPages).
                 joinJob.addDependingJob (loadUsers),
                JobConf group - new JobConf (MRE
                group.setJobName("Group URLs");
group.setInputFormat(KeyValueTextInputFormat.class);
                group.setOutputKeyClass(Text.class);
group.setOutputValueClass(clagWritable.class);
group.setOutputFormat(SequenceF: leOutputFormat.class);
group.astOutputFormat(SequenceFi laOut)
group.astMappentlass(loadfoined.class))
group.astReducerClass(ReduceVis.class))
group.astReducerClass(ReduceVis.class)
Fath("user/gates/tmp/joined"))
FileOutputFormat.astOutputFath(group.new
Path("user/gates/tmp/joined"))
 Path ("/was=/gates/tmp/g=owned"))
                group.setNumReduceTasks(50);
Job groupJob - new Job(group)
                g=oupJob.addDependingJob(joinJob);
                top100.setJobname("Top 100 sites"),
top100.setJobname("Top 100 sites"),
top100.setJoputFormat(SequenceFileInputFormat.class),
top100.setOutputKeyClass(LongWeitable.class),
top100.setOutputValusClass(Text.class),
                top100.setOutputFormat(SequenceFileOutputF
top100.setMapperClass(LoadClicks.class);
                top100.setCombine=Class(LimitClicks.class);
top100.setReduce=Class(LimitClicks.class),
FileInputFormat.addInputFath(top100, new
Path("/use=/gates/tmp/grouped")),
FileOutputFormat.satOutputPath(top100, new Path("/user/gates/top100sitesforusers18to25"));
                top100.setNumReduceTasks(1)
Job limit - new Job(top100)
                limit.addDependingJob(groupJob);
                JobControl jo - new JobControl("Find top
                                                                                                         100 sites for users
                o"),
jc.addJob(loadPages),
jc.addJob(loadUsers),
jc.addJob(jcinJob),
                jc.addJob(groupJob);
jc.addJob(limit);
                je.sun();
```

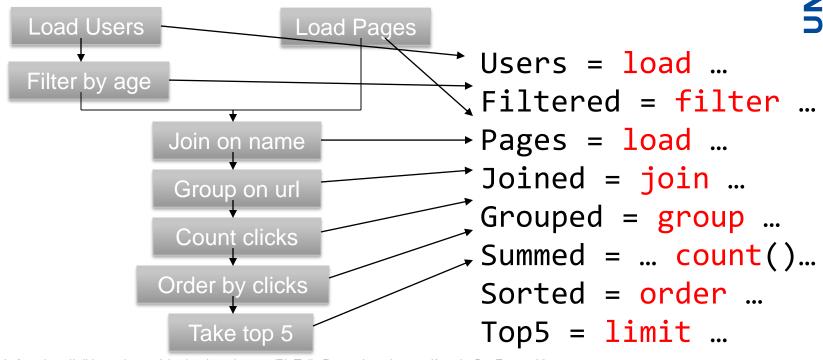
JNI REIBURG

In Pig Latin

```
Users = load 'users' as (name, age);
Filtered = filter Users by
                 age >= 18 and age <= 25;
Pages = load 'pages' as (user, url);
Joined = join Filtered by name, Pages by user;
Grouped = group Joined by url;
Summed = foreach Grouped generate group,
                  count(Joined) as clicks;
Sorted = order Summed by clicks desc;
Top5 = limit Sorted 5;
store Top5 into 'top5sites';
```

Ease of Translation





Hive



- Developed at Facebook
- Used for majority of Facebook jobs
- Written in Java, current version 2.1
- Data Warehouse infrastructure that provides data summarization and ad hoc querying on top of Hadoop
 - MapReduce for execution
 - HDFS for storage
- "Relational database" built on Hadoop
 - Maintains list of table schemas
 - SQL-like query language (HQL)
 - Supports table partitioning, clustering, complex data types, some optimizations



Sample Hive Query

Find top 5 pages visited by users aged 18-25:

```
SELECT p.url, COUNT(1) as clicks
FROM users u JOIN page_views p ON (u.name = p.user)
WHERE u.age >= 18 AND u.age <= 25
GROUP BY p.url
ORDER BY clicks
LIMIT 5;</pre>
```

HBase



Clone of BigTable (Google)

- See "Bigtable: A Distributed Storage System for Structured Data", 2006.
- Data is stored "column-oriented"
- Distributed over many servers
- Layered over HDFS
- Strong consistency (CAP Theorem)
- Scalable up to billions of rows and millions of columns





Sqoop

- Tool designed to help users of large data import existing relational databases into their Hadoop cluster
- Integrates with Hive

Zookeeper

High-performance coordination service for distributed applications

Avro

- Data serialization system

Chukwa

- Data collection system
- Displaying, monitoring and analyzing log files

Takeaways

- By providing a data-parallel programming model, MapReduce can control job execution in useful ways:
 - Automatic division of job into tasks
 - Automatic partition and distribution of data
 - Automatic placement of computation near data
 - Recovery from failures & stragglers
- Hadoop, an open source implementation of MapReduce, enriched by many useful subprojects
- User focuses on application, not on complexity of distributed computing

Resources



- Cloudera's Hadoop Distribution → recommended
 - http://www.cloudera.com/
 - Virtual Machine Image (VMWare)
- Video tutorials
 - http://www.cloudera.com/training/library/hadoop-essentials.html
 - http://www.cloudera.com/training/library/tutorials.html
 - http://www.cloudera.com/training/library/hadoop-ecosystem.html
- Apache Hadoop
 - http://hadoop.apache.org/core/
 - http://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html
- Project website:
 - MapReduce introduction slides
 - Eclipse projects with some MapReduce examples

Resources (2)

THE STATE OF THE S

- MapReduce: Simplified data processing on large clusters
 - Dean, J. and Ghemawat, S., ACM 2008
- The Google file system
 - Ghemawat, S. and Gobioff, H. and Leung, S.T., 2003
- Hadoop: The Definitive Guide, 4th Edition → recommended
 - Tom White
 - O'Reilly Media, 2015
- Data-Intensive Text Processing with MapReduce
 - Jimmy Lin, Chris Dyer, Graeme Hirst
 - Morgan and Claypool Publishers, 2010
- Cluster Computing and MapReduce Lecture Series
 - Google, 2007: Available on Youtube

Acknowledgement

A CONTROLLING THE PARTY OF THE

- Slides are partially from
 - Prof. Jeffrey D. Ullman & Dr. Jure Leskovec.
 - Martin Przyjaciel-Zablocki & Alexander Schätzle.