

Distributed Processing with Computer Clusters II

Advanced Databases and Information Systems
SS18

Albert-Ludwigs-Universität Freiburg

Prof. Dr. Georg Lausen
Patrick Philipp

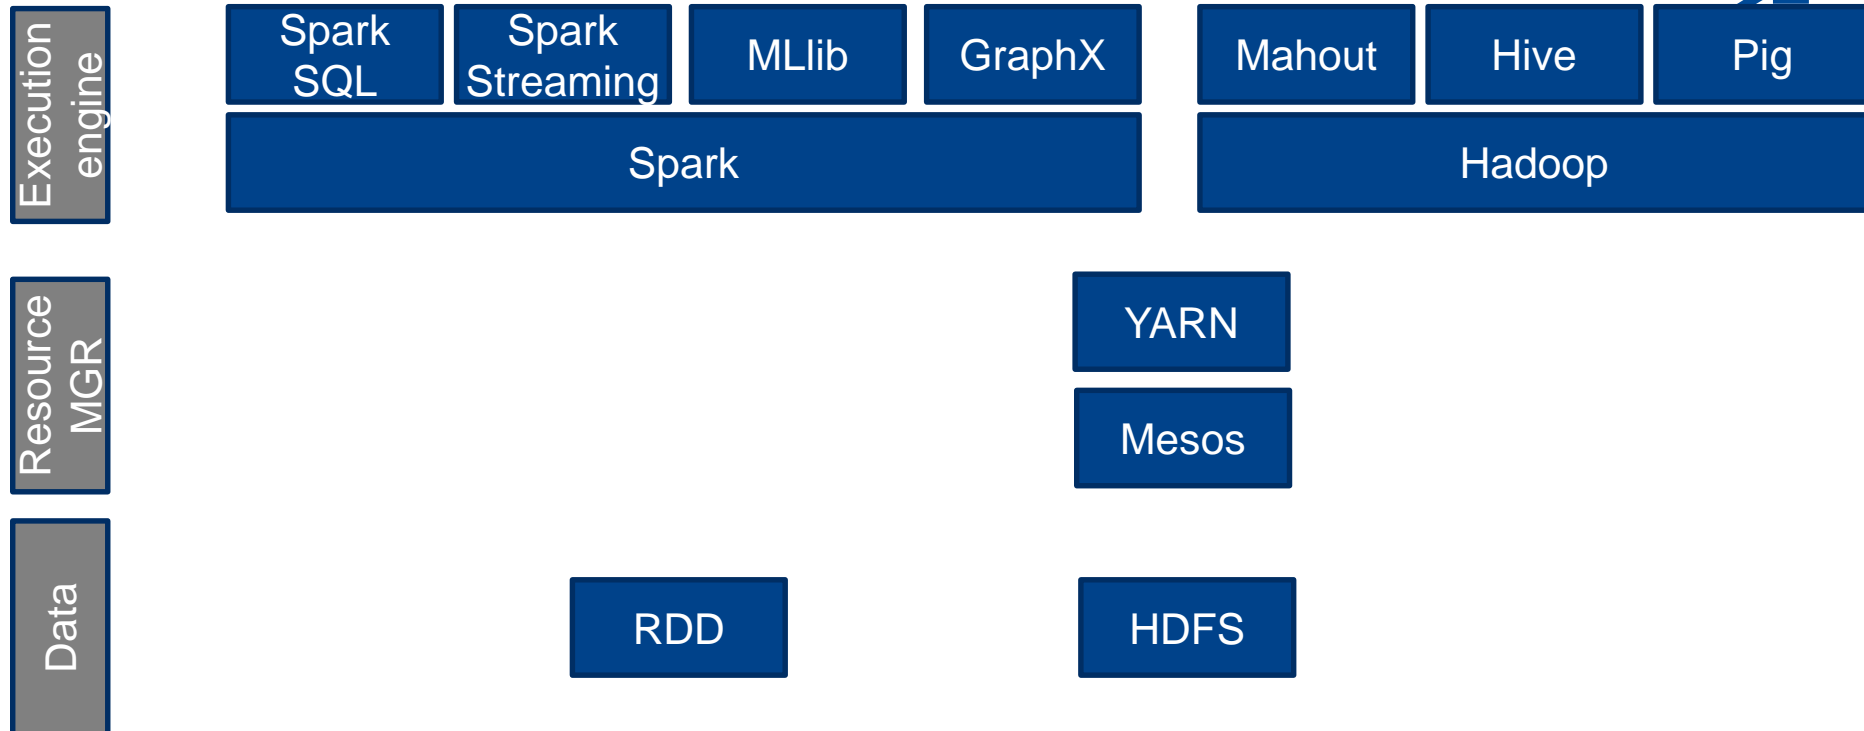


UNI
FREIBURG

- **Open-source, distributed, general-purpose cluster computing framework.**
- For processing large volumes of data, both *batch processing* and *streaming*.
- Supported features, among others:
 - analytics
 - machine learning
 - graph processing
- APIs for Scala, Python, Java, R, and SQL.
- Specialized modules that run on Spark core:
 - Spark SQL, Spark Streaming, MLlib (machine learning), GraphX (Graph processing)



Spark Components



Operations in Spark



- More general *functional* programming model than MapReduce:
Transformations and **actions**
- Examples of transformations: *map*, *filter*, *groupBy*
- Examples of actions: *count*, *collect*, *save*
- (In MapReduce, transformation \approx map(), action \approx reduce())

Data in Spark: RDDs



- **Resilient Distributed Datasets (RDDs)**
- RDD=Collection of elements spread across a cluster, which is
 - **immutable** (read-only) **Nothing will be deleted in future**
 - **resilient** (fault-tolerant: Automatically rebuilt in case of failure)
 - **distributed** (dataset **spread out to more than one node**)
- Stored in RAM or on disk.
- (Parallel) **Transformation**: Build a new RDD
- **Action**: Performed on RDD's elements or return result to program.

Cluster Architecture



- **Job:** bunch of transformations & actions on RDDs
- **Cluster manager:** Allocates worker nodes

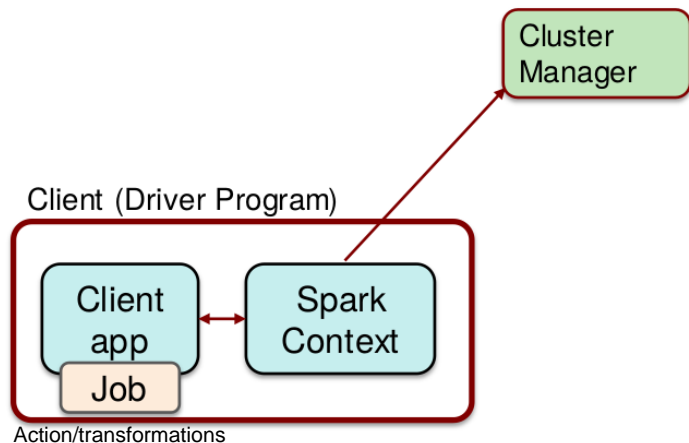


figure from Paul Krzyzanowski

Cluster Architecture (2)



- **Driver** breaks the job into tasks
- Sends tasks to worker nodes where the data lives

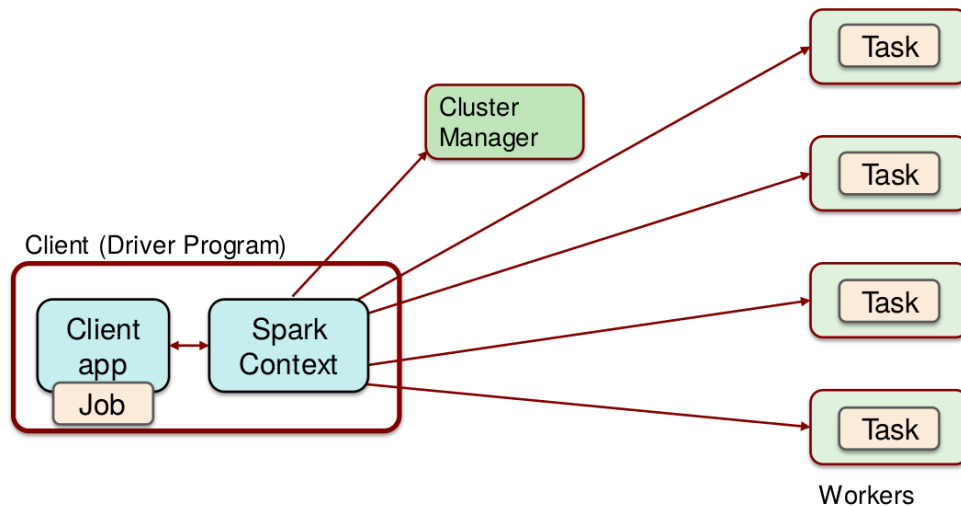
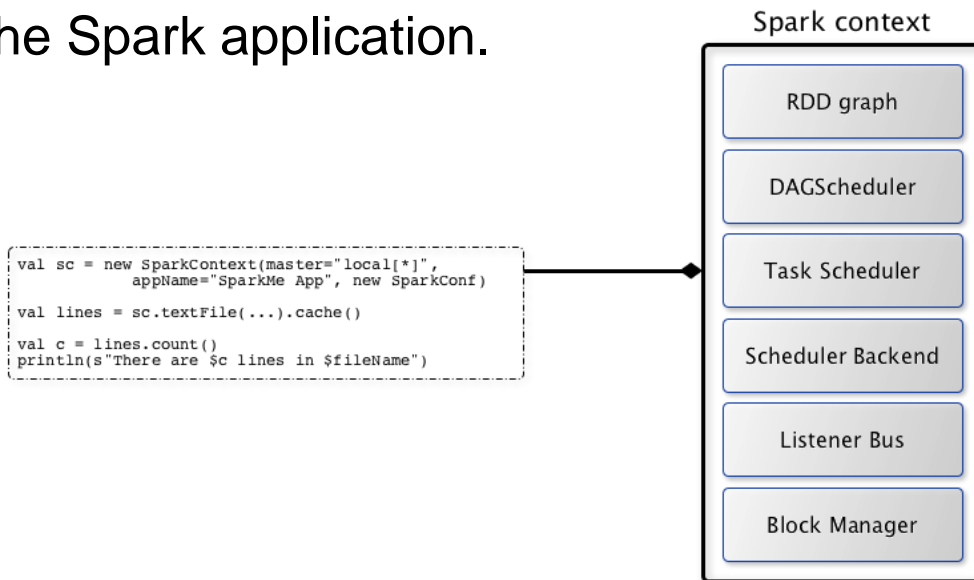


figure from Paul Krzyzanowski

Cluster Architecture (3)



- **Spark context:** sets up internal services and establishes a connection to a Spark execution environment.
- masters the Spark application.




from: <https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/spark-sparkcontext.html>

Worker Node



- Has one or more *executors*
- JVM process
- Talks with cluster manager
- Receives tasks
 - JVM code (e.g., compiled Java, Clojure, Scala, ...)
 - Task = *transformation* or *action*
- *RDD*
 - *Data to be processed*
 - *local to the node*
- *Cache*
 - *frequently used data is kept in memory* - for high performance!



- Organized in RDDs
- Idea: Partition (big) data across machines
- How are RDDs created:
 1. Create from any **file** stored in HDFS or other supported storage (Amazon S3, HBase, Cassandra, etc.)
 - Created externally (e.g., event stream, text files, database)
 - Examples:
 - Query a database & use results as RDD
 - Any Hadoop InputFormat, such as a list of files or a directory

Data Organization (2)



2. *Streaming sources* (via Spark Streaming)
 - Fault-tolerant stream with a sliding window
3. An RDD can be the *output of a Spark transformation function*.
 - Example: filter out data, select key-value pairs

Data Organization (3)



- Main properties of RDDs:
 - Immutable
 - You cannot change RDDs, only create new RDDs
 - The framework will eventually collect unused RDDs
 - Partitioned: RDDs distributed among servers
 - Default partitioning function: $\text{hash}(\text{key}) \bmod \text{server_count}$
- Optional properties of RDDs:
 - Typed: RDDs are not BLOBs
 - – Embedded data structure, e.g., key-value set
 - Ordered: Elements in an RDD can be sorted

Operations on RDDs



Two types of operations on RDDs:

1. *Transformations*

- **Lazy**: not computed immediately, only if result required by driver program.
- Transformed RDD is recomputed when an action is run on it
- RDD can be persisted into memory or disk storage


2. *Actions*

- *Finalizing* operations
 - e.g., reduce, count, grab samples, write to file

Operations on RDDs (2)



- Transformations and Actions are generalizations of `map()` and `reduce()` of Hadoop
- In Spark: Directed Acyclic Graph (DAG) created and submitted to DAG scheduler
 - Nodes are RDDs
 - Arrows are transformations
- => more efficient than MapReduce, since **DAG optimizer** rearranges the order of the operators and since data is cached.



Spark Transformations



Transformation	Description
map (func)	Pass each element through a function func
filter (func)	Select elements of the source on which func returns true
flatMap (func)	Each input item can be mapped to 0 or more output items
sample (withReplace ment, fraction, seed)	Sample a <i>fraction</i> of the data, with or without replacement, using a given random number generator seed
union (otherdataset)	Union of the elements in the source data set and otherdataset
distinct ([numtasks])	The distinct elements of the source dataset

Spark Transformations (2)



Transformation	Description
groupByKey ([numtasks])	When called on a dataset of (K,V) pairs, returns a dataset of (K,seq[V]) pairs
reduceByKey (func, [numtasks])	Aggregate the values for each key using the given <i>reduce</i> function
sortByKey ([ascending], [numtasks])	Sort keys in ascending or descending order
join (otherDataset, [numTasks])	Combines two datasets (K,V) and (K,W) into (K, (V,W))
cogroup (otherDataset, [numtasks])	Given (K,V) and (K,W), returns (K,Seq[V], Seq[W])
cartesian (otherDataset)	For two datasets T and U, returns a dataset of (T,U) pairs

Spark Actions



Action	Description
reduce (func)	Aggregate elements of the dataset using <i>func</i> .
collect (func, [numtasks])	Return all elements of the dataset as an array
count ()	Return the number of elements in the dataset
first ()	Return the first element of the dataset
take (n)	Return an array with the first n elements of the dataset
takeSample (withReplacement, fraction, seed)	Return an array with a random sample of <i>num</i> elements of the dataset.

Spark Actions (2)



Action	Description
saveAsTextFile (path)	Write dataset elements as a text file
saveAsSequenceFile (path)	Write dataset elements as a Hadoop SequenceFile
countByKey ()	For (K,V) RDDs, return a map of (K, Int) pairs with the count of each key
foreach (func)	Run <i>func</i> on each element of the dataset

Descriptions from Paul Krzyzanowski

Spark in Action – Some Examples



- Obtain number of BSD licenses written in license file “LICENSE”

```
val licLines = sc.textFile("/usr/local/spark/LICENSE")  
val bsdLines = licLines.filter(line => line.contains("BSD"))  
bsdLines.count  
bsdLines.foreach(bLine => println(bLine))
```

Download spark in action

Spark in Action – map



```
val numbers = sc.parallelize(10 to 50 by 10)
numbers.foreach(x => println(x))
val numbersSquared = numbers.map(num => num * num)
numbersSquared.foreach(x => println(x))
```

Spark in Action – distinct & flatMap



```
val lines = sc.textFile("/home/spark/client-ids.log")
val idsStr = lines.map(line => line.split(","))
ids.collect
idsStr.first
```

```
val ids = lines.flatMap(_.split(","))
ids.collect
ids.first
```

```
val intIds = ids.map(_.toInt)
intIds.collect
```

```
val uniqueIds = intIds.distinct
uniqueIds.collect
val finalCount = uniqueIds.count
```

Spark in Action – sample



Prepare sample of 30% of clientIDs:

without replacement:

```
val s = uniqueIds.sample(false, 0.3)
s.collect
s.count
```

with replacement:

```
val swr = uniqueIds.sample(true, 0.5)
swr.collect
swr.count
```

Spark in Action – take, takeSample



- **takeSample:**

Action!

```
val taken = uniqueIds.takeSample(false, 5)
```

- **take:**

```
uniqueIds.take(3)
```

Spark in Action – Double values



- Implicit conversion of int to double in Scala
- Can be used for total sum of all of an RDD 's elements and their mean value, standard deviation, and variance.
- Example code:

```
intIds.mean  
intIds.sum  
intIds.variance  
intIds.stdev
```


Note



- RDDs have new methods added to them automatically, depending on the type of data they hold.

- Spark does not care how source data is stored
 - RDD connector determines that
- RDD fault tolerance
 - RDDs track the sequence of transformations used to create them
 - Enables recomputing of lost data
 - Go back to the previous RDD and apply the transforms again

Example: Log processing



- Transform (create new RDDs):
 1. Retrieve error message from a log file.
 2. Retrieve only ERROR messages and extract the source of error.
- Actions to perform: Count mysql errors and php errors

```
val lines = sc.textFile("hdfs://...")  
val errors = lines.filter(_.startsWith("ERROR"))  
val messages = errors.map(_.split("\t")).map(r => r(1))  
messages.cache()  
messages.filter(_.contains("mysql")).count()  
messages.filter(_.contains("php")).count()
```

base RDD

transformed
RDDs


action 1

action 2

Spark Streaming




- MapReduce, Pregel, etc. work on static data
- Spark Streaming enables processing of live data streams
 - Same programming operations
 - Input data is chunked into batches
 - Time interval specified in implementation



Spark Streaming (2)

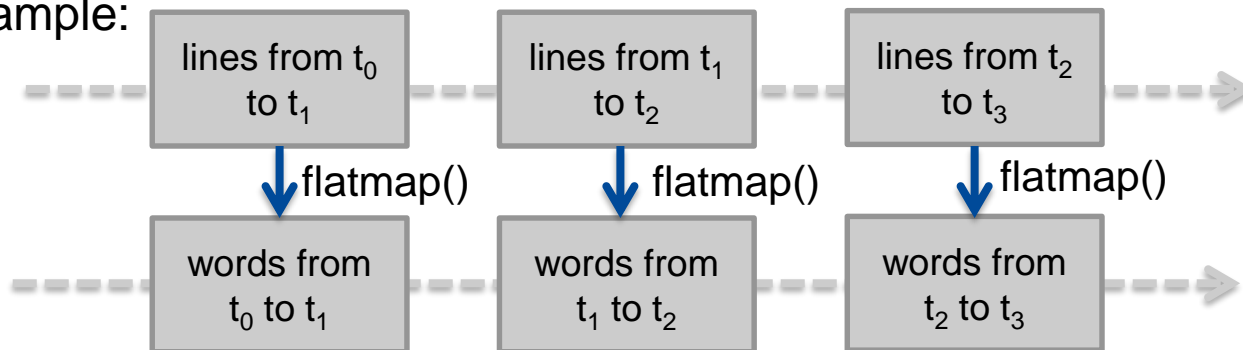


- Discretized Stream = DStream
 - Continuous stream of data (from source or a transformation)
 - Appears as a series of RDDs, each for a time interval



- Each operation on a DStream translates to operations on the RDDs

Example:






- Join operations allow combining multiple streams

Spark Core Summary



- Resilient Distributed Datasets (**RDDs**) as data collections.
 - RDDs are created from files (HDF etc.; storage agnostic), streaming sources, or output of Spark transformation function.
 - RDDs are immutable and distributed among servers.
 - Fault tolerant: RDDs can be regenerated.
- Spark operations: **Transformations** and **actions** on RDDs.
 - ensemble of different transformations and actions form a directed acyclic graph.
- Fast
 - Often up to 10x faster on disk and 100x faster in memory than MapReduce
 - General execution graph model
 - In-memory storage for RDDs
- Spark streaming: Handle continuous data streams via Spark Streaming



MLlib: Machine Learning in Spark



- Example: **Linear SVM learning**

from: <https://spark.apache.org/docs/1.2.0/mllib-linear-methods.html#linear-support-vector-machines-svms>

```
import ...

// Load training data in LIBSVM format.
val data = MLUtils.loadLibSVMFile(sc, "data/mllib/sample_libsvm_data.txt")

// Split data into training (60%) and test (40%).
val splits = data.randomSplit(Array(0.6, 0.4))
val training = splits(0).cache()
val test = splits(1)

// Run training algorithm to build the model
val numIterations = 100
val model = SVMWithSGD.train(training, numIterations)
```


MLlib: Machine Learning in Spark (2)



```
// Clear the default threshold.
model.clearThreshold()

// Compute raw scores on the test set.
val scoreAndLabels = test.map { point =>
    val score = model.predict(point.features)
    (score, point.label)
}

// Get evaluation metrics.
val metrics = new BinaryClassificationMetrics(scoreAndLabels)
val auROC = metrics.areaUnderROC()

// Print result (receiver operating characteristic curve)
println("Area under ROC = " + auROC)
```

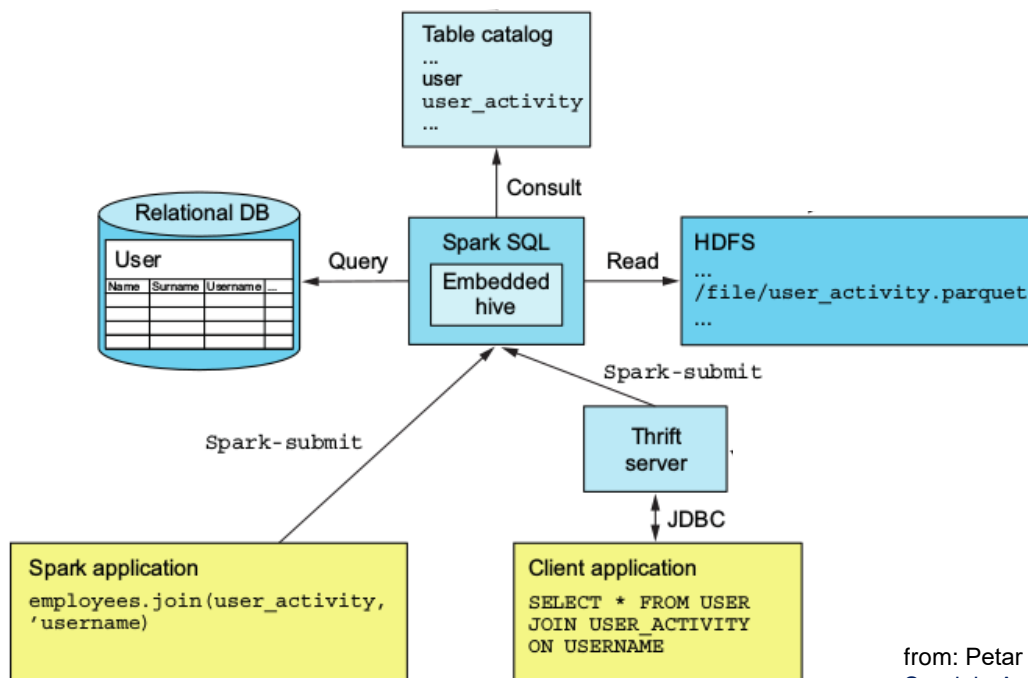
DataFrame in Spark



- Central component, supported since Spark 1.3
- Handle structured data in a table-like representation (with column names/types)
- Translate SQL code and domain-specific language (DSL) expressions into optimized low-level **RDD operations**
- Yield one common API for all languages (Scala, Java, Python, R)

- Since Spark 2.0: DataFrame is a row of a DataSet
- Registration of DataFrames in a table catalog to become visible/queryable
 - Only DataFrame name needed for querying.
- Spark exploits HiveQL, which is similar to SQL (see Hive in last lecture; not only used for MapReduce jobs, but also for Spark jobs)
 - It is also possible to query from remote via JDBC/ODBC

Spark SQL Overview



from: Petar Zečević and Marko Bonaći:
Spark in Action, Manning: 2016.

How to create DataFrames?



1. Converting existing RDDs
2. Running SQL queries
3. Loading external data

From RDDs to DataFrames



- Possibilities for creating a DataFrame from RDD:

- a) Using RDDs containing row data as tuples
- b) Using case classes
- c) Specifying a schema

Schema is only inferred for
Option a) and b).

Option c) mostly used; there,
explicitly specifying a
schema.

From RDDs to DataFrames



- Needed for transformation of RDDs to DataFrames:
 - **SparkSession**
 - **SparkSession** is a wrapper around **SparkContext** and **SQLContext** (which was directly used for creating DataFrames until Spark 2.0)
 - Implicit methods for transforming RDDs to DataFrames:
 - `import spark.implicits._`
 - allows to call additional method `toDF`.

a) Using RDDs containing row data as tuples



- Example in Scala:

```
scala> val itPostsRDD = itPostsSplit.map(x => (x(0),x(1),x(2),x(3),x(4),
x(5),x(6),x(7),x(8),x(9),x(10),x(11),x(12)))
itPostsRDD: org.apache.spark.rdd.RDD[(String, String, ...
scala> val itPostsDataFrame = itPostsRDD.toDF()
itPostsDF: org.apache.spark.sql.DataFrame = [_1: string, ...
```

- Then one can work with it:

```
scala> itPostsDataFrame.show(10)
```

```
+---+-----+-----+-----+-----+-----+
| _1|          _2| _3|          _4| _5|          _6|
+---+-----+-----+-----+-----+-----+
| 4|2013-11-11 18:21:...| 17|&lt;p>The infi...| 23|2013-11-10 19:37:...
| 5|2013-11-10 20:31:...| 12|&lt;p>Come cre...| 1|2013-11-10 19:44:...
| 2|2013-11-10 20:31:...| 17|&lt;p>Il verbo...| 5|2013-11-10 19:58:...
...
```

not very elegant

a) Using RDDs containing row data as tuples (2)



- Recommended: column names given automatically, rename column names with toDF function:

```
scala> val itPostsDF = itPostsRDD.toDF("commentCount", "lastActivityDate",  
"ownerUserId", "body", "score", "creationDate", "viewCount", "title",  
"tags", "answerCount", "acceptedAnswerId", "postTypeId", "id")
```

- Check (print) column names:

```
scala> itPostsDF.printSchema  
root  
|-- commentCount: string (nullable = true)  
|-- lastActivityDate: string (nullable = true)  
|-- ownerUserId: string (nullable = true)  
|-- body: string (nullable = true)  
|-- score: string (nullable = true)
```


b) RDDs to DataFrames using Case Classes



- In columns all strings so far => Introduce data types via case classes.
- Steps: (a) define case class (b) map each row in RDD to case class (c) use toDF method.
- Example for (a):

```
import java.sql.Timestamp
case class Post(
  commentCount:Option[Int],
  lastActivityDate:Option[java.sql.Timestamp],
  ownerUserId:Option[Long],
  body:String,
  score:Option[Int],
  creationDate:Option[java.sql.Timestamp],
  ...
  postTypeId:Option[Long],
  id:Long)
```

b) RDDs to DataFrames using Case Classes (2)



- For better conversion (of different data types, i.e., to not only have strings), define class and methods as follows:

```
object StringImplicits {  
  implicit class StringImprovements(val s: String) {  
    import scala.util.control.Exception.catching  
    def toIntSafe = catching(classOf[NumberFormatException]) opt s.toInt  
    def toLongSafe = catching(classOf[NumberFormatException]) opt s.toLong  
    def toTimestampSafe = catching(classOf[IllegalArgumentException]) opt  
      Timestamp.valueOf(s)  
  }  
}
```

b) RDDs to DataFrames using Case Classes (3)



- Actual transformation:

```
import StringImplicits._
def stringToPost(row:String):Post = {
  val r = row.split("~")
  Post(r(0).toIntSafe,
    r(1).toTimestampSafe,
    r(2).toLongSafe,
    r(3),
    r(4).toIntSafe,
    r(5).toTimestampSafe,
    r(6).toIntSafe,
    r(7),
    ...
    r(11).toLongSafe,
    r(12).toLong)
}
val itPostsDFCase = itPostsRows.map(x => stringToPost(x)).toDF()
```

b) RDDs to DataFrames using Case Classes (4)



- Check if transformation successful:

```
scala> itPostsDFCase.printSchema
root
|-- commentCount: integer (nullable = true)
|-- lastActivityDate: timestamp (nullable = true)
|-- ownerUserId: long (nullable = true)
|-- body: string (nullable = true)
|-- score: integer (nullable = true)
|-- creationDate: timestamp (nullable = true)
|-- viewCount: integer (nullable = true)
|-- title: string (nullable = true)
|-- tags: string (nullable = true)
|-- answerCount: integer (nullable = true)
|-- acceptedAnswerId: long (nullable = true)
|-- postTypeId: long (nullable = true)
|-- id: long (nullable = false)
```

c) From RDDs to DataFrames by Specifying a Schema



- Use `SparkSession`'s `createDataFrame` method (needed: objects of type `Row` and a `StructuredType` (=schema)).

- `StructuredType` definition example:

```
import org.apache.spark.sql.types._
val postSchema = StructType(Seq(
  StructField("commentCount", IntegerType, true),
  StructField("lastActivityDate", TimestampType, true),
  StructField("ownerUserId", LongType, true),
  StructField("body", StringType, true),
  StructField("score", IntegerType, true),
  StructField("creationDate", TimestampType, true),
  StructField("viewCount", IntegerType, true),
  ...
  StructField("id", LongType, false))
)
```

Supported data types:
strings, integers, shorts,
floats, doubles, bytes,
dates, timestamps, binary
values, arrays, maps,
structs.

c) From RDDs to DataFrames by Specifying a Schema (2)



- Actual transformation (now **Row**, not **Post** type; the rest is the same as above):

```
def stringToRow(row:String):Row = {  
  val r = row.split("~")  
  Row(r(0).toIntSafe.getOrElse(null),  
    r(1).toTimestampSafe.getOrElse(null),  
    r(2).toLongSafe.getOrElse(null),  
    r(3),  
    r(4).toIntSafe.getOrElse(null),  
    r(5).toTimestampSafe.getOrElse(null),  
    r(6).toIntSafe.getOrElse(null),  
    r(7),  
    r(8),  
    r(9).toIntSafe.getOrElse(null),  
    r(10).toLongSafe.getOrElse(null),  
    r(11).toLongSafe.getOrElse(null),  
    r(12).toLong)  
}
```

c) From RDDs to DataFrames by Specifying a Schema (3)



```
val rowRDD = itPostsRows.map(row => stringToRow(row))  
val itPostsDFStruct = spark.createDataFrame(rowRDD, postSchema)
```

- Check if everything is fine, e.g., by `.columns`-method (=print column names) or by `.dtypes`-method (column names plus data types):

```
scala> itPostsDFCase.columns  
res0: Array[String] = Array(commentCount, lastActivityDate, ownerUserId,  
body, score, creationDate, viewCount, title, tags, answerCount,  
acceptedAnswerId, postTypeId, id)
```

How to Use the DataFrame API



- Example:

```
scala> val postsDf = itPostsDFStruct
scala> val postsIdBody = postsDf.select("id", "body")
postsIdBody: org.apache.spark.sql.DataFrame = [id: bigint, body: string]
```

- Using col-method:

```
val postsIdBody = postsDf.select(postsDf.col("id"), postsDf.col("body"))
```

- To retrieve all columns except one: e.g.

```
val postIds = postsIdBody.drop("body")
```

- removes the **body** column from the **postsIdBody** DataFrame

How to Use the DataFrame API (2)



- Filtering data:
 - using **where** or **filter** functions (synonymous!)

- Examples:

```
scala> postsIdBody.filter('body contains "Italiano").count  
res0: Long = 46
```

Hyphon needed for scala parser

```
scala> val noAnswer = postsDf.filter(('postTypeId === 1) and  
('acceptedAnswerId isNull))
```

- Return top n elements, using limit:

```
scala> val firstTenQs = postsDf.filter('postTypeId === 1).limit(10)
```

SQL Functions in Spark



- SQL functions are available through
 1. the *DataFrame API* and
 2. *SQL expressions*.

SQL Functions Using the DataFrame API



- 4 kinds of SQL functions when using the DataFrame API:
 - a) **Scalar functions** return a single value for each row based on calculations on one or more columns.
 - b) **Aggregate functions** return a single value for a group of rows.
 - c) **Window functions** return several values for a group of rows.
 - d) **User-defined functions** include custom scalar or aggregate functions.

a) Scalar Functions



- **Math calculations:**
 - **abs** (calculates absolute value),
 - **hypot** (calculates hypotenuse based on two columns or scalar values),
 - **log** (calculates logarithm),
 - **cbrt** (computes cube root), and others
- **String operations:**
 - **length** (calculates length of a string),
 - **trim** (trims a string value left and right),
 - **concat** (concatenates several input strings), and others
- **Date-time operations:**
 - **year** (returns the year of a date column),
 - **date_add** (adds a number of days to a date column), and others.

c) Window Functions



- Window functions: for making selects and joins simpler.
- First import by

```
import org.apache.spark.sql.expressions.Window
```

Using SQL Commands in Spark



- Idea: Use SQL, since SQL is widely used and easier
- SQL commands get translated into DataFrames
- Spark supports (1) Spark's SQL dialect and (2) Hive Query Language (HQL).
- HQL is recommended, since richer functionalities and more support from community.
- One can store DataFrames as tables in a *table catalog* – either temporarily or permanently.

- **Temporary registration:**

```
postsDf.createOrReplaceTempView("posts_temp")
```

Afterwards, you are able to query the data with SQL commands

- **Permanent registration:**

```
postsDf.write.saveAsTable("posts")
```

```
votesDf.write.saveAsTable("votes")
```

SparkSession with Hive support can be used to register table definitions that will survive your application's restarts

Using SQL Commands in Spark (2)



- **Checking the table catalog:**

```
scala> spark.catalog.listTables().show()
```

- **Execute SQL queries then:**

```
val resultDf = sql("select * from posts")
```

- **Result is a DataFrame again.**

- **Spark SQL shell** (command: `spark-sql`), which provides additional functions beyond `spark-shell` and `spark-submit`.

- **Example:**

```
spark-sql> select substring(title, 0, 70) from posts where  
postTypeId = 1 order by creationDate desc limit 3;
```

JDBC/ODBC Connection



- Note: Besides executing SQL queries directly from programs or through SQL shell: SQL commands via JDBC (or ODBC) possible, namely via Thrift (which is a special Spark application).

Acknowledgements.



- Slides are partially based on
 - Petar Zečević and Marko Bonaći: [Spark in Action](#), Manning: 2016.
 - Paul Krzyzanowski, Rutgers University, 2016.