

Blockchain and Cryptocurrencies

Week 5 — Chapter 3 (part 2): Mechanics of Bitcoin

Prof. Dr. Peter Thiemann

Albert-Ludwigs-Universität Freiburg, Germany

SS 2020

Mechanics of Bitcoin

draws on material from

- Bitcoin and Cryptocurrency Technologies
- Bitcoin, Blockchain, and Cryptoassets
- Antonopoulos: Mastering Bitcoin

Contents

1 Bitcoin Blocks

2 Bitcoin Network


- Network Protocol
- SPV and Bloom Filters

Why blocks?

- transactions are grouped in blocks of 1 MB
- mainly for efficiency:
 - ▶ it keeps the chain shorter
 - ▶ it helps to maintain a reasonable transaction rate

(to r/w P2P Consensus)

Structure of a Block



hash chain of
block headers

Merkle tree of
transactions

Block Header (binary format)

Size	Field	Description
4 bytes	Version	A version number to track software/protocol upgrades
32 bytes	Previous Block Hash	A reference to the hash of the previous (parent) block in the chain
32 bytes	Merkle Root	A hash of the root of the Merkle-Tree of this block's transactions
4 bytes	Timestamp	The approximate creation time of this block (seconds from Unix Epoch)
4 bytes	Difficulty Target	The Proof-of-Work algorithm difficulty target for this block
4 bytes	Nonce	A counter used for the Proof-of-Work algorithm

- only the block header is hashed!

Coinbase Transactions

- ① minting new bitcoins
- ② single input, single output
- ③ input does not redeem previous output \Rightarrow hash all zeroes
- ④ as of June 2020, the output is 6.25 BTC plus transaction fees
- ⑤ there is a special `coinbase` parameter which is arbitrary When the nonce is not possible

Contents

1 Bitcoin Blocks

2 Bitcoin Network

- Network Protocol
- SPV and Bloom Filters

The Bitcoin Network

- Bitcoin relies on a **peer-to-peer network**

The Bitcoin Network

- Bitcoin relies on a **peer-to-peer network**
- It consists of peers (nodes) having equal rights
⇒ no notion of servers or clients


The Bitcoin Network

- Bitcoin relies on a **peer-to-peer network**
- It consists of peers (nodes) having equal rights
⇒ no notion of servers or clients
- Nodes are connected, not necessarily by a direct link

The Bitcoin Network

- Bitcoin relies on a **peer-to-peer network**
- It consists of peers (nodes) having equal rights
⇒ no notion of servers or clients
- Nodes are connected, not necessarily by a direct link
- An overlay network independent of the underlying network (i.e., the Internet)

Nodes Types and Roles



The **reference client** (Bitcoin core) includes all roles


Wallet safe keeping of private keys for end users

Miner minting new bitcoins by creating new blocks


Blockchain verification of block traffic, keeping a copy of the blockchain

Network routing messages, maintaining connectivity needed for every node

Full node




Solo miner



- Alternative: miner farms
- clients do just mining and networking
- coordinated by a (full node) server
- miner farms use a different protocol

Light Weight (SPV) Wallet



- SPV = Simplified Payment Verification
 - only block headers are kept locally, not the full chain
- ⇒ 1000x decrease in storage
- do not have full picture of all UTXOs
 - need to obtain transaction data from surrounding full nodes to verify, but is prone to double spending attacks
 - asking for specific transactions can compromise anonymity ⇒ Bloom filters


Contents

1 Bitcoin Blocks

2 Bitcoin Network

- Network Protocol
- SPV and Bloom Filters

Network Discovery




the version message for initial contact with at least one bitcoin node

```
1 {"PROTOCOL_VERSION": 70015,  
2   "nLocalServices": 1033,  
3   "nTime": 1591903131,  
4   "addrYou": "122.51.104.28",  
5   "addrMe": "82.38.163.188",  
6   "subver": "/Satoshi:0.20.0/",  
7   "BestHeight": 634259  
8 }
```


second part is optional

Address Propagation and Discovery




- When B receives A's `getaddr` message, it randomly selects up to 2500 recent addresses from its pool and sends them to A
- A connection times out after 90 minutes

Node Synchronization



- (new node start with just the genesis block)
- after connection establishment, peers exchange hash of top block
- peer with the longer chain sends `inv` (up to 300 hashes)
- `getdata` requests blocks by hash, spreading the load

Header Synchronization (SPV Nodes)



- SPV nodes use `getheaders` to obtain only the headers
- The `headers` message contains up to 2000 headers

Contents

1 Bitcoin Blocks

2 Bitcoin Network

- Network Protocol
- SPV and Bloom Filters

SPV can compromise anonymity

- SPV nodes need to retrieve blocks / transactions in connection with the payment they are processing
- This selective download can reveal the addresses in the node's wallet
- To avoid revelation, SPV nodes request transaction data using **Bloom filters**
- A Bloom filter narrows down the set of transactions without revealing the selection criterion

Bloom Filters

Definition

A Bloom filter `bf` is a probabilistic datastructure that provides a fuzzy encoding of a set. Its operations are

- `bf.add(x)` includes `x` in the Bloom filter.
- `bf.elem(x)` checks membership of `x`. If it returns `False`, then `x` has never been added to the filter. If `x` has been added, then it returns `True`.

Bloom Filters

Definition

A Bloom filter `bf` is a probabilistic datastructure that provides a fuzzy encoding of a set. Its operations are

- `bf.add(x)` includes `x` in the Bloom filter.
- `bf.elem(x)` checks membership of `x`. If it returns `False`, then `x` has never been added to the filter. If `x` has been added, then it returns `True`.

Remark

- If `bf.elem(x)` returns `True`, `x` may or may not have been added.
- A negative result is definitive.


Implementation of a Bloom Filter

- bitvector of size m , initialized to all 0
- k different independent hash functions $h_i : A \rightarrow \{0, \dots, m - 1\}$
- to add x
 - ▶ compute $y_i = h_i(x)$ for all k hash functions
 - ▶ set bits y_1, \dots, y_k
- to check (potential) presence of x
 - ▶ compute $y_i = h_i(x)$ for all k hash functions
 - ▶ return False if any of the bits y_i is 0
 - ▶ otherwise return True
- the false positive rate


$$(1 - e^{-kn/m})^k$$

depends on the number n of elements that have been added to the filter


Example




- Bloom filter with three hash functions




Example



- Bloom filter with three hash functions
- add A




Example




- Bloom filter with three hash functions
- add A


Example



- Bloom filter with three hash functions
- add A
- add B




Example



- Bloom filter with three hash functions
- add A
- add B


Example



- Bloom filter with three hash functions
- add A
- add B
- check for X



1	0	1	0	0	0	1	0	0	0	0	0	0	1	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Example




- Bloom filter with three hash functions
- add A
- add B
- check for X
- may be an element

Example



- Bloom filter with three hash functions
- add A
- add B
- check for X
- may be an element
- check for Y

Example



- Bloom filter with three hash functions
- add A
- add B
- check for X
- may be an element
- check for Y
- definitely not an element

Pools

- (in main memory of the node)
- Every node maintains a **transaction pool** of unconfirmed, validated transactions
- These are propagated to the neighbors and may be different on each node
- Every node maintains a pool of UTXOs (confirmed, valid outputs)
- They represent the consensus of the longest chain

Thanks!