

Foundations of Artificial Intelligence

Riccardo Salvalaggio

19th of April, 2021

Contents

1	Introduction	4
2	Rational Agents	4
2.1	Structure of Rational Agents	4
2.2	Classes of agents	4
2.2.1	Table-Driven (the simplest)	4
2.2.2	Interpretative Reflex	4
2.2.3	Model-Based Reflex	4
2.2.4	Learning agents	4
2.3	Types of environments	5
3	Solving Problems by Searching	5
3.1	Problem-solving agents	5
3.2	Problem Formulation	6
3.3	Problem Types	6
3.4	Search strategies	6
3.4.1	Uninformed or blind searches	7
4	Informed Search Methods	8
4.1	Greedy Search	8
4.2	Heuristics	8
4.3	A* and IDA*	9
4.3.1	Optimality of A*	9
4.3.2	Graph- vs. Tree-search	9
4.3.3	Variants of A*	9
4.4	Local Search Methods	9
4.5	Genetic Algorithms	10
5	Board games	10
5.1	Minimax search	10
5.2	Alpha-Beta Search	10
5.3	Games with an Element of Chance	11
6	Constraint satisfaction problems	11
6.1	Backtracking search for CSPs	12
6.2	CSP Heuristics	12
6.3	Constraint Propagation	12
6.4	Problem Structure	13
7	Propositional logic	13
7.1	Agents that Think Rationally	13
7.1.1	Syntax	14
7.1.2	Normal forms	14
8	Satisfiability and Model Construction	15
8.1	Davis-Putnam-Logemann-Loveland (DPLL) Procedure	16
8.1.1	DPLL on Horn Clauses	16
8.2	Local Search	17
8.3	State of the Art	17
8.3.1	Improvements of DPLL Algorithms	17
8.3.2	Algorithm configuration	18
8.3.3	Algorithm selection	18
8.3.4	Automated construction of portfolios from a single algorithm	18

9 Predicate Logic (First-Order Predicate Logic (PL-1))	19
9.1 Semantics of PL1-Logic	19
9.1.1 Variable assignment	20
9.1.2 Satisfiability	20
9.2 Free and Bound Variables	20
9.3 Derivation in PL1	20
9.3.1 Finite universes	20
9.3.2 Instantiation	21
10 Action Planning	21
10.1 Planning Formalisms	21
10.1.1 Domain-Independent Action Planning	21
10.1.2 Planning as Logical Inference	21
10.1.3 Basic STRIPS Formalism	21
10.1.4 PDDL: The Planning Domain Description Language	23
10.2 Basic Planning Algorithms	23
10.3 Computational Complexity	24
10.4 Current Algorithmic Approaches	25
11 Making simple decisions under uncertainty (Probability)	25
11.1 Foundations of Probability Theory	26
11.2 Probabilistic Inference	26
11.2.1 Joint Probability	26
11.2.2 Bayes' Rule	27
11.2.3 Relative Probability	28
11.2.4 Normalization	28
11.2.5 Multiple Evidence	29
11.2.6 Recursive Bayesian Updating	30
11.3 Bayesian Network	30
12 Acting under Uncertainty Maximizing Expected Utility	32
12.1 Introduction to Utility theory	32
12.2 Choosing Individual Actions	32
12.3 Sequential Decision Problems	33
12.4 Markov Decision Processes	33
12.5 Value Iteration	34
12.5.1 Policy Iteration	36
13 Machine Learning	37
13.1 The learning agent	37
13.2 Types of learning	37
13.3 Decision trees	38
13.3.1 Boolean Decision Trees	38

1 Introduction

Artificial Intelligence is the attempt to make computers more "intelligent" to better understand human intelligence.

2 Rational Agents

It is a model that perceives the environment through sensors and act through actuators. In order to evaluate their performance use performance measure (e.g. vacuum -> level of cleanliness etc.) even if optimal behaviour is often unattainable because it is quite impossible to reach the goal in every aspect.

Omniscient if it knows the effects of its actions.

Rational agent behaves according to its percepts and knowledge and attempts to maximize the expected performance.

Ideal: for each possible percept sequence, selects an action that is expected to maximize its performance measure.

2.1 Structure of Rational Agents

The mapping is realised through an agent program executed on an Architecture which also provides an interface to the environment (percepts, actions) **Agent = Architecture + Program**

2.2 Classes of agents

2.2.1 Table-Driven (the simplest)

```
function TABLE-DRIVEN-AGENT(percept) returns an action
  persistent: percepts, a sequence, initially empty
               table, a table of actions, indexed by percept sequences, initially fully specified

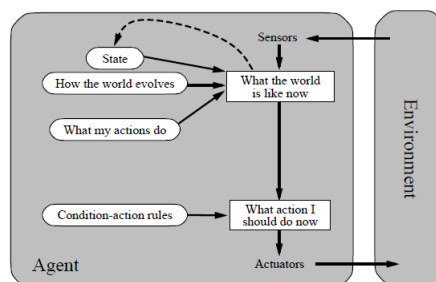
  append percept to the end of percepts
  action ← LOOKUP(percepts, table)
  return action
```

Problem: need a huge table to fulfill all the possible perceptions.

2.2.2 Interpretative Reflex

```
function SIMPLE-REFLEX-AGENT(percept) returns an action
  persistent: rules, a set of condition-action rules

  state ← INTERPRET-INPUT(percept)
  rule ← RULE-MATCH(state, rules)
  action ← rule.ACTION
  return action
```



Interpretation of the input, matching to a rule to extract an action.

2.2.3 Model-Based Reflex

Introduction of a utility function that maps a state onto a real number in order to compute the best action to do and to weigh the importance of competing goals.

2.2.4 Learning agents

Agents that improve over time starting from an empty knowledge and unknown environments.

Components:

Learning element: responsible for making improvements.

function MODEL-BASED-REFLEX-AGENT(*percept*) **returns** an action
persistent: *state*, the agent's current conception of the world state
model, a description of how the next state depends on current state and action
rules, a set of condition-action rules
action, the most recent action, initially none

state ← UPDATE-STATE(*state*, *action*, *percept*, *model*)
rule ← RULE-MATCH(*state*, *rules*)
action ← *rule*.ACTION
return *action*

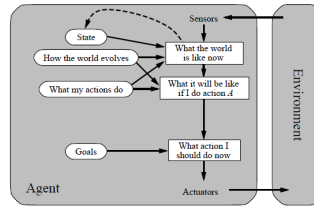


Figure 1: Goal-based

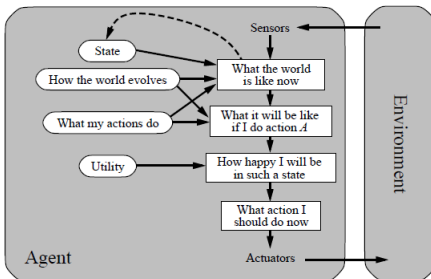
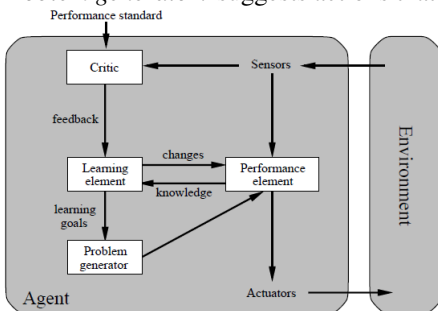


Figure 2: Utility-based

Performance element: select external actions.

Critic: determines performance of the agent.

Problem generator: suggests actions that lead to informative experiences.



2.3 Types of environments

Accessible vs. inaccessible, Deterministic vs. stochastic, Episodic vs. sequential, static vs. dynamic, discrete vs. continuous, single vs. multi agent.

3 Solving Problems by Searching

3.1 Problem-solving agents

Formulation: *problem* as a *state-space* and *goal* as a *particular condition on states*

Given: *initial state*

Goal: To reach the specified goal (a state) through the *execution of appropriate actions*

- Properties: Fully-observable, Deterministic/static env., discrete states, single-agent.

```

function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  persistent: seq, an action sequence, initially empty
               state, some description of the current world state
               goal, a goal, initially null
               problem, a problem formulation

  state  $\leftarrow$  UPDATE-STATE(state, percept)
  if seq is empty then
    goal  $\leftarrow$  FORMULATE-GOAL(state)
    problem  $\leftarrow$  FORMULATE-PROBLEM(state, goal)
    seq  $\leftarrow$  SEARCH(problem)
    if seq = failure then return a null action
  action  $\leftarrow$  FIRST(seq)
  seq  $\leftarrow$  REST(seq)
  return action

```

Figure 3: Simple Problem-solving Agent

3.2 Problem Formulation

Goal formulation, definition of: State space, actions, problem type, search and execution costs.

3.3 Problem Types

Based on knowledge of States and Actions: Observability, completeness of knowledge about world state and actions. (e.g. If the environment is completely observable, the vacuum cleaner always knows where it is and where the dirt is.)

Transition Model: Description of the outcome of an action.

Solution: Path from the initial to a goal state.

Search Costs: Time and storage requirements to find a solution.

Total Costs: Search costs + path costs.

Alternative formulations can influence a lot number of states, e.g. 8 queens problem: Naive - billions of state, Better - 2057 states.

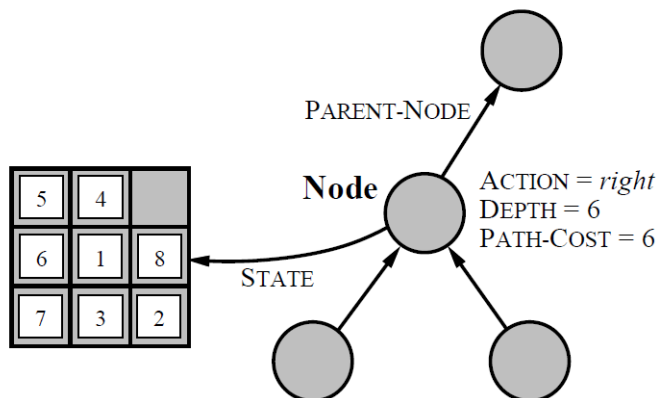
Examples of Real-World Problems: Route planning, shortest path problem, TSP, VLSI Layout, Robot nav., Assembly sequencing.

3.4 Search strategies

E.g.: node expansion, frontier, search strategy, tree-based search, graph-based search.

- Search Tree Data structure:

state, parent, action, path-cost.



```

function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    ⇒ choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier

```

```

function GRAPH-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  initialize the explored set to be empty
  loop do
    if the frontier is empty then return failure
    → choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    add the node to the explored set
    expand the chosen node, adding the resulting nodes to the frontier
    → only if not in the frontier or explored set

```

- Criteria for Search Strategies: **Completeness, Time complexity, Space Complexity, Optimality.**

3.4.1 Uninformed or blind searches

- Breadth-First Search:

Nodes are expanded in the order they were produced (first siblings, then children) (frontier = FIFO queue). Completeness is obvious, the solution is optimal. **Time complexity:** Let b be the maximal branching factor and d the depth of a solution path. Then the maximal number of nodes expanded is $O(b^d)$. **Space Complexity:** $O(b \cdot d)$

```

function BREADTH-FIRST-SEARCH(problem) returns a solution, or failure
  node ← a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  frontier ← a FIFO queue with node as the only element
  explored ← an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node ← POP(frontier) /* chooses the shallowest node in frontier */
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child ← CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)
        frontier ← INSERT(child, frontier)

```

- Uniform-Cost Search:

If step costs are different, uniform cost is better. It expands node with lowest path costs $g(n)$. It uses a Priority queue. Always finds the cheapest solution, given that $g(\text{successor}(n)) \geq g(n)$ for all n .

- Depth-First Search:

Always expands an unexpanded node at the greatest depth (frontier = a LIFO queue, first children, then siblings). Usually implemented recursively.

Generally, optimal is not guaranteed. Completeness only for graph-based search. **Time complexity:** in graph-based is bounded by the space, so it can be infinite, in tree-based: $O(b^m)$ (m max length of a path). **Space Complexity:** tree-based: $O(b \cdot m)$, graph-based: worst-case, all states need to be stored. (no better than breadth-first).

```

function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff
  return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  else if limit = 0 then return cutoff
  else
    cutoff_occurred? ← false
    for each action in problem.ACTIONS(node.STATE) do
      child ← CHILD-NODE(problem, node, action)
      result ← RECURSIVE-DLS(child, problem, limit - 1)
      if result = cutoff then cutoff_occurred? ← true
      else if result ≠ failure then return result
    if cutoff_occurred? then return cutoff else return failure

```

- Iterative Deepening Search:

Like depth-limited search and in every iteration increase search depth by one. Combines depth and breadth-first. Optimal and complete like breadth-first, but requires much less memory: $O(b \cdot d)$.

```

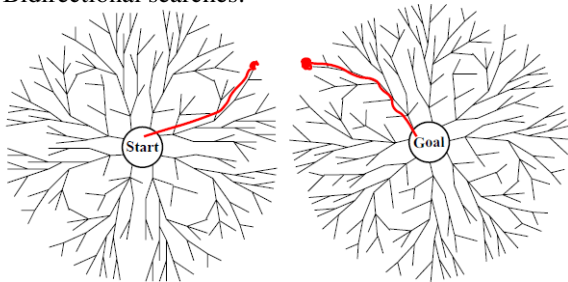
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result

```

Iterative deepening in general is the preferred uninformed search method when there is a large search space and the depth of the solution is not known.

For small space it is worse than breadth-first.

- Bidirectional searches:



As long as forward and backward searches are symmetric, search times of $O(2 \cdot b \text{ at } d/2) = O(b \text{ at } d/2)$ can be obtained. The operators are not always reversible, there must be an efficient way to check if a new node already appears in the search tree of the other half of the search.

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

4 Informed Search Methods

- Uninformed: rigid procedure with no knowledge of the cost of a given node to the goal.
- Informed: knowledge of the worth of expanding a node n is given in the form of an evaluation function $f(n)$ which assigns a real number to each node. Mostly, $f(n)$ includes as a component a heuristic function $h(n)$, which estimates the costs of the cheapest path from n to the goal.
- Best-first: informed that expands with the best f -value.

```

function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
     $\rightarrow$  choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier

```

Instance of tree-search algorithm in which frontier is a priority queue. When f is always correct, we don't need to search.

4.1 Greedy Search

$h(n)$ = estimated path-costs from n to the goal. A best-first search using $h(n)$ (heuristic function) as evaluation function is called greedy search.

4.2 Heuristics

Heuristics are fast but in certain situations incomplete methods for problem-solving, they improve the search in the average-case and the time complexity. In general, not optimal and incomplete; graph-search version is complete only in

finite spaces.

4.3 A* and IDA*

A* combines greedy search with the uniform-cost search: always expand node with lowest $f(n) = g(n)$ (actual cost from start to n) + $h(n)$ (estimated cost to goal/optimistic estimate of the costs). A new h is *admissible* iff: $h(n) \leq h^*(n)$.

4.3.1 Optimality of A*

Claim: The first solution found has the minimum path cost.

Proof: Suppose there exists a goal node G with optimal path cost f^* , but A* has first found another node $G2$ with $g(G2) < f^*$. Let n be a node on the path from the start to G that has not yet been expanded. Since h is admissible, we have: $f(n) \leq f^* \rightarrow f(G2) \leq f(n) \rightarrow f(G2) \leq f^* \Rightarrow g(G2) \leq f^*$ Contradiction.

- Completeness: If a solution exists, A* will find it provided that (1) every node has a finite number of successor nodes, and (2) there exists a positive constant $\delta > 0$ such that every step has at least cost δ .
- Complexity: In general, still exponential in the path length of the solution (space, time), it depends on the choice of Heuristic used.

4.3.2 Graph- vs. Tree-search

For the graph-based variant, either needs to consider re-opening nodes from the explored set, when a better estimate becomes known, or needs to require stronger restrictions on the heuristic estimate: it needs to be consistent (iff for all actions a leading from s to s' : $h(s) - h(s') \leq c(a)$, where $c(a)$ denotes the cost of action a). Consistency implies admissibility, A* can still be applied if heuristic is not consistent but optimality is lost.

4.3.3 Variants of A*

In general suffers from exponential memory growth. - Iterative-deepening A*: f -costs are used to define the cut-off (IDA*).

- Recursive Best First Search (RBFS): introduces a variable *f-limit* to keep track of the best alternative path, if the limit is exceeded opt for the alternative path.
- MA* and SMA*.

4.4 Local Search Methods

In many problems we are interested only to solving it, not how. If there is also a quality measure, local search can be used. It uses a Hill Climbing/ Gradient descent mechanisms (improvements step by step), requires few memory because it operates using just the current node.

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

```
current ← MAKE-NODE(problem.INITIAL-STATE)
loop do
  neighbor ← a highest-valued successor of current
  if neighbor.VALUE > current.VALUE then return current.STATE
  current ← neighbor
```

- Problem: the algorithm can stop in a sub-optimal solution (local maxima), the algorithm explore at random (plateaus), requires also suboptimal moves (ridges).
- Solution: restart if no improvements, inject noise.

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state
inputs: *problem*, a problem
schedule, a mapping from time to "temperature"

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

for *t* = 1 **to** ∞ **do**

T \leftarrow *schedule*(*t*)

if *T* = 0 **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ *next*.VALUE - *current*.VALUE

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$



4.5 Genetic Algorithms

Similar to evolution, we search for solutions by three operators: "mutation", "crossover", and "selection". Need of coding a solution into a string of symbols or bit-string, fitness function to judge the worth of configurations.

Example: Represent an individual (one 8-queens configuration) as a concatenation of eight x-y coordinates. Its fitness is judged by the number of non-attacks. The population consists of a set of configurations.

5 Board games

The states of a game are easy to represent. The possible actions of the players are well-defined.

The game can be implemented as a kind of search problem, the individual states are fully accessible and is nonetheless a contingency problem, because the actions of the opponent are not under the control of the player.

Board games are not only difficult because they are contingency problems, but also because the state space can become astronomically large.

5.1 Minimax search

In contrast to regular searches, where a path from beginning to end is a solution, max must come up with a strategy to reach a favorable terminal state regardless of what min does – > all of min moves must be considered and reactions to them must be computed.

When it is possible to produce the full game tree, the minimax algorithm delivers an optimal strategy for max.

- Algorithm:

1. Generate the complete tree with depth-first search.

2. Apply utility function.

3. Determine the utility: if predecessor is MIN, assign minimum value, if predecessor is MAX, assign maximum value.

Predecessor node is a max-node: Value is the maximum of its child nodes From the initial state (root of the game tree), max chooses the move that leads to the highest value (minimax decision).

function MINIMAX-DECISION(*state*) **returns** an action
return arg max_{*a* ∈ ACTIONS(*s*)} MIN-VALUE(RESET(*state*, *a*))

function MAX-VALUE(*state*) **returns** a utility value
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

v $\leftarrow -\infty$

for each *a* **in** ACTIONS(*state*) **do**

v \leftarrow MAX(*v*, MIN-VALUE(RESET(*s*, *a*)))

return *v*

function MIN-VALUE(*state*) **returns** a utility value
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

v $\leftarrow \infty$

for each *a* **in** ACTIONS(*state*) **do**

v \leftarrow MIN(*v*, MAX-VALUE(RESET(*s*, *a*)))

return *v*

5.2 Alpha-Beta Search

When the tree is becoming too large, we have to evaluate correctly where to expand it. The design of the evaluation function is fundamental!

Utility function for chance nodes C over MAX:

d_i : possible dice roll

$P(d_i)$: probability of obtaining that roll

$S(C, d_i)$: attainable positions from C with roll d_i

UTILITY(s): Evaluation of s

$$\text{EXPECTIMAX}(C) = \sum_i P(d_i) \max_{s \in S(C, d_i)} (\text{UTILITY}(s))$$

EXPECTIMIN likewise

Figure 4: Expected value

Standard evaluation functions are weighted linear - Alpha-Beta pruning: α : value of the best choice for MAX, β : value of the best choice for MIN.

(1) Prune below the min node whose β

-bound is less than or equal to the α -bound of its max-predecessor node. (2) Prune below the max node whose α

-bound is greater than or equal to the β

-bound of its min-predecessor node.

function ALPHA-BETA-SEARCH($state$) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(state, -\infty, +\infty)$
return the action in ACTIONS($state$) with value v

function MAX-VALUE($state, \alpha, \beta$) **returns** a utility value
if TERMINAL-TEST($state$) **then** **return** UTILITY($state$)
 $v \leftarrow -\infty$
for each a **in** ACTIONS($state$) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
if $v \geq \beta$ **then** **return** v
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
return v

function MIN-VALUE($state, \alpha, \beta$) **returns** a utility value
if TERMINAL-TEST($state$) **then** **return** UTILITY($state$)
 $v \leftarrow +\infty$
for each a **in** ACTIONS($state$) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
if $v \leq \alpha$ **then** **return** v
 $\beta \leftarrow \text{MIN}(\beta, v)$
return v

The alpha-beta search cuts the largest amount off the tree when we examine the best move first. Best case: $O(b \text{ at } d/2)$. Average case: $O(b \text{ at } 3d/4)$. Practical case: A simple ordering heuristic brings the performance close to the best case.

5.3 Games with an Element of Chance

In addition to min- and max nodes, we need chance nodes (e.g. for the dice).

Search costs increase: Instead of $O(b \text{ at } d)$, we get $O((b \times n) \text{ at } d)$, where n is the number of possible dice outcomes.

6 Constraint satisfaction problems

Such problem is defined by:

- variables, a set of value domains, a set of constraints, an assignment.

Main idea is to exploit the constraints to delete large portions of search space.

A constraint graph can be used to visualize binary constraints, Nodes = variables, arcs = constraints.

State: a variable assignment.

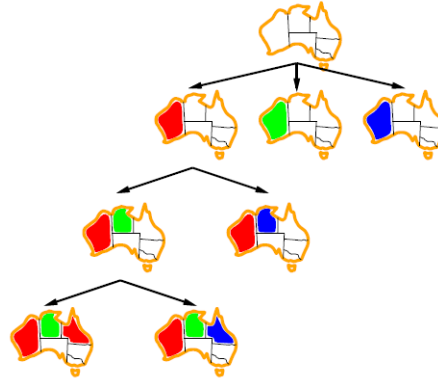
6.1 Backtracking search for CSPs

It assigns values to variables step by step using DFS single-variable (Backtracking search).

```

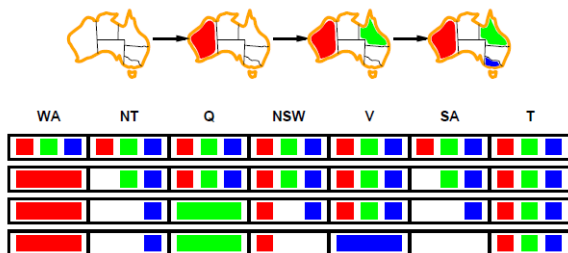
function BACKTRACKING-SEARCH(csp) returns a solution, or failure
    return BACKTRACK({ }, csp)

function BACKTRACK(assignment, csp) returns a solution, or failure
    if assignment is complete then return assignment
    var ← SELECT-UNASSIGNED-VARIABLE(csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
        if value is consistent with assignment then
            add {var = value} to assignment
            inferences ← INFERENCE(csp, var, value)
            if inferences ≠ failure then
                add inferences to assignment
                result ← BACKTRACK(assignment, csp)
                if result ≠ failure then
                    return result
            remove {var = value} and inferences from assignment
    return failure
    
```



6.2 CSP Heuristics

- Variable ordering: e.g.: most constrained first: choose the variable with the fewest remaining legal value.
- Value ordering: e.g.: Least constraining value first - Rule out failures early: e.g. Forward checking:



6.3 Constraint Propagation

A problem of Forward Checking is that constraints of unassigned variables are not propagated.

A directed arc $X \rightarrow Y$ is consistent iff for every choice of x there exists y that satisfies constraints; It detects failures earlier, can be used as preprocessing.

```

function AC-3(csp) returns false if an inconsistency is found and true otherwise
    inputs: csp, a binary CSP with components ( $X, D, C$ )
    local variables: queue, a queue of arcs, initially all the arcs in csp

    while queue is not empty do
        ( $X_i, X_j$ ) ← REMOVE-FIRST(queue)
        if REVISE(csp,  $X_i, X_j$ ) then
            if size of  $D_i$  = 0 then return false
            for each  $X_k$  in  $X_i$ .NEIGHBORS - { $X_j$ } do
                add ( $X_k, X_i$ ) to queue
    return true

    function REVISE(csp,  $X_i, X_j$ ) returns true iff we revise the domain of  $X_i$ 
        revised ← false
        for each  $x$  in  $D_i$  do
            if no value  $y$  in  $D_j$  allows ( $x, y$ ) to satisfy the constraint between  $X_i$  and  $X_j$  then
                delete  $x$  from  $D_i$ 
                revised ← true
        return revised
    
```

Check whether I have to remove a value from this domain.

Time complexity: $O($

$$d^3 * n^2)$$

$)$. Of course, AC-3 does not detect all inconsistencies (which is an NP-hard problem).

6.4 Problem Structure

If the CSP graph is a tree, then it can be solved in $O($

$$nd^2$$

$)$ (general CSPs need in the worst case $O($

$$d^n$$

$)$). Pick root, order nodes, apply arc consistency bottom-up, assign starting at root. This algorithm is linear in n .

Another solution is to reduce the graph structure by fixing values in a reasonably chosen subset. Instantiate a variable and prune values in neighboring variables (**conditioning**).

Almost Tree construction:

Algorithm [Cutset Conditioning](#):

- 1 Choose a subset S of the CSPs variables such that the constraint graph becomes a tree after removal of S . The set S is called a [cycle cutset](#).
- 2 For each possible assignment of variables in S that satisfies all constraints on S
 - 1 remove from the domains of the remaining variables any values that are inconsistent with the assignments for S , and
 - 2 if the remaining CSP has a solution, return it together with the assignment for S



- Tree Decomposition:

Decompose the problem into a set of connected (*share a constraint*) sub-problems and solve them independently. If a variable appears in two sub-problems, it must appear in all sub-problems on the path between the two sub-problems.

Consider sub-problems as new mega-variables, use tree-structured CSP to find solution.

The tree width w of a tree decomposition is the size of largest sub-problem minus 1, If a graph has tree width w and we know a tree decomposition with that width, we can solve the problem in $O($

$$nd^{(w+1)}$$

$)$.

7 Propositional logic

Logic is a universal tool with many powerful applications.

7.1 Agents that Think Rationally

Rational action requires rational (logical) thought on the agent's part so they must know a portion of the world in its near (**Knowledge base, KB**). KB is composed of sentences in a logic language.

- **Levels:**

1. *Knowledge level*: most abstract level, concerns the total knowledge.
2. *Logical level*: encoding of knowledge in a formal language.
3. *Implementation level*: the internal representation of the sentences. (as a string or as a value).

```

function KB-AGENT(percept) returns an action
  persistent: KB, a knowledge base
               t, a counter, initially 0, indicating time

  TELL(KB, MAKE-PERCEPT-SENTENCE(percept, t))
  action ← ASK(KB, MAKE-ACTION-QUERY(t))
  TELL(KB, MAKE-ACTION-SENTENCE(action, t))
  t ← t + 1
  return action

```

A logic also defines the semantics or meaning of sentences, defines the truth of a sentence with respect to each possible world. If a sentence α is true in a possible world m , we say that m satisfies α or m is a model of α .

- **Logical entailment:** $\alpha \models \beta$ iff in every model in which α is true, β is also true. (follow from KB) e.g.: $x = 0 \mid = xy = 0$.

- **Inference:** we can derive α with an inference method i . This is written as: $KB \vdash_i \alpha$. (Derivation)

- **Declarative Languages:** we state what we want to compute, not how. System believes P iff it consider P true. We must know symbols, when a sentence is true etc.

We'd like to have inference algorithms that derive only sentences that are entailed (*soundness*) and **all** of them (*completeness*). The building blocks of propositional logic are indivisible, atomic statements.

7.1.1 Syntax

An Interpretation I is called model of ω if $I \models \omega$. A truth assignment of the atoms in Σ , or an interpretation I over Σ , is a function: $I: \Sigma \rightarrow \{T, F\}$.

A formula ϕ can be: satisfiable, unsatisfiable, falsifiable (there exists I that doesn't satisfy ω), valid (tautology: $I \models \omega \forall I$), logically equivalent ($I \models \omega$ iff $I \models \chi \forall I$). A method to decide if a formula is satisfiable is: Truth Table.

7.1.2 Normal forms

- **conjunctive normal form (CNF):** conjunction of disjunctions:

$$\bigwedge_{i=1}^n \left(\bigvee_{j=1}^{m_i} l_{i,j} \right)$$

- **disjunctive normal form (DNF):** disjunction of conjunctions:

$$\bigvee_{i=1}^n \left(\bigwedge_{j=1}^{m_i} l_{i,j} \right)$$

It is always possible to transform a formula in a CNF/DNF but the conversion can make the sentence growing in an exponential way.

Some properties of logical implication:

- **Deduction theorem:** $KB \cup \{\varphi\} \models \psi$ iff $KB \models \varphi \Rightarrow \psi$
- **Contraposition theorem:** $KB \cup \{\varphi\} \models \neg\psi$ iff $KB \cup \{\psi\} \models \neg\varphi$
- **Contradiction theorem:** $KB \cup \{\varphi\}$ is unsatisfiable iff $KB \models \neg\varphi$

We can often derive new formulae from formulae in the KB. These new formulae should follow logically from the syntactical structure of the KB formulae.

Calculus: Set of inference rules (potentially including so-called logical axioms). A calculus C is sound (or correct) if all formulae that are derivable from a KB actually follow logically. A calculus is complete if every formula that follows logically from the KB is also derivable with C from the KB.

Resolution

We want a way to derive new formulae that does not depend on testing every interpretation (idea: to prove that $KB \models \phi$, we can prove that $KB \cup \neg\phi$ is unsatisfiable (contradiction theorem)).

$$\{(P \vee Q) \wedge (R \vee \neg P) \wedge S\} \text{ by } \{\{P, Q\}, \{R, \neg P\}, \{S\}\}$$

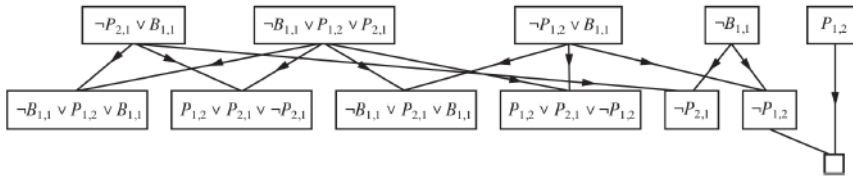
$$\frac{C_1 \dot{\cup} \{l\}, C_2 \dot{\cup} \{\bar{l}\}}{C_1 \cup C_2}$$

$C_1 \cup C_2$ are called **resolvents** of the **parent clauses** $C_1 \dot{\cup} \{l\}$ and $C_2 \dot{\cup} \{\bar{l}\}$. l and \bar{l} are the **resolution literals**.

function PL-RESOLUTION(KB, α) **returns** *true* or *false*
inputs: KB , the knowledge base, a sentence in propositional logic
 α , the query, a sentence in propositional logic

$clauses \leftarrow$ the set of clauses in the CNF representation of $KB \wedge \neg\alpha$
 $new \leftarrow \{ \}$
loop do
 for each pair of clauses C_i, C_j **in** $clauses$ **do**
 $resolvents \leftarrow$ PL-RESOLVE(C_i, C_j)
 if $resolvents$ contains the empty clause **then return** *true*
 $new \leftarrow new \cup resolvents$
 if $new \subseteq clauses$ **then return** *false*
 $clauses \leftarrow clauses \cup new$

Example:



Lemma (soundness) If $\Delta \vdash D$, then $\Delta \models D$.

Proof idea: Since all $D \in R(\Delta)$ follow logically from Δ , the Lemma results through induction over the length of the derivation.

Is resolution also complete? So the opposite implications are also valid. Not in general.

However, it can be shown that resolution is refutation-complete: Δ is unsatisfiable implies $\Delta \vdash \text{empty}$.

Theorem: Δ is unsatisfiable iff $\Delta \vdash \text{empty}$.

We can now infer new facts, but how do we translate knowledge into action?

Negative selection: Excludes any provably dangerous actions.

$A1,1 \wedge \text{EastA} \wedge W2,1 \Rightarrow \neg \text{Forward}$

Positive selection: Only suggests actions that are provably safe.

$A1,1 \wedge \text{EastA} \wedge \neg W2,1 \Rightarrow \text{Forward}$

8 Satisfiability and Model Construction

SAT solving is the best available technology for practical solutions to many NP-hard problems. Differently from Logical deduction, SAT returns a model of the theory (solution) given a logical theory. SAT can be formulated as a Constraint-Satisfaction-Problem (\rightarrow search):

- CSP-variables: alphabet
- Domain: True, False

- Constraints given by clauses

8.1 Davis-Putnam-Logemann-Loveland (DPLL) Procedure

It corresponds to backtracking with inference in CSPs. Inference in DPLL:

Simplify: if variable v is assigned a value d , then all clauses containing v are simplified immediately (corresponds to forward checking), variables in unit clauses are immediately assigned.

Given a set of clauses Δ defined over a set of variables Σ , return "satisfiable" if Δ is satisfiable. Otherwise return "unsatisfiable".

1. If $\Delta = \emptyset$ return "satisfiable"
2. If $\square \in \Delta$ return "unsatisfiable"
3. **Unit-propagation Rule:** If Δ contains a **unit-clause** C , assign a truth-value to the variable in C that satisfies C , simplify Δ to Δ' and return **DPLL**(Δ').
4. **Splitting Rule:** Select from Σ a variable v which has not been assigned a truth-value. Assign one truth value t to it, simplify Δ to Δ' and call **DPLL**(Δ')
 - a. If the call returns "satisfiable", then return "satisfiable".
 - b. Otherwise assign the other truth-value to v in Δ , simplify to Δ'' and return **DPLL**(Δ'').

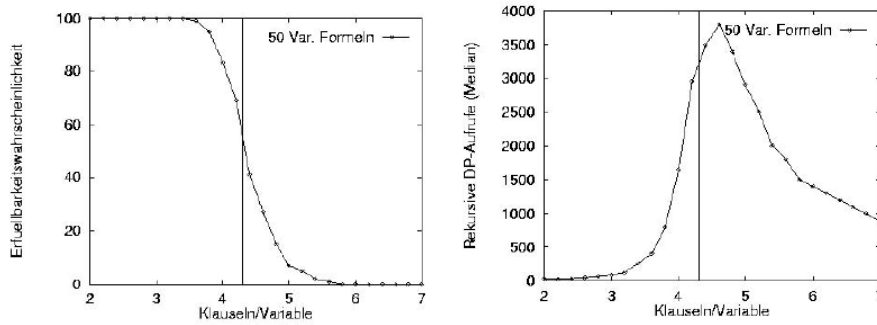
DPLL is complete, correct, and guaranteed to terminate. DPLL constructs a model, if one exists. In general, DPLL requires exponential time (splitting rule!), Heuristics are needed. DPLL is polynomial on Horn clauses. Horn Clauses constitute an important special case, since they require only polynomial runtime of DPLL. A **Horn clause** is a clause with maximally one positive literal.

8.1.1 DPLL on Horn Clauses

- **1.** The simplifications in DPLL on Horn clauses always generate Horn clauses
- **2.** If unit propagation rule in DPLL does not lead to termination, a set of Horn clauses without unit clauses is generated.
- **3.** A set of Horn clauses without unit clauses and without the empty clause is satisfiable, since:
All clauses have at least one negative literal.
Assigning false to all variables satisfies formula.
- **4.** It follows from 3.:
 - a. every time the splitting rule is applied, the current formula is satisfiable.
 - b. every time, when the wrong decision is made, this will be immediately detected.
- **5.** Therefore, the search trees for n variables can only contain a maximum of n nodes, in which the splitting rule is applied (and the tree branches).
- **6.** Therefore, the size of the search tree is only polynomial in n and therefore the running time is also polynomial.

For CNF-formulae, in which the probability for a positive appearance, negative appearance and non-appearance in a clause is $1/3$, DPLL needs on average quadratic time. All NP-complete problems have at least one order parameter and the hard to solve problems are around a critical value of this order parameter. This critical value (a phase transition) separates one region from another, such as over-constrained and under-constrained regions of the problem space.

When the probability of a solution is close to 1 (**under-constrained**), there are many solutions. If the probability of a solution is close to 0 (**over-constrained**).



8.2 Local Search

In many cases, we can sacrifice completeness if we can "solve" much larger instances this way. Standard process for optimization problems: **Local Search**. Problem: local minima. However: By restarting and/or injecting noise, we can often escape local maxima. Local search can perform very well for SAT solving.

Procedure GSAT

INPUT: a set of clauses α , MAX-FLIPS, and MAX-TRIES

OUTPUT: a satisfying truth assignment of α , if found

begin

for $i := 1$ to MAX-TRIES

$T :=$ a randomly-generated truth assignment

for $j := 1$ to MAX-FLIPS

if T satisfies α **then return** T

$v :=$ a propositional variable such that a change in its truth assignment gives the largest increase in the number of clauses of α that are satisfied by T

$T := T$ with the truth assignment of v reversed

end for

end for

return "no satisfying assignment found"

end

8.3 State of the Art

8.3.1 Improvements of DPLL Algorithms

Branching on variables can cause stopping, we can "learn" (**here: logically infer**) a new clause (negation of a variable). Leads to conflict-directed clause learning (CDCL).

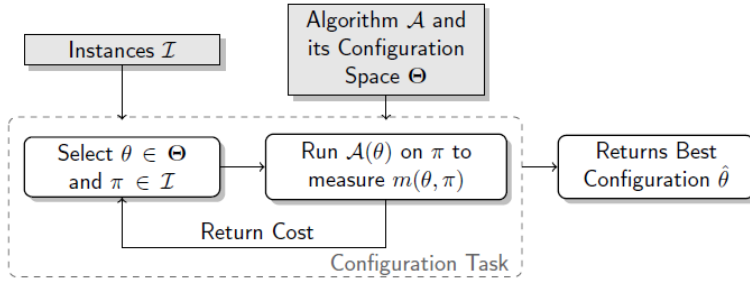
Both for DPLL/CDCL algorithms and local search algorithms

- Randomization and restarts
- Efficient data structures and indexing
- Engineering ingenious heuristics

Meta-algorithmic advances

- Automated parameter tuning and algorithm configuration
- Selection of the best-fitting algorithm based on instance characteristics
- Selection of the best-fitting parameters based on instance characteristics
- Use of machine learning to pinpoint what factors most affects performance.

8.3.2 Algorithm configuration



Definition: algorithm configuration

Given:

- a parameterized algorithm \mathcal{A} with possible parameter settings Θ ;
- a distribution \mathcal{D} over problem instances with domain \mathcal{I} ; and
- a cost metric $m : \Theta \times \mathcal{I} \rightarrow \mathbb{R}$,

Find: $\theta^* \in \arg \min_{\theta \in \Theta} \mathbb{E}_{\pi \sim \mathcal{D}}(m(\theta, \pi))$.

- Strength:

find a single configuration with strong performance for a given cost metric.

- Weakness:

for heterogeneous instance sets, there is often no configuration that performs great for all instances.

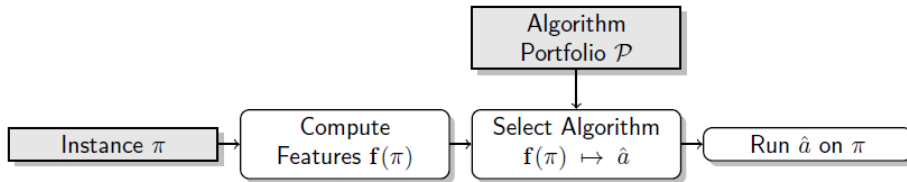
8.3.3 Algorithm selection

Definition: algorithm selection

Given

- a set \mathcal{I} of problem instances,
- a portfolio of algorithms \mathcal{P} ,
- and a cost metric $m : \mathcal{P} \times \mathcal{I} \rightarrow \mathbb{R}$,

the per-instance algorithm selection problem is to find a mapping $s : \mathcal{I} \rightarrow \mathcal{P}$ that optimizes $\sum_{\pi \in \mathcal{I}} m(s(\pi), \pi)$, the sum of cost measures achieved by running the selected algorithm $s(\pi)$ for instance π .



- **Strength:** for heterogeneous instance sets, pick the right algorithm from a set. - **Weakness:** the set to choose from typically only contains a few algorithms.

8.3.4 Automated construction of portfolios from a single algorithm

Putting the two together: Use algorithm configuration to determine useful configurations. Use algorithm selection to select from them based on instance characteristics.

Hydra

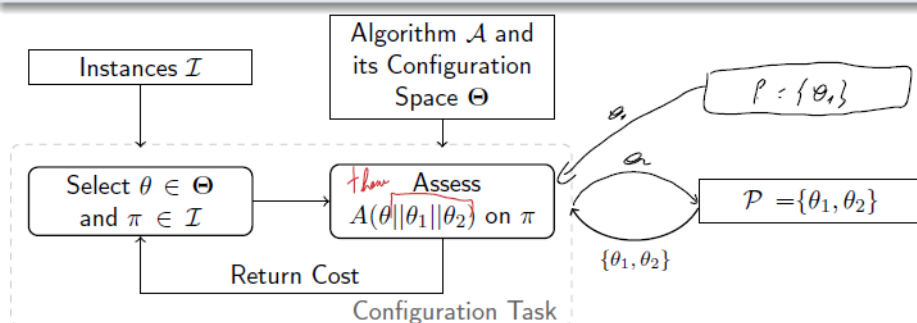
Idea

Iteratively add configurations to a portfolio \mathcal{P} , starting with $\mathcal{P} = \emptyset$

In each iteration, determine configuration that is complementary to \mathcal{P}

Maximize marginal contribution of configuration θ to current portfolio \mathcal{P} :

$$m(\mathcal{P}) - m(\mathcal{P} \cup \{\theta\})$$



9 Predicate Logic (First-Order Predicate Logic (PL-1))

Propositional logic has no structure in the atomic propositions.

In addition to Operators, Variables, Brackets we have Predicate and function. They have an arity (number of arguments): 0-ary predicate: propositional logic atoms, 0-ary function = constants.

Terms: Every variable is a term, If t_1, t_2, \dots, t_n are terms and f is an n -ary function, then $f(t_1, t_2, \dots, t_n)$ is also a term. Terms without variables: ground terms.

Atomic Formulae: Represent statements about objects.

- If t_1, t_2, \dots, t_n are terms and P is an n -ary predicate, then $P(t_1, t_2, \dots, t_n)$ is an atomic formula.
- If t_1 and t_2 are terms, then $t_1 = t_2$ is an atomic formula. Atomic formulae without variables: ground atoms (contain only ground terms).

9.1 Semantics of PL1-Logic

Interpretation: $\mathcal{I} = \langle D, x^{\mathcal{I}} \rangle$ D domain, $x^{\mathcal{I}}$ is a function that:

maps n -ary function symbols to functions over D :

$$f^{\mathcal{I}} \in [D^n \mapsto D]$$

maps individual constants to elements of D :

$$a^{\mathcal{I}} \in D$$

maps n -ary predicate symbols to relations over D :

$$P^{\mathcal{I}} \subseteq D^n$$

Interpretation of ground terms: $(f(t_1, \dots, t_n))^{\mathcal{I}} = f^{\mathcal{I}}(t_1^{\mathcal{I}}, \dots, t_n^{\mathcal{I}})$.

Satisfaction of ground atoms $P(t_1, \dots, t_n)$: $\mathcal{I} \models P(t_1, \dots, t_n)$ iff $\langle t_1^{\mathcal{I}}, \dots, t_n^{\mathcal{I}} \rangle \in P^{\mathcal{I}}$.

9.1.1 Variable assignment

Set of all variables V . Function $\alpha : V \mapsto D$

Notation: $\alpha[x/d]$ is the same as α apart from point x .

For $x : \alpha[x/d](x) = d$.

Interpretation of terms under I, α :

$$\begin{aligned} x^{I, \alpha} &= \alpha(x) \\ a^{I, \alpha} &= a^I \\ (f(t_1, \dots, t_n))^{I, \alpha} &= f^I(t_1^{I, \alpha}, \dots, t_n^{I, \alpha}) \end{aligned}$$

Satisfaction of atomic formulae:

$$I, \alpha \models P(t_1, \dots, t_n) \text{ iff } \langle t_1^{I, \alpha}, \dots, t_n^{I, \alpha} \rangle \in P^I$$

9.1.2 Satisfiability

A formula φ is **satisfied** by an **interpretation** I and a variable assignment α , i.e., $I, \alpha \models \varphi$:

$$\begin{aligned} I, \alpha &\models \top \\ I, \alpha &\not\models \perp \\ I, \alpha &\models \neg\varphi \text{ iff } I, \alpha \not\models \varphi \\ &\dots \end{aligned}$$

and all other propositional rules as well as

$$\begin{aligned} I, \alpha \models P(t_1, \dots, t_n) &\text{ iff } \langle t_1^{I, \alpha}, \dots, t_n^{I, \alpha} \rangle \in P^I \\ I, \alpha \models \forall x \varphi &\text{ iff for all } d \in D, I, \alpha[x/d] \models \varphi \\ I, \alpha \models \exists x \varphi &\text{ iff there exists a } d \in D \text{ with } I, \alpha[x/d] \models \varphi \end{aligned}$$

9.2 Free and Bound Variables

$$\forall x [R(\boxed{y}, \boxed{z}) \wedge \exists y ((\neg P(y, x) \vee R(y, \boxed{z})))]$$

When boxed, is called *free*, otherwise **bound**. Formulae with no free variables are called *closed formulae* or *sentences*.

With closed formulae, α can be left out on the left side of the model relationship symbol: $I \models \omega$. An interpretation I is called a **model** of ω under α if: $I, \alpha \models \omega$

9.3 Derivation in PL1

Reduction to propositional logic by instantiation based on the so-called **Herbrand Universe** (all possible terms) – \rightarrow infinite propositional theories.

Simple way for special case: If the number of objects is **finite**, instantiate all variables by possible objects.

9.3.1 Finite universes

Domain closure axiom (DCA): $\forall [x = c_1 \text{ or } \dots \text{ or } x = c_n]$

unique name assumption/axiom or UNA: $\text{and}_{i \neq j} [c_i \neq c_j]$

Eliminate quantification by instantiating all variables with all possible values.

9.3.2 Instantiation

Notation: if φ is a formula, then $\varphi[x/a]$ is the formula with all free occurrences of x replaced by a .

Universally quantified formulas are replaced by a conjunction of formulas with the variable instantiated to all possible values (from DCA):

$$\forall x \varphi \rightsquigarrow \bigwedge_i \varphi[x/c_i]$$

Existentially quantified variables are replaced by a disjunction of formulas with the variable instantiated to all possible values (from DCA):

$$\exists x \varphi \rightsquigarrow \bigvee_i \varphi[x/c_i]$$

Note: does blow up the formulas exponentially in the **arity** of the predicates!

10 Action Planning

Planning is the art and practice of thinking before acting.

Planning is the process of generating (possibly partial) representations of future behavior prior to the use of such plans to constrain or control that behavior.

Planning is not searching, synthesis or scheduling.

10.1 Planning Formalisms

10.1.1 Domain-Independent Action Planning

- Start with a declarative specification of the planning problem
- Use a domain-independent planning system to solve the planning problem
- Domain-independent planners are generic problem solvers

10.1.2 Planning as Logical Inference

Initial state:

$$At(truck1, loc1, s_0) \wedge At(package1, loc3, s_0)$$

Operators (successor-state axioms):

$$\forall a, s, l, p, t \quad At(t, p, Do(a, s)) \Leftrightarrow \{a = Drive(t, l, p) \wedge Poss(Drive(t, l, p), s) \\ \vee At(t, p, s) \wedge (a \neq \neg Drive(t, p, l, s) \vee \neg Poss(Drive(t, p, l, s), s))\}$$

Goal conditions (query):

$$\exists s \quad At(package1, loc2, s)$$

Inefficient for bigger sets.

10.1.3 Basic STRIPS Formalism

STRIPS: STanford Research Institute Problem Solver

S is a *first-order vocabulary* (predicate and function symbols) and Σ_S denotes the set of *ground atoms* over the signature (also called **facts** or **fluents**).

$\Sigma_{S,V}$ is the set of atoms over S using variable symbols from the set of variables V .

A **first-order STRIPS state** S is a subset of Σ_S denoting a *complete theory* or *model* (using CWA).

A **planning task** (or **planning instance**) is a 4-tuple

$\Pi = \langle S, O, I, G \rangle$, where

- O is a set of **operator** (or *action types*)
- $I \subseteq \Sigma_S$ is the **initial state**
- $G \subseteq \Sigma_S$ is the **goal specification**

No domain constraints (although present in original formalism)

- **Operator:** $o = \langle \text{para}, \text{pre}, \text{eff} \rangle$
- **Operator instance or action:** operator with empty parameter list

$$App(S, o) = \begin{cases} S \cup \text{eff}^+(o) - \neg \text{eff}^-(o) & \text{if } \text{pre}(o) \subseteq S \text{ \& } \text{eff}(o) \text{ is cons.} \\ \text{undefined} & \text{otherwise} \end{cases}$$

- **State change:**

Logical atoms: $at(O, L)$, $in(O, V)$, $airconn(L1, L2)$, $street(L1, L2)$, $plane(V)$, $truck(V)$

Load into truck: *load*

Parameter list: (O, V, L)

Precondition: $at(O, L), at(V, L), truck(V)$

Effects: $\neg at(O, L), in(O, V)$

Drive operation: *drive*

Parameter list: $(V, L1, L2)$

Precondition: $at(V, L1), truck(V), street(L1, L2)$

Effects: $\neg at(V, L1), at(V, L2)$

...

Some constant symbols: $v1, s, t$ with $truck(v1)$ and $street(s, t)$

Action: $drive(v1, s, t)$

A plan Δ is a sequence of actions:

$$Res(S, \langle \rangle) = S$$

$$Res(S, (o; \Delta)) = \begin{cases} Res(App(S, o), \Delta) & \text{if } App(S, o) \text{ is defined} \\ \text{undefined} & \text{otherwise} \end{cases}$$

A plan is successful or solves a planning task if has a result that respects goal specifications.

$$\text{Initial state: } S = \left\{ \begin{array}{l} at(p1, c), at(p2, s), at(t1, c), \\ at(t2, c), street(c, s), street(s, c) \end{array} \right\}$$

$$\text{Goal: } G = \{ at(p1, s), at(p2, c) \}$$

$$\text{Successful plan: } \Delta = \langle \text{load}(p1, t1, c), \text{drive}(t1, c, s), \text{unload}(p1, t1, s), \text{load}(p2, t1, s), \text{drive}(t1, s, c), \text{unload}(p2, t1, c) \rangle$$

```

(define (domain logistics)
  (:types truck airplane - vehicle
    package vehicle - physobj
    airport location - place
    city place physobj - object)

  (:predicates (in-city ?loc - place ?city - city)
    (at ?obj - physobj ?loc - place)
    (in ?pkg - package ?veh - vehicle))

  (:action LOAD-TRUCK
    :parameters (?pkg - package ?truck - truck ?loc - place)
    :precondition (and (at ?truck ?loc) (at ?pkg ?loc))
    :effect (and (not (at ?pkg ?loc)) (in ?pkg ?truck)))
    ...))

```

Figure 5: Logistics Example

STRIPS as described above allows for unrestricted first-order terms.

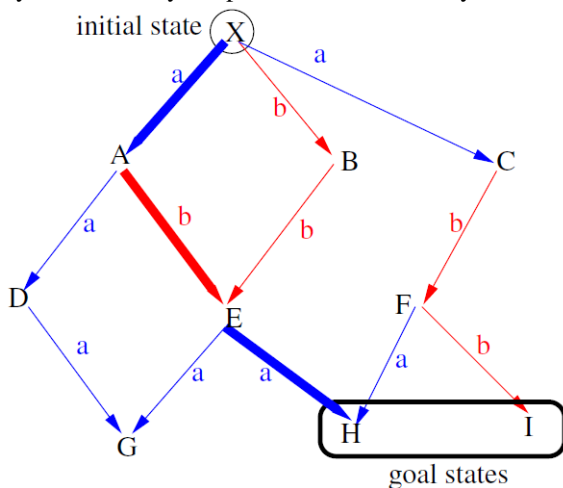
Simplifications:

- **1. Infinite state space:** No function terms (only 0-ary = constants).
- **2. DATALOG-STRIPS:** No variables in operators (= actions).
- **3. Propositional STRIPS:** used in planning algorithms nowadays.

10.1.4 PDDL: The Planning Domain Description Language

10.2 Basic Planning Algorithms

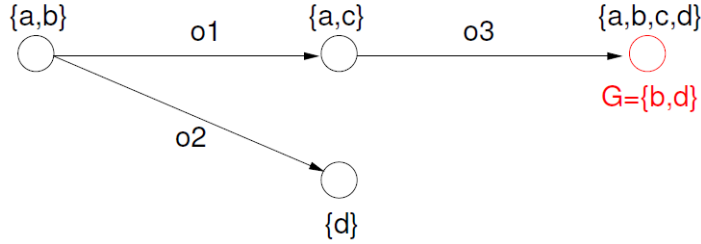
We can view planning problems as searching for goal nodes in a large labeled graph (transition system). Nodes are defined by the value assignment to the fluents (states). Labeled edges are defined by actions. Create the transition system on the fly and visit only the parts that are necessary.



- **Progression Planning: Forward Search** Start at initial state:
 1. Initialize: $\Delta = \langle \rangle$, add the initial state and make it S.
 2. Test whether we have reached goal: G in S? Return.
 3. Select one applicable action o_i non deterministically, compute successor $S = \text{App}(S, o_i)$, extend plan adding it and **come back to step 2**.

This algorithm can be easily extended to more expressive planning languages (not only boolean: goal or not).

$$\begin{aligned}
\mathcal{S} &= \{a, b, c, d\}, \\
\mathbf{O} &= \{ o_1 = \langle \emptyset, \{a, b\}, \{\neg b, c\} \rangle, \\
&\quad o_2 = \langle \emptyset, \{a, b\}, \{\neg a, \neg b, d\} \rangle, \\
&\quad o_3 = \langle \emptyset, \{c\}, \{b, d\} \rangle, \\
\mathbf{I} &= \{a, b\} \\
\mathbf{G} &= \{b, d\}
\end{aligned}$$



• **Regression Planning: Backward Search** Start from the goal:

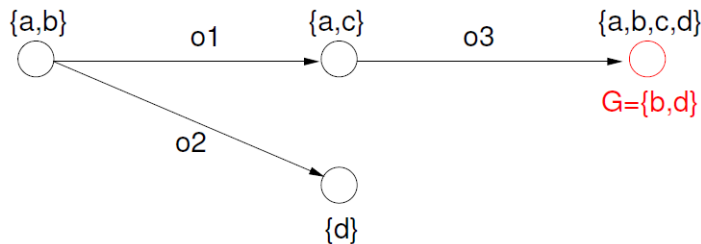
1. Initialize: $\Delta = \langle \rangle$, add the goal state and make it S.
2. Test whether we have reached initial: I in S? Return.
3. Select one applicable action o_i non deterministically which does not make sub-goals false, $S \cap \neg eff^-(o_i) = \emptyset$ and compute the regression of the description S through o_i :

$$S = S - eff^+(o_i) \cup pre(o_i)$$

then extend plan $\Delta = \langle o_i, \Delta \rangle$, back to step 2.

(Instead of non-deterministic we can use some search strategy).

$$\begin{aligned}
\mathcal{S} &= \{a, b, c, d\}, \\
\mathbf{O} &= \{ o_1 = \langle \emptyset, \{a, b\}, \{\neg b, c\} \rangle, \\
&\quad o_2 = \langle \emptyset, \{a, b\}, \{\neg a, \neg b, d\} \rangle, \\
&\quad o_3 = \langle \emptyset, \{c\}, \{b, d\} \rangle, \\
\mathbf{I} &= \{a, b\} \\
\mathbf{G} &= \{b, d\}
\end{aligned}$$



10.3 Computational Complexity

- **Definition (Plan existence problem (PLANEX)):** Does exist a plan that solve it?
- **Definition (Bounded plan existence problem (PLANLEN)):** Does there exist a plan of length n or less that solves it?

The state space for STRIPS with general first-order terms is infinite. The existence of a plan is then equivalent to the existence of a successful computation on the Turing machine.

Theorem: *PLANEX for STRIPS with first-order terms is undecidable.*

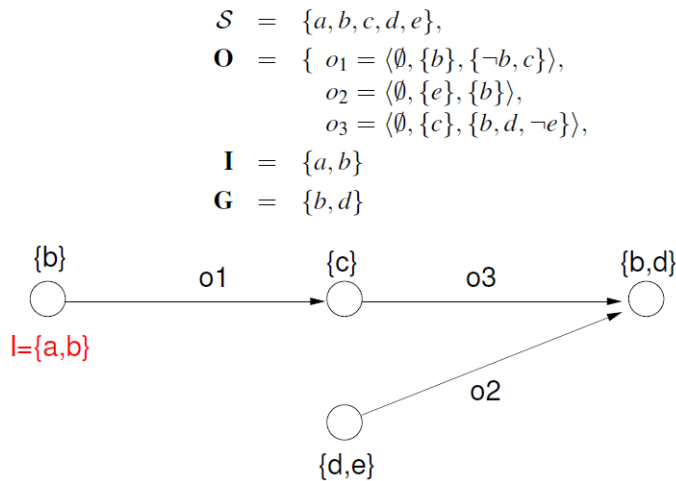
Theorem: *PLANEX is PSPACE-complete for propositional STRIPS.*

- **Restrictions on Plans** If we restrict the length of the plans to be short, i.e., only polynomial in the size of the planning task, PLANEX becomes NP-complete. If we use a unary representation of the natural number k , then PLANLEN becomes NP-complete. We can use methods for NP-complete problems if we are only looking for “short” plans.

10.4 Current Algorithmic Approaches

- **Heuristic Search Planning:** use an automatically generated heuristic estimator in order to select the next action or state. It is often easier to go for sub-optimal solutions (remember Logistics)
- **Deriving Heuristics: Relaxations** Define a simplification (relaxation) of the problem and take the difficulty of a solution for the simplified problem as an heuristic estimator. (E.g.: Straight line distance on a map to estimate the travel distance).

- **Ignoring Negative Effects: Example**



11 Making simple decisions under uncertainty (Probability)

In many cases, our knowledge of the world is incomplete (not enough information) or uncertain (sensors are unreliable). Without perfect knowledge, logical rules do not help much!

One possibility for expressing the degree of belief is to use probabilities. Probabilities quantify the uncertainty that stems from lack of knowledge.

$$\begin{aligned}
 \text{Decision Theory} &= \text{Utility Theory} + \text{Probability Theory} \\
 \text{argmax}_a \sum_{\omega} p(\omega|a) [U(\omega) - c(a)]
 \end{aligned}$$

- **Decision-Theoretic Agent**

In **Logistics**: The negative effects in *load* and *drive* are ignored:

Simplified load operation: $load(O, V, P)$

Precondition: $at(O, P), at(V, P), truck(V)$

Effects: $\neg at(O, P), in(O, V)$

After loading, the package is still at the place and also inside the truck

Simplified drive operation: $drive(V, P1, P2)$

Precondition: $at(V, P1), truck(V), street(P1, P2)$

Effects: $\neg at(V, P1), at(V, P2)$

After driving, the truck is in two places!

We want the length of the shortest **relaxed** plan $\rightsquigarrow h^+(s)$

How difficult is **monotonic planning**?

11.1 Foundations of Probability Theory

We use random variables such as Weather (capitalized word), which has a domain of ordered values. In our case that could be sunny, rain, cloudy, snow (lower case words).

- **Unconditional Probabilities:** $P(a)$ denotes the unconditional probability that it will turn out that $A = true$ in the absence of any other information.
- $P(a | b) = \frac{P(a \wedge b)}{P(b)}$ is the conditional or posterior probability of a given that all we know is b : $P(cavity|toothache) = 0.8$
- **Product rule:** $P(a \wedge b) = P(a|b)P(b)$

11.2 Probabilistic Inference

11.2.1 Joint Probability

The joint probability distribution $P(X1, ..., Xn)$ assigns a probability to every atomic event. Example of such a complete instantiation:

	<i>toothache</i>	\neg <i>toothache</i>
<i>cavity</i>	0.04	0.06
\neg <i>cavity</i>	0.01	0.89

All relevant probabilities can be computed using the joint probability by expressing them as a disjunction of atomic events.

Examples:

$$P(\text{cavity} \vee \text{toothache}) = P(\text{cavity} \wedge \text{toothache}) + P(\neg \text{cavity} \wedge \text{toothache}) + P(\text{cavity} \wedge \neg \text{toothache})$$

We obtain marginal probabilities by adding across a row or column:

$$P(\text{cavity}) = P(\text{cavity} \wedge \text{toothache}) + P(\text{cavity} \wedge \neg \text{toothache})$$

We obtain conditional probabilities by using a marginal probability:

$$P(\text{cavity} \mid \text{toothache}) = \frac{P(\text{cavity} \wedge \text{toothache})}{P(\text{toothache})} = \frac{0.04}{0.04 + 0.01} = 0.80$$

- **Total probability:** $P(Y) = \sum_z P(Y, z) = \sum_z P(Y|z)P(z)$

We can easily obtain all probabilities from the joint probability. The joint probability, however, involves kn values, if there are n random variables with k values.

$$\begin{aligned} P(x_1, \dots, x_n) &= P(x_n, \dots, x_1) = P(x_n \mid x_{n-1} \dots, x_1) P(x_{n-1}, \dots, x_1) \\ &= P(x_n \mid x_{n-1} \dots, x_1) P(x_{n-1} \mid x_{n-2} \dots, x_1) P(x_{n-2}, \dots, x_1) \\ &= P(x_n \mid x_{n-1} \dots, x_1) P(x_{n-1} \mid x_{n-2} \dots, x_1) P(x_{n-2} \mid x_{n-3} \dots, x_1) \\ &\quad P(x_{n-3}, \dots, x_1) \\ &= \dots \\ &= P(x_n \mid x_{n-1} \dots, x_1) P(x_{n-1} \mid x_{n-2} \dots, x_1) \dots P(x_2 \mid x_1) P(x_1) \\ &= \prod_{i=1}^n P(x_i \mid x_{i-1} \dots x_1) \end{aligned}$$

In order to save memory and time, Modern systems work directly with conditional probabilities and make assumptions on the independence of variables (!conditional independence) to simplify calculations.

11.2.2 Bayes' Rule

Deriving from Product rule:

$$P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)} \quad (1)$$

Generalizing:

$$P(Y|X, e) = \frac{P(X|Y, e)P(Y|e)}{P(X|e)} \quad (2)$$

- Example:

$$\begin{aligned}
 P(\text{toothache} \mid \text{cavity}) &= 0.4 \\
 P(\text{cavity}) &= 0.1 \\
 P(\text{toothache}) &= 0.05 \\
 \rightarrow P(\text{cavity} \mid \text{toothache}) &= \frac{0.4 \times 0.1}{0.05} = 0.8
 \end{aligned}$$

$P(\text{toothache} \mid \text{cavity})$ (causal) is more robust than $P(\text{cavity} \mid \text{toothache})$ (diagnostic).

11.2.3 Relative Probability

Let's say we would also like to consider the probability that our patient has gum disease.

$$\begin{aligned}
 P(\text{toothache} \mid \text{gumdisease}) &= 0.7 \\
 P(\text{gumdisease}) &= 0.02
 \end{aligned}$$

Which diagnosis is more probable? Cavity or gum disease?

$$P(c \mid t) = \frac{P(t \mid c)P(c)}{P(t)} \quad \text{or} \quad P(g \mid t) = \frac{P(t \mid g)P(g)}{P(t)}$$

If we are only interested in the **relative probability**, we need not assess $P(t)$:

$$\begin{aligned}
 \frac{P(c \mid t)}{P(g \mid t)} &= \frac{P(t \mid c)P(c)}{P(t)} \times \frac{P(t)}{P(t \mid g)P(g)} = \frac{P(t \mid c)P(c)}{P(t \mid g)P(g)} \\
 &= \frac{0.4 \times 0.1}{0.7 \times 0.02} = 2.857
 \end{aligned}$$

→ We elegantly excluded other possible diagnoses for toothache.

11.2.4 Normalization

To compute absolute probability of $P(c \mid t)$ without knowing $P(t)$ we can do a complete case analysis using the property that $P(c \mid t) + P(\neg c \mid t) = 1$.

$$\begin{aligned}
 P(c \mid t) &= \frac{P(t \mid c)P(c)}{P(t)} \\
 P(\neg c \mid t) &= \frac{P(t \mid \neg c)P(\neg c)}{P(t)} \\
 P(c \mid t) + P(\neg c \mid t) &= \frac{P(t \mid c)P(c)}{P(t)} + \frac{P(t \mid \neg c)P(\neg c)}{P(t)} \\
 P(t) &= P(t \mid c)P(c) + P(t \mid \neg c)P(\neg c)
 \end{aligned}$$

In general: $P(Y|X) = \alpha P(X|Y)P(Y)$

α is the normalization constant needed to make the entries in $P(X|Y)$ sum to 1 for each X . (e.g.: $\alpha(.1, .1, .3) = (.2, .2, .6)$)

Your doctor tells you that you have tested positive for a serious but rare (1/10000) disease. This test (t) is correct to 99% (1% false positive & 1% false negative results).

What does this mean for you?

$$P(d | t) = \frac{P(t | d)P(d)}{P(t)} = \frac{P(t | d)P(d)}{P(t | d)P(d) + P(t | \neg d)P(\neg d)}$$

$$P(d) = 0.0001 \quad P(t | d) = 0.99 \quad P(t | \neg d) = 0.01$$

$$\begin{aligned} P(d | t) &= \frac{0.99 \times 0.0001}{0.99 \times 0.0001 + 0.01 \times 0.9999} = \frac{0.000099}{0.000099 + 0.009999} \\ &= \frac{0.000099}{0.010088} \approx 0.01 \end{aligned}$$

Your doctor tells you that you have tested positive for a serious but rare (1/10000) disease. This test (t) is correct to 99% (1% false positive & 1% false negative results).

What does this mean for you?

$$P(d | t) = \frac{P(t | d)P(d)}{P(t)} = \frac{P(t | d)P(d)}{P(t | d)P(d) + P(t | \neg d)P(\neg d)}$$

$$P(d) = 0.0001 \quad P(t | d) = 0.99 \quad P(t | \neg d) = 0.01$$

$$\begin{aligned} P(d | t) &= \frac{0.99 \times 0.0001}{0.99 \times 0.0001 + 0.01 \times 0.9999} = \frac{0.000099}{0.000099 + 0.009999} \\ &= \frac{0.000099}{0.010088} \approx 0.01 \end{aligned}$$

11.2.5 Multiple Evidence

A probe by the dentist catches (Catch = true) in the aching tooth (Toothache = true) of a patient. We already know that $P(\text{cavity} \mid \text{toothache}) = 0.8$. Furthermore, using Bayes' rule, we can calculate: **$P(\text{cavity} \mid \text{catch}) = 0.95$**

$$\begin{aligned} P(\text{cav} | \text{tooth} \wedge \text{catch}) &= \frac{P(\text{tooth} \wedge \text{catch} | \text{cav}) \times P(\text{cav})}{P(\text{tooth} \wedge \text{catch})} \\ &= \alpha P(\text{tooth} \wedge \text{catch} | \text{cav}) \times P(\text{cav}) \end{aligned}$$

The dentist needs $P(\text{tooth} \wedge \text{catch} | \text{cav})$, i.e., diagnostic knowledge of all combinations of symptoms in the general case. They are conditionally independent given that we know whether the tooth has a cavity:

$$P(\text{tooth} | \text{catch}, \text{cav}) = P(\text{tooth} | \text{cav})$$

If one already knows that there is a cavity, then the additional knowledge of the probe catching does not change the probability.

$$\begin{aligned} P(\text{tooth} \wedge \text{catch} | \text{cav}) &= P(\text{tooth} | \text{catch}, \text{cav})P(\text{catch} | \text{cav}) = \\ &P(\text{tooth} | \text{cav})P(\text{catch} | \text{cav}) \end{aligned}$$

The general definition of conditional independence of two variables X and Y given a third variable Z (a common cause) is:

$$P(X, Y|Z) = P(X|Z)P(Y|Z)$$

11.2.6 Recursive Bayesian Updating

Assuming conditional independence, multiple evidence can be reduced to prior probabilities and conditional probabilities. The general combination rule, if Z1 and Z2 are independent given X is $P(X|Z1, Z2) = \alpha P(X)P(Z1|X)P(Z2|X)$ where α is the normalization constant.

Generalization: $P(X|Z_1, \dots, Z_n) = \alpha P(X) / \prod P(Z_i|X)$

11.3 Bayesian Network

Example of Burglary/Earthquake.

- Domain knowledge/ assumptions:

Events *Burglary* and *Earthquake* are independent. *Alarm* might be activated by burglary or earthquake. *John* calls if and only if he heard the alarm. His call probability is not influenced by the fact, that there is an earthquake at the same time. Same for *Mary*.

The *random variables* are the *nodes*. Directed edges between nodes represent **direct influence**. A **table of conditional probabilities (CPT)** is associated with every node, in which the effect of the parent nodes is quantified. The graph is acyclic (a **DAG**).

$$P(\text{maryCalls}|\text{alarm}, \text{burglary}) =$$

$$P(\text{maryCalls}|\text{alarm}) \text{ and } P(\text{maryCalls}|\text{alarm}, \text{burglary}, \text{johnCalls}, \text{earthquake})$$

$$= P(\text{maryCalls}|\text{alarm})$$

- > Bayesian Networks can be considered as sets of (conditional) independence assumptions.

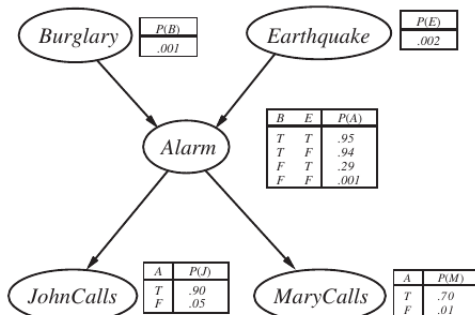
- General formula:

$$P(x_1, \dots, x_n) = P(x_n|x_{n-1}, \dots, x_1) \cdot \dots \cdot P(x_2|x_1)P(x_1) =$$

$$\prod_{i=1}^n P(x_i|x_{i-1}, \dots, x_1)$$

This is equivalent to:

$$P(x_1, \dots, x_n) = \prod_{i=1}^n P(x_i|\text{parents}(x_i))$$



$$P(j, m, a, b, e) = P(j|m, a, b, e)P(m|a, b, e)P(a|b, e)P(b|e)P(e)$$

$$= P(j|a)P(m|a)P(a|b, e)P(b|e)P(e)$$

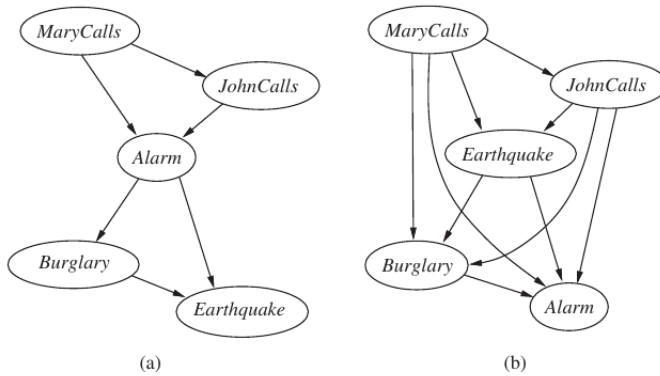
$$= 0.90.70.0010.9990.998 = 0.00062$$

In general, we need a table of size 2^n where n is the number of variables. The size depends on the application domain (local vs. global interaction) and the skill of the designer.

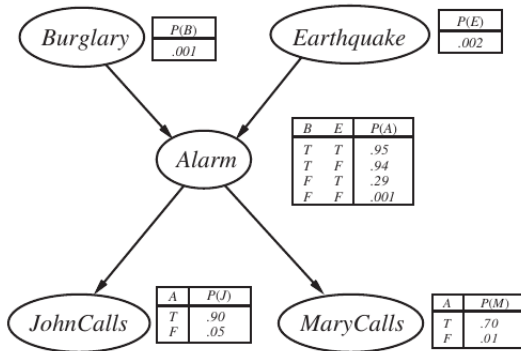
- Naive Networks method :

- Order all variables
- Take the first from those that remain
- Assign all direct influences from nodes already in the network to the new node (Edges + CPT).
- If there are still variables in the list, repeat from step 2.

left = M, J, A, B, E , right = M, J, E, B, A



- Inference in Bayesian Network



$P(\text{burglary} \mid \text{johncalls})$

$P(\text{burglary} \mid \text{johncalls}, \text{maryCalls})?$

A node is conditionally independent of its non-descendants given its parents.

- Example:

Consider $\mathbf{P}(\text{Burglary} \mid \text{JohnCalls} = \text{true}, \text{MaryCalls} = \text{true})$

The evidence variables are *JohnCalls* and *MaryCalls*.

The hidden variables are *Earthquake* and *Alarm*.

We have: $\mathbf{P}(B \mid j, m) = \alpha \mathbf{P}(B, j, m)$

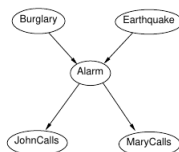
$$= \alpha \sum_e \sum_a \mathbf{P}(B, j, m, e, a)$$

If we consider the independence of variables, we obtain for $B = \text{true}$

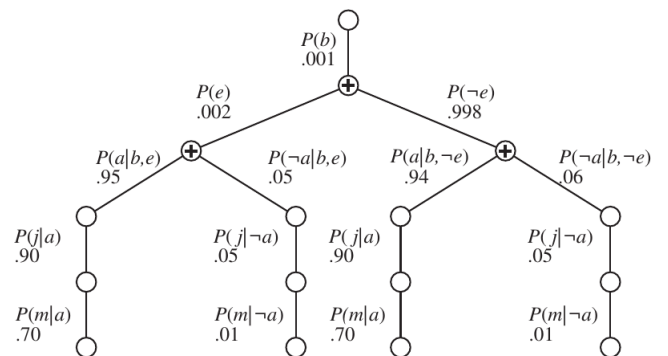
$$P(b \mid j, m) = \alpha \sum_e \sum_a P(j \mid a) P(m \mid a) P(a \mid e, b) P(e) P(b)$$

Reorganization of the terms yields:

$$P(b \mid j, m) = \alpha P(b) \sum_e P(e) \sum_a P(a \mid e, b) P(j \mid a) P(m \mid a)$$



$$P(b \mid j, m) = \alpha P(b) \sum_e P(e) \sum_a P(a \mid e, b) P(j \mid a) P(m \mid a)$$



$$\mathbf{P}(B \mid j, m) = \alpha(0.0006, 0.0015) = (0.284, 0.716)$$

```

function ENUMERATION-ASK( $X, \mathbf{e}, bn$ ) returns a distribution over  $X$ 
  inputs:  $X$ , the query variable
            $\mathbf{e}$ , observed values for variables  $\mathbf{E}$ 
            $bn$ , a Bayes net with variables  $\{X\} \cup \mathbf{E} \cup \mathbf{Y}$  /*  $\mathbf{Y}$  = hidden variables */

   $Q(X) \leftarrow$  a distribution over  $X$ , initially empty
  for each value  $x_i$  of  $X$  do
     $Q(x_i) \leftarrow$  ENUMERATE-ALL( $bn.VARS, \mathbf{e}_{x_i}$ )
    where  $\mathbf{e}_{x_i}$  is  $\mathbf{e}$  extended with  $X = x_i$ 
  return NORMALIZE( $Q(X)$ )

```

```

function ENUMERATE-ALL( $vars, \mathbf{e}$ ) returns a real number
  if EMPTY?( $vars$ ) then return 1.0
   $Y \leftarrow$  FIRST( $vars$ )
  if  $Y$  has value  $y$  in  $\mathbf{e}$ 
    then return  $P(y | \text{parents}(Y)) \times$  ENUMERATE-ALL( $\text{REST}(vars), \mathbf{e}$ )
    else return  $\sum_y P(y | \text{parents}(Y)) \times$  ENUMERATE-ALL( $\text{REST}(vars), \mathbf{e}_y$ )
    where  $\mathbf{e}_y$  is  $\mathbf{e}$  extended with  $Y = y$ 

```

Depth-first algorithm, linear space complexity, $O(2^n)$ time complexity.

- Variable elimination:

The algorithm can be improved by eliminating repeating or unnecessary calculations. The key idea is to evaluate expressions from right to left (bottom-up) and to save results for later use.

Let us consider the query $P(\text{JohnCalls} \mid \text{Burglary} = \text{true})$.

The nested sum is

$$P(j, b) = \alpha P(b) \sum_e P(e) \sum_a P(a \mid b, e) P(j, a) \sum_m P(m \mid a)$$

general observation: variables, that are not query or evidence variables and not ancestor nodes of query or evidence variables can be removed. Variable elimination repeatedly removes these variables and this way speeds up computation.

within example: Alarm and Earthquake are ancestor nodes of query variable JohnCalls and cannot be removed. MaryCalls is neither a query nor an evidence variable and no ancestor node. Therefore it can be removed.

- Complexity of exact inference:

If the network is singly connected or a polytree, the time and space complexity of is linear.

For multiply connected networks inference in Bayesian Networks is NP-hard.

12 Acting under Uncertainty Maximizing Expected Utility

12.1 Introduction to Utility theory

The utility function rates states and thus formalizes the desirability of a state by the agent. A non-deterministic action A can lead to the outcome states $\text{Result}_i(A)$. How high is the probability that the outcome state $\text{Result}_i(A)$ is reached, if A is executed in the current state with evidence E ?

$$\beta P(\text{Result}_i(A) \mid \text{Do}(A), E)$$

- Expected utility:

$$EU(A \mid E) = \sum_i P(\text{Result}_i(A) \mid \text{Do}(A), E) U(\text{Result}_i(A))$$

The principle of **maximum expected utility (MEU)** says that a rational agent should choose an action that maximizes $EU(A \mid E)$.

- Problem: Utility function requires search or planning, because an agent needs to know the possible future states in order to assess the worth of the current state ("effect of the state on the future").

12.2 Choosing Individual Actions

Preferences:

$L_1 > L_2$ Agent prefers 1 over 2

$L_1 \sim L_2$ Agent is indifferent between 1 or 2

$L_1 \succeq L_2$ Agent prefers 1 or is indifferent.

Orderability

$$(A > B) \vee (B > A) \vee (A \sim B)$$

Transitivity

$$(A > B) \vee (B > C) \Rightarrow (A > C)$$

Continuity

$$A > B > C \Rightarrow \exists p[p, A; 1 - p, C] \sim B$$

Substitutability

$$A \sim B \rightarrow [p, A; 1 - p, C] \sim [p, B; 1 - p, C]$$

Monotonicity

$$A > B \Rightarrow (p > q \Leftrightarrow [p, A; 1 - p, B] > [q, A; 1 - q, B])$$

Decomposability

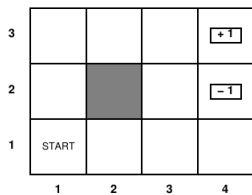
$$[p, A; 1 - p, [q, B; 1 - q, C]] \sim [p, A; (1 - p)q, B; (1 - p)(1 - q), C]$$

Expected Utility of a Lottery:

$$U[p_1, S_1, \dots, p_n, S_n] = \sum_i p_i U(S_i)$$

Since the outcome of a nondeterministic action is a lottery, an agent can act rationally only by following the Maximum Expected Utility (MEU) principle.

Given a utility scale between u_{min} and u_{max} we can assess the utility of any particular outcome S by asking the agent to choose between S and a standard lottery $[p, u_{max}; 1 - p, u_{min}]$. We adjust p until they are equally preferred.

12.3 Sequential Decision Problems

- **Deterministic version:**

All actions always lead to the next square in the selected direction, except that moving into a wall results in no change in position.

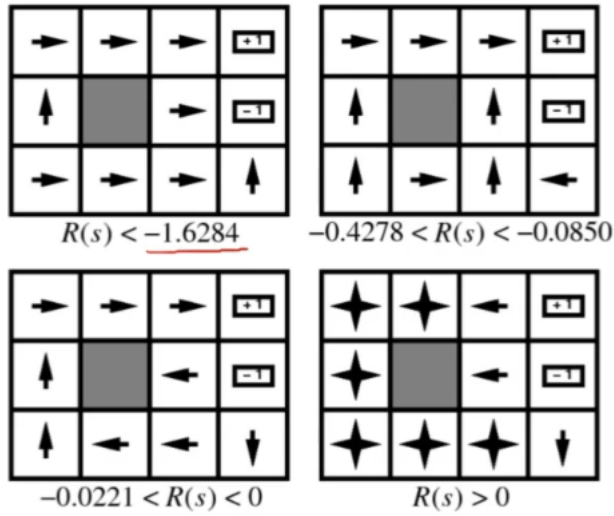
- **Stochastic version:**

Each action achieves the intended effect with probability 0.8, but the rest of the time, the agent moves at right angles to the intended direction.

12.4 Markov Decision Processes

Stochastic environment defined by: States, Actions, Transition model $P(s'|s, a)$ (probability), Reward function $R(s)$, Policy: mapping π that specifies for each state which action to take. We, obviously, want the optimal policy π^* that maximizes the future expected reward.

Given the optimal policy, it executes the action $\pi^*(s)$. In this way we obtain a simple reflex agent.



Performance is measured by the sum of rewards. To determine the optimal policy, first calculate the utility of each state, it also depends on the horizon, if it is finite or infinite. For finite horizons the optimal policy is called nonstationary, instead for infinite it is stationary.

- Stationary:

Two ways to reward:

- Additive rewards: $U_h([s_0, s_1, \dots, s_n]) = R(s_0) + R(s_1) \dots$
- Discounted rewards: $U_h([s_0, s_1, \dots, s_n]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) \dots$

$\gamma \in [0, 1]$ is the discount factor. With discount, the utility is always finite. The discount factor expresses that future rewards have less than current rewards.

12.5 Value Iteration

$U^\pi(s)$ is the utility of a state under policy π . s_t be the state of the agent after executing π for t steps. The utility of s under π is:

$$U^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t R(s_t) | \pi, s_0 = s\right] \quad (3)$$

$R(s)$ is the short-term reward for being in s and $U(s)$ is the long-term total reward from s onwards.

The utilities of the states in our 4×3 world with $\gamma = 1$ and $R(s) = -0.04$ for non-terminal states:

3	0.812	0.868	0.918	+ 1
2	0.762		0.660	- 1
1	0.705	0.655	0.611	0.388
	1	2	3	4

The agent simply chooses the action that maximizes the expected utility of the subsequent state:

$$\pi(s) = \operatorname{argmax}_a \sum_{s'} P(s'|s, a) U(s') \quad (4)$$

The utility of a state is the immediate reward for that state plus the expected discounted utility of the next state, - **Bellman-Equation**:

$$U(s) = R(s) + \gamma \max_a \sum_{s'} P(s'|s, a) U(s') \quad (5)$$

- **Example**:

In our 4×3 world the equation for the state (1,1) is

$$\begin{aligned}
 U(1, 1) &= -0.04 + \gamma \max \{ \begin{array}{l} 0.8U(1, 2) + 0.1U(2, 1) + 0.1U(1, 1), \\ 0.9U(1, 1) + 0.1U(1, 2), \\ 0.9U(1, 1) + 0.1U(2, 1), \\ 0.8U(2, 1) + 0.1U(1, 2) + 0.1U(1, 1) \end{array} \quad \begin{array}{l} (Up) \\ (Left) \\ (Down) \\ (Right) \end{array} \\
 &= -0.04 + \gamma \max \{ \begin{array}{l} 0.8 \cdot 0.762 + 0.1 \cdot 0.655 + 0.1 \cdot 0.705, \\ 0.9 \cdot 0.705 + 0.1 \cdot 0.762, \\ 0.9 \cdot 0.705 + 0.1 \cdot 0.655, \\ 0.8 \cdot 0.655 + 0.1 \cdot 0.762 + 0.1 \cdot 0.705 \end{array} \quad \begin{array}{l} (Up) \\ (Left) \\ (Down) \\ (Right) \end{array} \\
 &= -0.04 + 1.0 (0.6096 + 0.0655 + 0.0705), \quad (Up) = -0.04 + 0.7456 = 0.7056
 \end{aligned}$$

Up is the optimal action in (1,1).

3	0.812	0.868	0.918	+ 1
2	0.762		0.660	- 1
1	0.705	0.655	0.611	0.388
	1	2	3	4

- Value Iteration

An algorithm to calculate an optimal strategy. Basic Idea: Calculate the utility of each state. A sequence of actions generates a branch in the tree of possible states (histories). A utility function on histories U_h is separable iff there exists a function f such that:

$$U_h([s_0, s_1, \dots, s_n]) = f(s_0, U_h([s_1, \dots, s_n])) \quad (6)$$

The simplest form is an additive reward function R :

$$U_h([s_0, s_1, \dots, s_n]) = R(s_0) + U_h([s_1, \dots, s_n]) \quad (7)$$

Typical problems contain cycles, which means the length of the histories is potentially infinite. Solution: Use

$$U_{t+1}(s) = R(s) + \gamma \max_a \sum_{s'} P(s'|s, a) U_t(s') \quad (8)$$

The Bellman equation is the basis of value iteration, thanks to the max-operator the n equations for the n states are nonlinear. We can apply an iterative approach in which we replace the equality by an assignment:

$$U(s') \leftarrow R(s) + \gamma \max_a \sum_{s'} P(s'|s, a) U(s') \quad (9)$$

- Algorithm

function VALUE-ITERATION(mdp, ϵ) **returns** a utility function

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$,
rewards $R(s)$, discount γ

ϵ , the maximum error allowed in the utility of any state

local variables: U, U' , vectors of utilities for states in S , initially zero

δ , the maximum change in the utility of any state in an iteration

repeat

$U \leftarrow U'; \delta \leftarrow 0$

for each state s **in** S **do**

$U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

if $|U'[s] - U[s]| > \delta$ **then** $\delta \leftarrow |U'[s] - U[s]|$

until $\delta < \epsilon(1 - \gamma)/\gamma$

return U

- Convergence:

if $\|U_{t+1} - U_t\| < \epsilon(1 - \gamma)/\gamma$ then $\|U_{t+1} - U\| < \epsilon$

if $\|U_t - U\| < \epsilon$ then $\|U^{\pi_t} - U\| < 2\epsilon\gamma/(1 - \gamma)$

12.5.1 Policy Iteration

Policy iteration alternates the following two steps beginning with an initial policy π_0 :

- **Policy evaluation:** given π_t , calculate $U_t = U^{\pi_t}$ the utility if π_t was executed.

- **Policy improvement:** calculate a new maximum expected utility policy π_{t+1} according to:

$$\pi_{t+1}(s) = \operatorname{argmax}_a \sum_{s'} P(s'|s, a) U_t(s') \quad (10)$$

function POLICY-ITERATION(mdp) **returns** a policy

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$

local variables: U , a vector of utilities for states in S , initially zero

π , a policy vector indexed by state, initially random

repeat

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$

$unchanged? \leftarrow \text{true}$

for each state s **in** S **do**

if $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ **then do**

$\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

$unchanged? \leftarrow \text{false}$

until $unchanged?$

return π

13 Machine Learning

An agent learns when it improves its performance w.r.t. a specific task with experience.

13.1 The learning agent

Performance element: Processes percepts and chooses actions.

Learning element: Carries out improvements.

Critic: Evaluation of the agent's behavior based on a given external behavioral measure.

Problem generator: Suggests explorative actions that lead the agent to new experiences.

13.2 Types of learning

Supervised learning: Involves learning a function from examples of its inputs and outputs.

An example is a pair $(x, f(x))$. The complete set of examples is called the **training set**. Pure inductive inference: for a collection of examples for f , return a function h (hypothesis) that approximates f . A good hypothesis should predict unseen examples correctly. A **hypothesis is consistent** with the data set if it agrees with all the data. **Ockham's razor:** prefer the simplest hypothesis consistent with the data.

Unsupervised learning: The agent has to learn patterns in the input when no specific output values are given.

Reinforcement learning: The most general form of learning in which the agent is not told what to do by a teacher. Rather it must learn from a reinforcement or reward. It typically involves learning how the environment works.

13.3 Decision trees

Input: set of attributes.

Output: a decision.

Continuous (regression) or discrete (classification), binary or not.

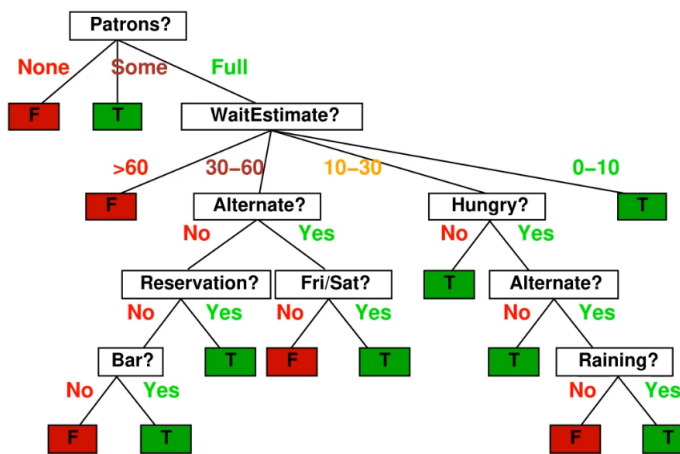
13.3.1 Boolean Decision Trees

Input: set of vectors of input attributes.

Output: Yes/No decision based on a goal predicate.

Properties: an internal node represents a test of a property, branches are labeled with the possible values, each leaf specifies the returning value.

Example:



Each decision tree hypothesis can be seen as an assertion of the form:

$$\forall s \text{WillWait}(s) \leftrightarrow (P_1(s) \vee P_2(s) \vee \dots P_n(s)) \quad (11)$$

$P_i(s)$ are the tests along the path. **Limitation:** trees always involve just one variable.

Compact representations:

We can construct a decision tree by translating every row of a truth table to a path in the tree. This can lead to exponential trees (parity, majority).

- **Training Set example:**

Classification of an example = Value of the goal predicate

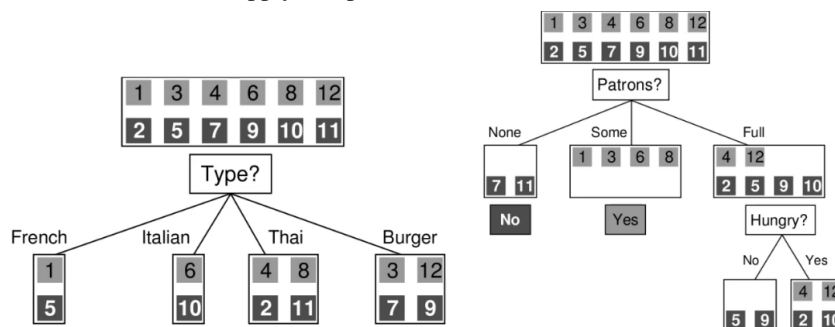
$T \rightarrow$ positive example

$F \rightarrow$ negative example

Example	Attributes										Target
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	WillWait
X_1	T	F	F	T	Some	\$\$\$	F	T	French	0-10	T
X_2	T	F	F	T	Full	\$	F	F	Thai	30-60	F
X_3	F	T	F	F	Some	\$	F	F	Burger	0-10	T
X_4	T	F	T	T	Full	\$	F	F	Thai	10-30	T
X_5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X_6	F	T	F	T	Some	\$\$	T	T	Italian	0-10	T
X_7	F	T	F	F	None	\$	T	F	Burger	0-10	F
X_8	F	F	F	T	Some	\$\$	T	T	Thai	0-10	T
X_9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X_{10}	T	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F
X_{11}	F	F	F	F	None	\$	F	F	Thai	0-10	F
X_{12}	T	T	T	T	Full	\$	F	F	Burger	30-60	T

- Creation possibilities:

- **Naive:** simply construct with one path to a leaf for each example. Test all attributes. Even if it will classify the examples, It will not say anything about other classes. It just memorizes the observation and doesn't generalize.
- **Smallest solution:** Applying Ockham's razor we find the smallest decision tree. unfortunately, the smallest definition is intractable.
- **Divide and Conquer:** Choose an attribute, Split the training set such that each corresponds to a particular value of that attribute, recursive apply this process to the smaller sets.



Type is a poor attribute since it leaves us with state full of true and false.
Patrons is a better choice.

In each recursive step there are 4 cases:

- **Positive and negative:** choose a new attribute.
- **Only positive:** finished.
- **No examples:** Answer the majority
- **No attributes left:** Too noisy.

```

function DTL(examples, attributes, default) returns a decision tree
  if examples is empty then return default
  else if all examples have the same classification then return the classification
  else if attributes is empty then return MODE(examples)
  else
    best ← CHOOSE-ATTRIBUTE(attributes, examples)
    tree ← a new decision tree with root test best
    for each value  $v_i$  of best do
       $examples_i \leftarrow \{\text{elements of } examples \text{ with } best = v_i\}$ 
      subtree ← DTL( $examples_i$ , attributes – best, MODE(examples))
      add a branch to tree with label  $v_i$  and subtree subtree
  return tree

```

Original tree:

