

# Image Processing and Computer Graphics

## Image Processing

---


### Class 6

### Matching and local descriptors

- Key problem in computer vision appearing in:
  - Motion estimation
  - Camera calibration
  - Stereo
  - Image retrieval
  - Object recognition




Object recognition: training image on the left, test image on the right. Matching here is quite hard.



Stereo pair: point matching needed to compute depth


1. Sparse matching
  - What are good regions to match?
  
2. Manual descriptor design
  - What is required for a good descriptor?
  
3. Feature learning
  - How can we optimize descriptors for a particular task?

- Straightforward way to match points in images:
  - Regard the image patch around each point in image 1
  - Compare it to the image patches around all points in image 2



- Computationally expensive  
 $O(kN^2)$ ,  $k$ : size of patch,  $N$ : size of image (in pixels)
- Normal image patches are not **invariant** to typical appearance changes  
→ normalization or invariant descriptors

- Often we need only a limited number of matches
- Idea: Do not match all points in the images, but only promising subsets  
→ significantly reduced complexity
- Requirements for good interest points:
  1. Points must come with enough information for unique matching



2. Subset in image 2 must contain matches from subset in image 1

- Choose points with high information content and clear localization  
→ typically corner points
- Corner detection with the structure tensor:  
(Förstner-Gülch 1987, Harris-Stevens 1988)

$$J_\rho = K_\rho * (\nabla I \nabla I^\top) = \begin{pmatrix} K_\rho * I_x^2 & K_\rho * I_x I_y \\ K_\rho * I_x I_y & K_\rho * I_y^2 \end{pmatrix}$$


- Measure of cornerness (unintuitive, but fast to compute):

$$c = \det J_\rho - \alpha \operatorname{tr} J_\rho$$

↑  
= gradient magnitude

- Eigenvalue decomposition of the structure tensor:

$$J_\rho = T \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix} T^\top$$
$$c = \lambda_2$$



- Interpretation:
  - Smoothing of  $J$  integrates gradients from the neighborhood
  - Eigenvectors in  $T$  yield the dominant orientation in this neighborhood and the perpendicular orientation
  - Eigenvalues yield the structure magnitude in these directions
  - A large second eigenvalue indicates strong structures in multiple directions  $\rightarrow$  corners

- Corners: local maxima of the second eigenvalue





- Problem: Detected corners depend on the image scale




- Solution: Create and compare descriptors at multiple scales

- Considers local maxima of the Laplacian in scale space:



$$x^*, y^*, \sigma^* = \operatorname{argmax}_{x, y, \sigma} \left( \sigma^2 \cdot |\partial_{xx}(K_\sigma * I(x, y)) + \partial_{yy}(K_\sigma * I(x, y))| \right)$$

- Laplacian is rather a blob detector than a corner detector
- Detects regions with multiple sizes  
→ good to normalize out scale changes



Authors: Krystian Mikolajczyk and Cordelia Schmid

- Maximally stable extremal regions  
(Matas et al. 2002)
  - Regions encircled by large gradients
  - Obtained by low-level segmentation algorithm




Maximally stable extremal regions and fitted ellipses. Author: Andrea Vedaldi

- Apart from scale also yields elongation of fitted ellipses  
→ allows for affine invariance

- Affine transformation (6 degrees of freedom in 2D):


$$f(x) = Ax + t$$

- Maps a circle to an ellipse (or vice-versa):



- Approximation of a projective transformation
- Parameters can be estimated from a region detector

- Significantly reduced complexity: with 100 detected points in both images, one has to compare only 10000 patches instead of 96 billion(!) in 640x480 images
- Allows us to efficiently normalize out some variations (scale, affine transformations)
- Non-dense displacement fields (important matches might be missed)
- Corresponding patches can be slightly shifted
- Other variations (than affine transformations) not covered



















- **Local descriptors:** vectors that contain information about the local neighborhood of a point or region of interest
- Goal: design local descriptors that are **invariant** under the mentioned selected transformations
- Be Careful: If a descriptor is invariant to all sorts of transformations, it may not be descriptive anymore

- Alternative to a normalized neighborhood:  
derive invariant features within the fixed block
- Gray value histogram:
  - Rotational invariance (good)
  - Invariant to blurring (good)
  - Sensitive to lighting changes (bad)
  - Significant loss of information (very bad)
- Histogram of the gradient direction (**orientation histograms**)
  - Invariant to (additive) lighting changes
  - Building block of many successful descriptors






- Popular local descriptor (several variants exist)
- Based on local assembly of orientation histograms and adaptive local neighborhoods





Author: David Lowe

- Extract SIFT feature points
    - Strongest responses of Laplacian in scale space  
→ position and scale
    - Fit quadratic function to obtain subpixel accuracy
  - Create orientation histogram at selected scale
    - Peak of smoothed histogram estimates orientation
    - In case of two peaks, create two feature points
- Estimation of position, scale, and orientation
- Affine invariance can be provided with MSER
  - In object recognition: dense sampling of such points at all positions and all scales, no rotation invariance




Author: David Lowe

1. Compute gradient orientation and magnitude at each pixel
2. Compute orientation indicator at each pixel
  - Create  $N \times M \times 8$  array and initialize with zero
  - Quantize the orientation at each pixel (here 8 bins) and add the respective magnitude to the respective entry in the array
3. Local integration  $\rightarrow$  orientation histogram  
Smooth array with a Gaussian kernel
4. Smooth in orientation direction  
(among neighboring channels)




5. Sample feature vectors from the histogram image
  - Original SIFT:  
4 pixel spacing, 4x4 histogram array  
→ 128-D vector
6. Normalize the feature vector to unit length



Author: Utkarsh Sinha


<http://aishack.in/tutorials/sift-scale-invariant-feature-transform-features/>

- Instead of manual descriptor design, let the computer find the optimal descriptor for a defined task and training set
- Example task: object classification  
→ training set consists of images and their class labels



- Shallow modeling of the function  $f(I)$  is not efficient to cover all the variation that appears in an object class  
→ hierarchy of functions, “deep” representation





- Classification networks are trained on large datasets with class labels (e.g. ImageNet with 1M images)  
→ network learns a representation that is good for object classification
- Intermediate layer outputs turn out to be also good generic descriptors (still a bit mystic what is represented)


- Train CNN to discriminate **surrogate classes** (Dosovitskiy et al. 2015)




Seed patch and transformed versions of it make up a surrogate class

- Applied transformations:  
translation, rotation, scaling, color, contrast, brightness, blur
- Transformations define invariance properties of the features to be learned by the network → early version of today's contrastive learning


Exemplar-CNN-blur vs SIFT



Exemplar-CNN-blur vs AlexNet



- Trained directly on matching and non-matching patches



Learns the  
metric

Learns the  
descriptor

Example from Zagoruyko&Komodakis 2015


- Issue class imbalance: far more non-matching patches than matching ones

- The principle of surrogates can be generalized to a contrastive loss to learn a feature embedding.
- Main idea: for **a sample** define another **positive sample** that should be close in embedding space, and multiple **negative samples** that should be far

- Contrastive loss:

$$\ell_{i,j} = -\log \frac{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)}{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)}$$

where  $\tau$  is a scaling parameter that can be reduced over time



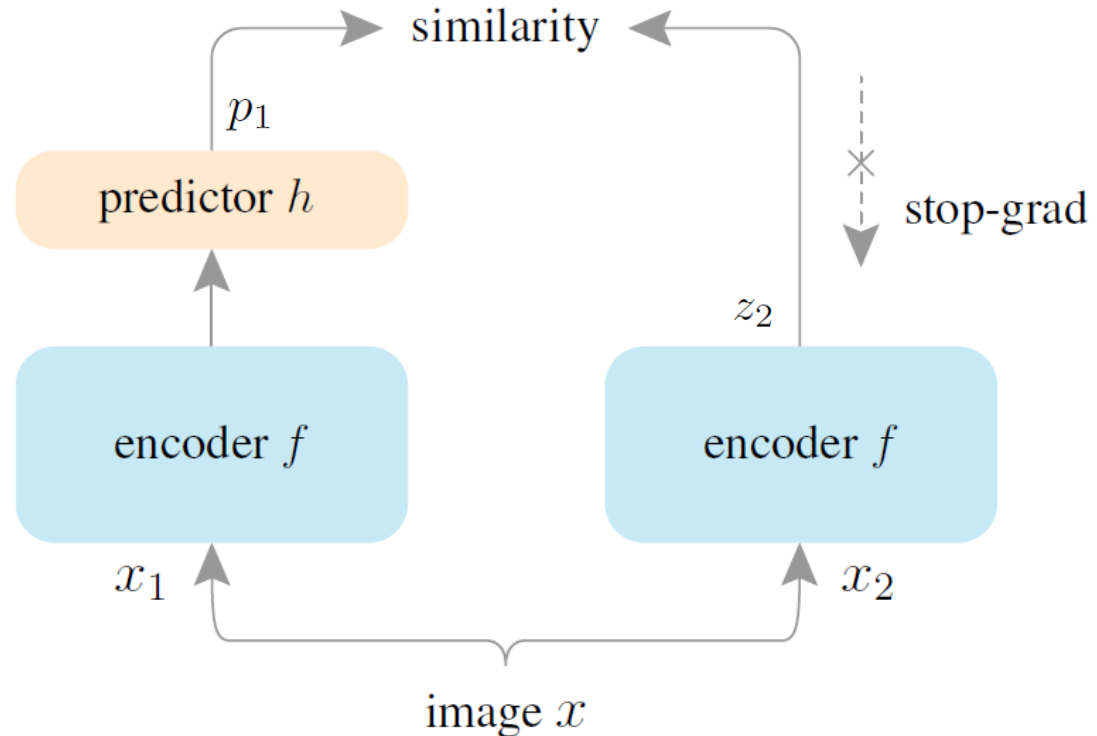
- Positives and negatives can often be defined without supervision  
→ self-supervised learning
- Are negative samples needed?

- Put two matching samples (positives) as  $x_1$  and  $x_2$

- Train the encoder with a similarity loss

$$\mathcal{D}(p_1, z_2) = -\frac{p_1}{\|p_1\|_2} \cdot \frac{z_2}{\|z_2\|_2}$$

- Predictor is an additional 2-layer network



- Predictor is flipped alternately between  $x_1$  and  $x_2$

$$\mathcal{L} = \frac{1}{2}\mathcal{D}(p_1, z_2) + \frac{1}{2}\mathcal{D}(p_2, z_1)$$

- This strategy (predictor and stop-gradient switching permanently) avoids collapse to a trivial representation

- Interest points are distinctive points in an image with a significant information content in their neighborhood
- Interest point detection can help establish invariance to certain image transformations.
- Local descriptors describe a local area in the image for the purpose of matching.
- The SIFT descriptor is based on a grid of orientation histograms
- Intermediate layers of convolutional networks yield good descriptors
- Unsupervised learning strategies can learn targeted invariance

- D. G. Lowe: Distinctive image features from scale-invariant keypoints, *International Journal of Computer Vision* 60(2):91-110, 2004.
- J. Matas, O. Chum, M. Urban, T. Pajdla: Robust wide baseline stereo from maximally stable extremal regions, *Proc. British Machine Vision Conference*, 2002.
- A. Dosovitskiy, P. Fischer, T. Springenberg, M. Riedmiller, T. Brox: Discriminative unsupervised feature learning with convolutional neural networks, *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2016.
- S. Zagoruyko, N. Komodakis: Learning to Compare Image Patches via Convolutional Neural Networks, *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2015.
- X. Chen, K. He: Exploring Simple Siamese Representation Learning, *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021.



- Implement the corner detector based on the second eigenvalue of the structure tensor
  - For computing derivatives and for smoothing images you can make use of the predefined filter masks as well as the convolution operations in `CFilter.h`.
  - The structure tensor of a color image is the sum of tensors over all channels
  - See an online math lecture if you do not remember how to compute the eigenvalues of a matrix <http://www.khanacademy.org/>
- Apply the corner detector to the images in `ImageProcessing08Ex03.zip` and play with the parameters
- Implement the dense SIFT descriptor (without the detector). Use a 4 pixel spacing and a 3x3 grid of histograms. You can ignore scale and rotation invariance and even skip normalization for this exercise.
- Run your corner detector on `tennis500.ppm` and manually select among the interest points the 10 visually most interesting ones. Extract SIFT descriptors for these points.
- Compute SIFT descriptors for all points in `tennis505.ppm`. For each descriptor in `tennis500.ppm` find the best match in `tennis505.ppm` and visualize the correspondences in your result image.
- Play with the amount of smoothing, the spacing, and the number of histograms per descriptor.