

Sapienza University of Rome

Master in Artificial Intelligence and Robotics  
Master in Engineering in Computer Science

## Machine Learning

A.Y. 2020/2021

Prof. L. Iocchi, F. Patrizi

## 8. Linear models for regression

L. Iocchi, F. Patrizi

# Overview

- Linear models for regression
- Maximum likelihood and Least squares
- Sequential learning
- Regularization

## References

C. Bishop. Pattern Recognition and Machine Learning. Sect. 3.1

## Linear Models for Regression

Learning a function  $f : X \rightarrow Y$ , with

- $X \subseteq \mathbb{R}^d$
- $Y = \mathbb{R}$

from data set  $D = \{(\mathbf{x}_n, t_n)_{n=1}^N\}$

# Linear Models for Regression

Define a model  $y(\mathbf{x}; \mathbf{w})$  with parameters  $\mathbf{w}$  to approximate the target function  $f$ . *f. → linear*

Linear model for linear function

*linear comb.*

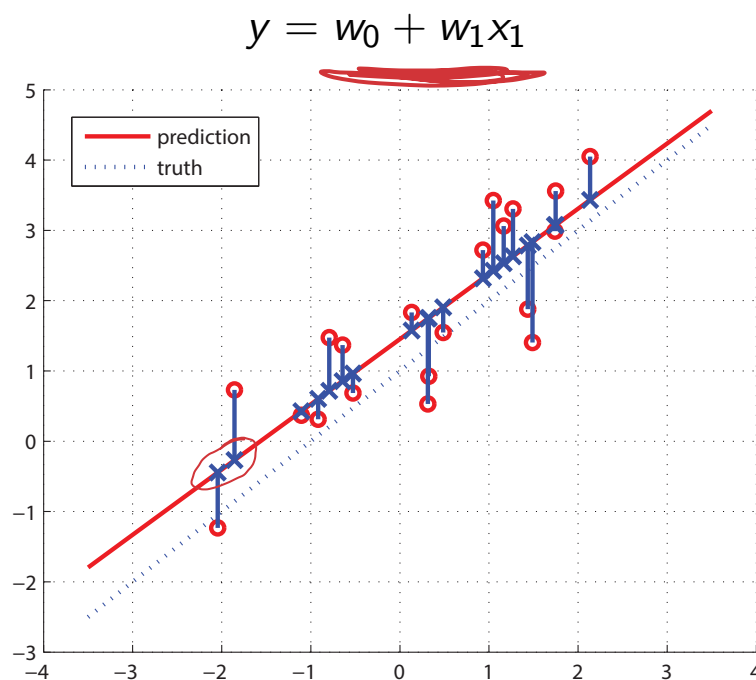
$$y(\mathbf{x}; \mathbf{w}) = w_0 + w_1 x_1 + \dots + w_d x_d = \mathbf{w}^T \mathbf{x}$$

*no ~ because includes  $w_0$  & 1.*

with  $\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix}$  and  $\mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix}$

*Why 1? from ab 104?*

## Example: 2D line fitting



# Linear Models for Regression

## Linear Basis Function Models

Using nonlinear functions of input variables:

$$y(\mathbf{x}; \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}),$$

*the same of  
classification:  
nonlinear inputs  
data.*

*no need to  
be invertible*

with  $\mathbf{w} = \begin{bmatrix} w_0 \\ \vdots \\ w_{M-1} \end{bmatrix}$   $\phi(\mathbf{x}) = \begin{bmatrix} \phi_0(\mathbf{x}) \\ \vdots \\ \phi_{M-1}(\mathbf{x}) \end{bmatrix}$ , and  $\phi_0(\mathbf{x}) = 1$ .

*Basis function*

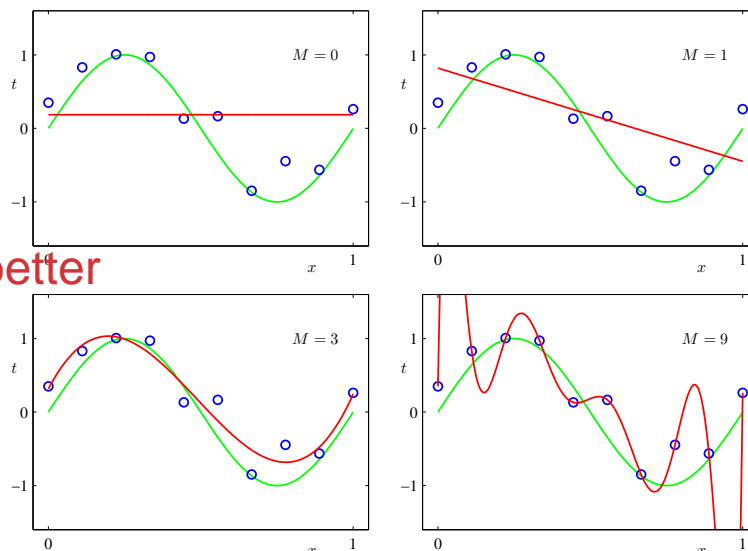
- Still linear in the parameters  $\mathbf{w}$ !

*Screen 10:26  
10:28  
10:31  
w: ) )*

## Example: Polynomial curve fitting

*problem is to find w.*

$$y = w_0 + w_1 x + w_2 x^2 + \dots + w_M x^M = \sum_{j=0}^M w_j x^j$$




Right degree fit better


Too high degree for polynomial you will overfit, 0 training error, but very high test error.

# Linear Regression Basis Functions


## Examples of basis functions



Polynomial



Radial



Sigmoid / Tanh

# Linear Regression - Algorithms

## Maximum likelihood and least squares

Target value  $t$  is given by  $y(\mathbf{x}; \mathbf{w})$  affected by additive noise  $\epsilon$

$$t = y(\mathbf{x}; \mathbf{w}) + \epsilon$$

*noise (error)*

Assume Gaussian noise  $P(\epsilon|\beta) = \mathcal{N}(\epsilon|0, \beta^{-1})$ , with precision (inverse variance)  $\beta$ .

*centered at 0*

We have:

$$P(t|\mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t|y(\mathbf{x}; \mathbf{w}), \beta^{-1})$$

# Linear Regression - Algorithms

Assume observations independent and identically distributed (i.i.d.)

We seek the maximum of the likelihood function:

$$P(\{t_1, \dots, t_N\} | \mathbf{x}_1, \dots, \mathbf{x}_N, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}).$$

*hint is that max this P.*

or equivalently:

$$\ln P(\{t_1, \dots, t_N\} | \mathbf{x}_1, \dots, \mathbf{x}_N, \mathbf{w}, \beta) = \sum_{n=1}^N \ln \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$$

$$= -\beta \frac{1}{2} \underbrace{\sum_{n=1}^N [t_n - \mathbf{w}^T \phi(\mathbf{x}_n)]^2}_{E_D(\mathbf{w})} - \frac{N}{2} \ln(2\pi\beta^{-1}).$$

*Squared error.*

Maximizing the probability means minimizing error

# Linear Regression - Algorithms

Maximum likelihood (zero-mean Gaussian noise assumption)

$$\operatorname{argmax} P(\{t_1, \dots, t_N\} | \mathbf{x}_1, \dots, \mathbf{x}_N, \mathbf{w}, \beta)$$

corresponds to least square error minimization

$$\operatorname{argmin} E_D(\mathbf{w}) = \operatorname{argmin} \frac{1}{2} \sum_{n=1}^N [t_n - \mathbf{w}^T \phi(\mathbf{x}_n)]^2$$

# Linear Regression - Algorithms

Note:

$$E_D(\mathbf{w}) = \frac{1}{2}(\mathbf{t} - \Phi\mathbf{w})^T(\mathbf{t} - \Phi\mathbf{w}),$$

with  $\mathbf{t} = \begin{bmatrix} t_1 \\ \vdots \\ t_N \end{bmatrix}$  and  $\Phi = \begin{bmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \cdots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \cdots & \phi_{M-1}(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \cdots & \phi_{M-1}(\mathbf{x}_N) \end{bmatrix}.$

Optimality condition:

$$\nabla E_D = 0 \iff \Phi^T \Phi \mathbf{w} = \Phi^T \mathbf{t}.$$

Why gradient instead of the solution using phi matrix? If the input is too large this may be not practical

Hence:

$$\mathbf{w}_{ML} = \underbrace{(\Phi^T \Phi)^{-1} \Phi^T}_{\Phi^\dagger: \text{pseudo-inverse}} \mathbf{t}.$$

# Linear Regression - Algorithms

## Sequential Learning

Stochastic gradient descent algorithm:

$$\hat{\mathbf{w}} \leftarrow \hat{\mathbf{w}} - \eta \nabla E_n$$

$\eta$ : learning rate parameter

Therefore:

$$\hat{\mathbf{w}} \leftarrow \hat{\mathbf{w}} + \eta [t_n - \hat{\mathbf{w}}^T \phi(\mathbf{x}_n)] \phi(\mathbf{x}_n)$$

Algorithm converges for suitable small values of  $\eta$ .

# Linear Regression - Regularization

Regularization is a technique to control over-fitting.

$$\operatorname{argmin} \overbrace{E_D(\mathbf{w})}^{\text{error}} + \underbrace{\lambda E_W(\mathbf{w})}_{\text{artificial error}}$$

with  $\lambda > 0$  being the regularization factor

A common choice:

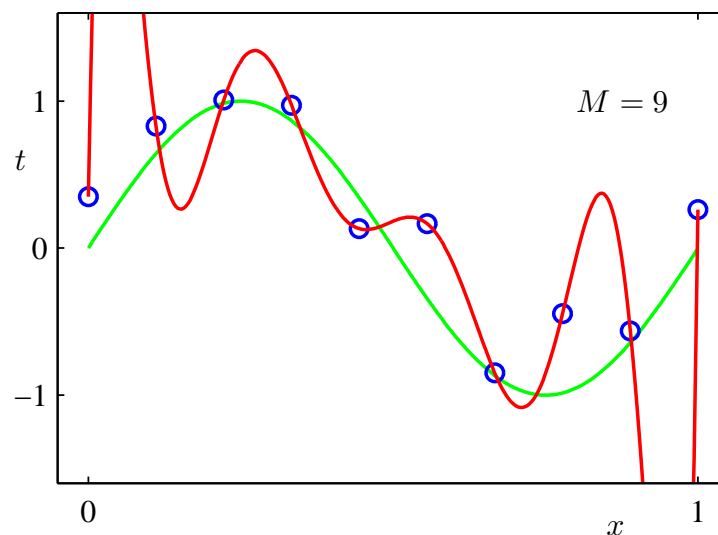
$$E_W(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}.$$

Other choices:

$$E_W(\mathbf{w}) = \sum_{j=0}^{M-1} |w_j|^q.$$

# Linear Regression - Regularization

$$\operatorname{argmin} E_D(\mathbf{w})$$


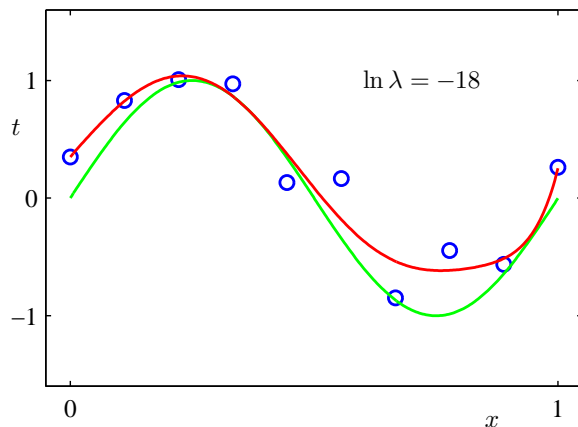




# Linear Regression - Regularization

$$\operatorname{argmin} E_D(\mathbf{w}) + \lambda \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

*Handwritten:  $\lambda E_w$  above the regularization term*



## Linear Regression - Multiple outputs

$$\mathbf{y}(\mathbf{x}; \mathbf{W}) = \mathbf{W}^T \phi(\mathbf{x})$$

Target variable is given by:

$$\mathbf{T} = \mathbf{y}(\mathbf{x}; \mathbf{W}) + \epsilon$$

with  $P(\epsilon|\beta) = \mathcal{N}(\epsilon|0, \beta^{-1}\mathbf{I})$ .

Similarly with before we obtain:

$$\mathbf{W}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{T}.$$