

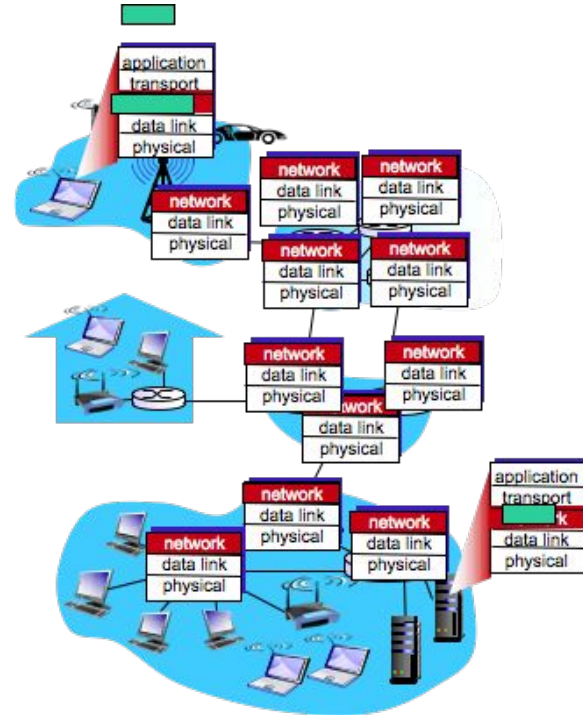
# Software Defined Networking

Network Infrastructures A.A 2020/21

- Control Plane and Data Plane
- Generalized Packet Forwarding
- Flow Table Entry
- Software Defined Networking
- OpenFlow Protocol
- SDN control/data plane interaction

# Network Layer

- **transport segment** from sending to receiving host
- on sending side **encapsulates segments into datagrams**
- on receiving side, **delivers segments to transport layer**
- network layer protocols in **every host, router**
- **router examines header** fields in all IP datagrams passing through it




# Two key network-layer functions

- network-layer functions:
  - **forwarding:** move packets from router's input to appropriate router output
  - **routing:** determine route taken by packets from source to destination
    - routing algorithms
- analogy: taking a trip
  - **forwarding:** process of getting through single interchange
  - **routing:** process of planning trip from source to destination

# Network layer: data plane, control plane

- **Data plane**

- **local**, per-router function
- determines **how datagram** arriving on router input port is **forwarded to router output port**
- forwarding function




- **Control plane**

- **network-wide logic**
- determines **how datagram is routed** among routers along end-end path from source host to destination host
- **two** control-plane **approaches**:
  - **traditional routing algorithms**: implemented in routers
  - **software-defined networking (SDN)**: implemented in (remote) servers


# Per-router control plane

- Individual routing algorithm components in each and every router interact in the control plane



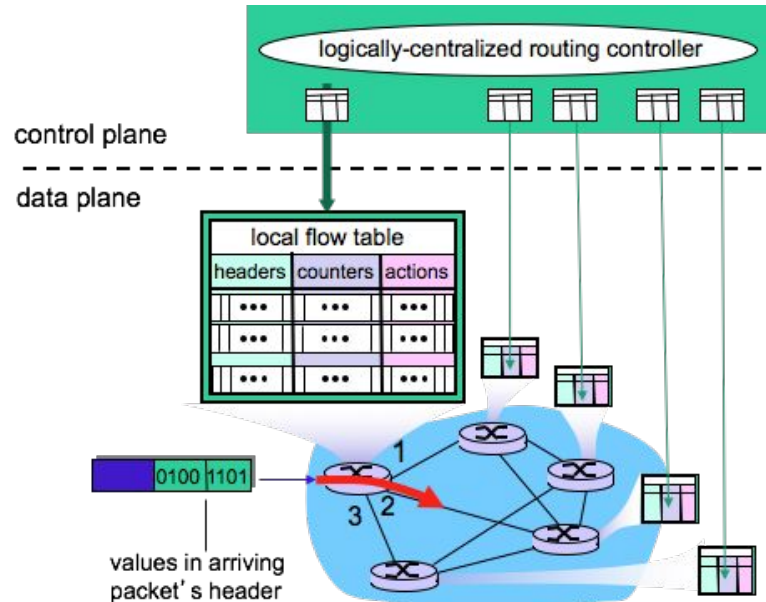
# Logically centralized control plane

- A distinct (typically remote) controller interacts with local control agents (CAs)



# Generalized Forwarding and SDN

- Each router contains a flow table that is computed and distributed by a logically centralized routing controller





# OpenFlow data plane abstraction

- **flow:** defined by header fields
- **generalized forwarding:** simple packet-handling rules
  - **Pattern:** match values in packet header fields
  - **Actions:** for matched packet: drop, forward, modify, matched packet or send matched packet to controller
  - **Priority:** disambiguate overlapping patterns
  - **Counters:** #bytes and #packets



Flow table in a router (computed and distributed by controller) define router's match+action rules

# OpenFlow data plane abstraction

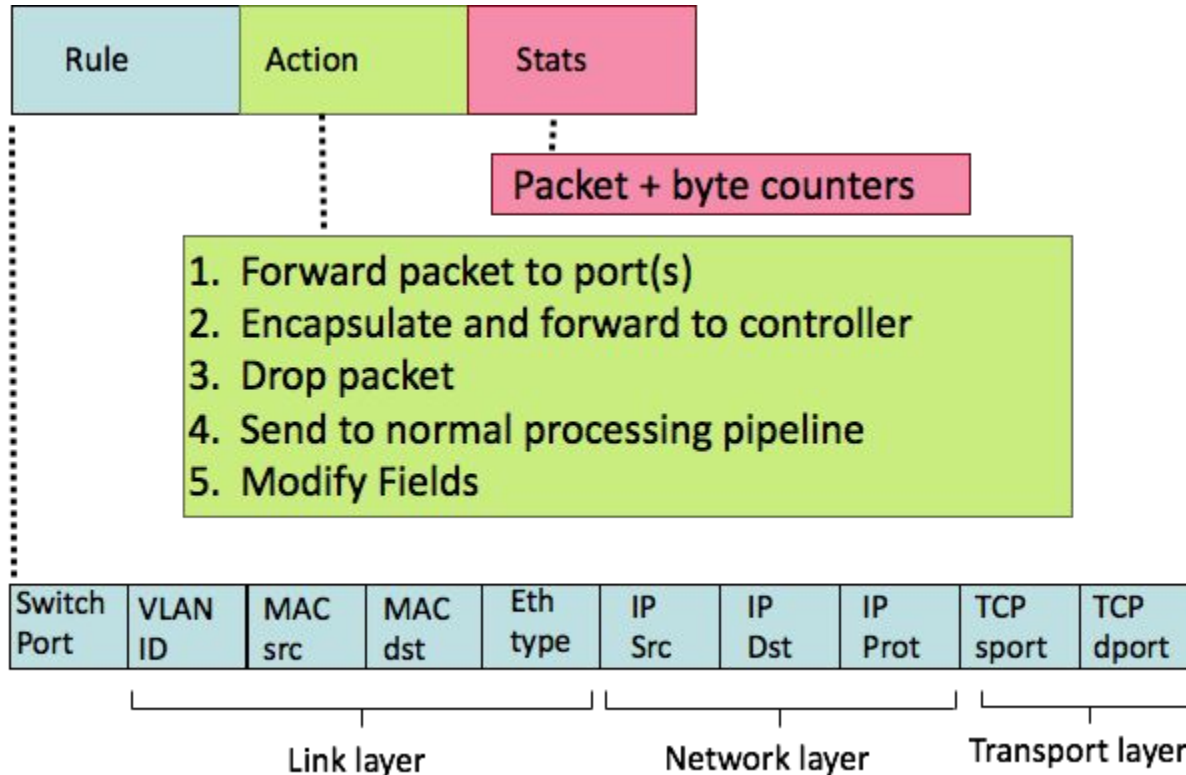
- **flow:** defined by header fields
- **generalized forwarding:** simple packet-handling rules
  - **Pattern:** match values in packet header fields
  - **Actions:** for matched packet: drop, forward, modify, matched packet or send matched packet to controller
  - **Priority:** disambiguate overlapping patterns
  - **Counters:** #bytes and #packets



\* : wildcard

1. src=1.2.\*.\*, dest=3.4.5.\* → drop
2. src = \*.\*.\*.\*, dest=3.4.\*.\* → forward(2)
3. src=10.1.2.3, dest=\*.\*.\*.\* → send to controller

# OpenFlow: Flow Table Entries



# Examples

## Destination-based forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	51.6.0.8	*	*	*	port6

*IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6*

## Firewall:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Forward
*	*	*	*	*	*	*	*	*	22	drop

*do not forward (block) all datagrams destined to TCP port 22*

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Forward
*	*	*	*	*	128.119.1.1	*	*	*	*	drop

*do not forward (block) all datagrams sent by host 128.119.1.1*

# Examples

## Destination-based layer 2 (switch) forwarding:

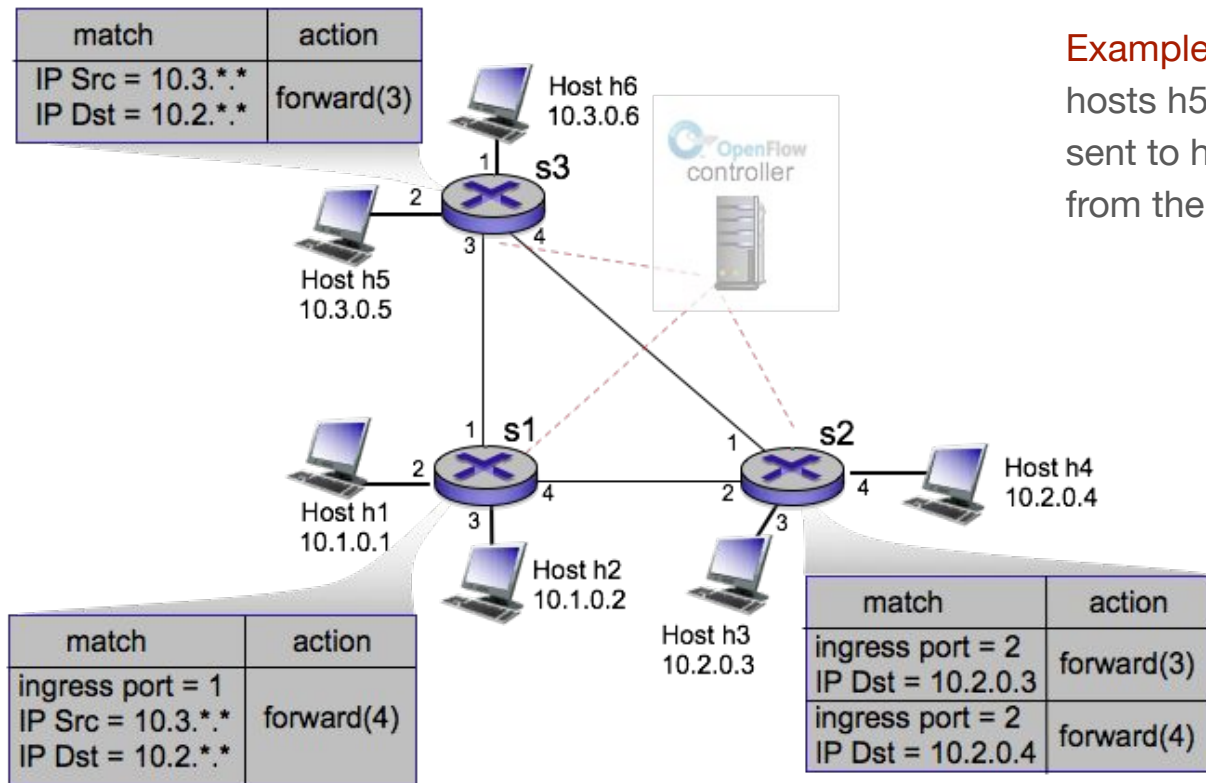
Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	22:A7:23:11:E1:02	*	*	*	*	*	*	*	*	port3

*layer 2 frames from MAC address 22:A7:23:11:E1:02  
should be forwarded to output port 6*

# OpenFlow abstraction

- **match+action:** unifies different kinds of devices
- Router
  - **match:** longest destination IP prefix
  - **action:** forward out a link
- Switch
  - **match:** destination MAC address
  - **action:** forward or flood
- Firewall
  - **match:** IP addresses and TCP/UDP port numbers
  - **action:** permit or deny
- NAT
  - **match:** IP address and port
  - **action:** rewrite address and port

# OpenFlow example



**Example:** datagrams from hosts h5 and h6 should be sent to h3 or h4, via s1 and from there to s2

# Software defined networking (SDN)

---

- Internet network layer: historically has been implemented via distributed, per-router approach
  - **monolithic router contains switching hardware**, runs proprietary implementation of Internet standard protocols (IP, RIP, IS-IS, OSPF, BGP) in proprietary router OS (e.g., Cisco IOS)
  - different “middleboxes” for different network layer functions: firewalls, load balancers, NAT boxes, ..
- ~2005: renewed interest in rethinking network control plane





# Software defined networking (SDN)

- Why a logically centralized control plane?
  - easier network management: avoid router misconfigurations, greater flexibility of traffic flows
  - table-based forwarding (recall OpenFlow API) allows “programming” routers
    - centralized “programming” easier: compute tables centrally and distribute
    - distributed “programming”: more difficult: compute tables as result of distributed algorithm (protocol) implemented in each and every router
  - open (non-proprietary) implementation of control plane


# Analogy: mainframe to PC evolution




Vertically integrated  
Closed, proprietary  
Slow innovation  
Small industry



Horizontal  
Open interfaces  
Rapid innovation  
Huge industry



# Traffic engineering: difficult traditional routing




Q: what if network operator wants u-to-z traffic to flow along uvwz, x-to-z traffic to flow xwyz?

A: need to define link weights so traffic routing algorithm computes routes accordingly (or need a new routing algorithm)!

Link weights are only control “knobs”: wrong!


# Traffic engineering: difficult



Q: what if network operator wants to split u-to-z traffic along uvwz and uxyz (load balancing)?

A: can't do it (or need a new routing algorithm)


# Traffic engineering: difficult



Q: what if w wants to route blue and red traffic differently?

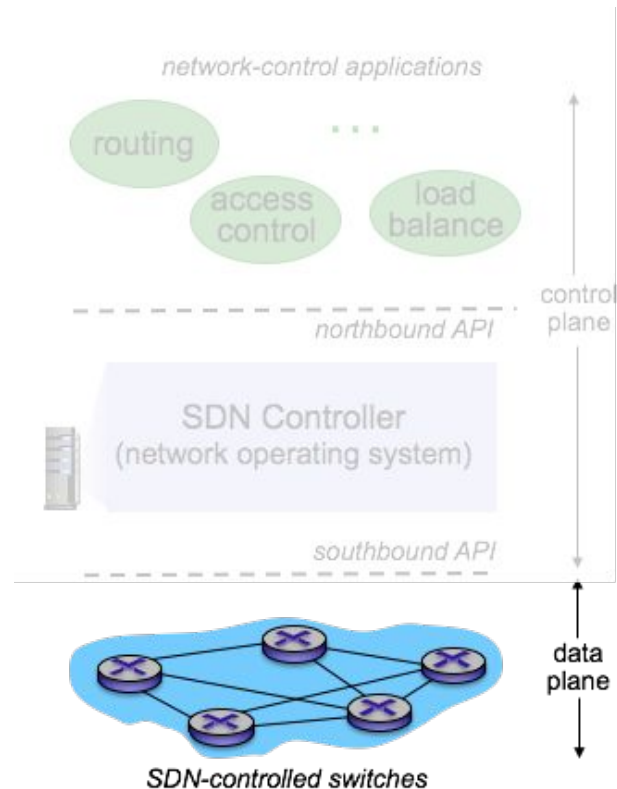
A: can't do it (with destination based forwarding, and LS, DV routing)

# Software defined networking (SDN)



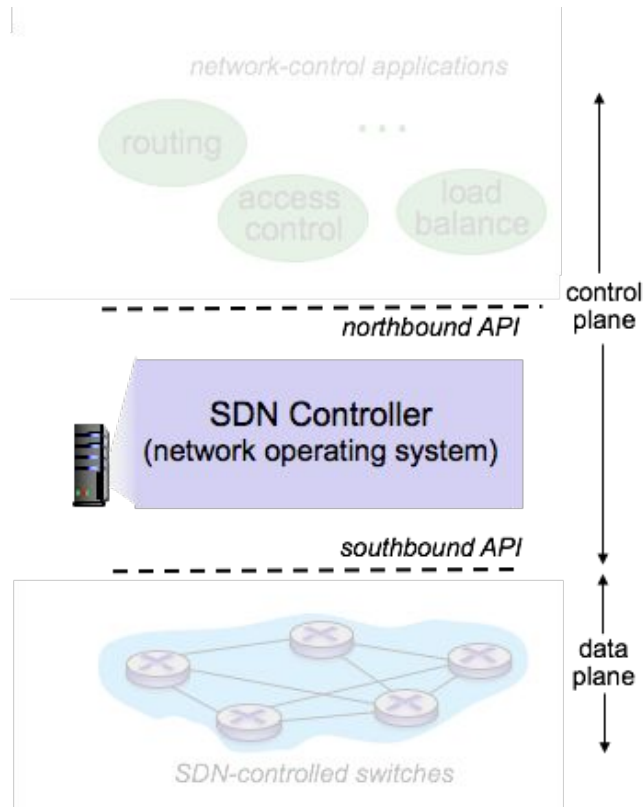
# SDN perspective: data plane switches

- Data plane switches
  - fast, simple, commodity switches implementing generalized data-plane forwarding in hardware
  - switch flow table computed, installed by controller
  - API for table-based switch control (e.g., OpenFlow)
    - defines what is controllable and what is not
  - protocol for communicating with controller (e.g., OpenFlow)



# SDN perspective: SDN controller

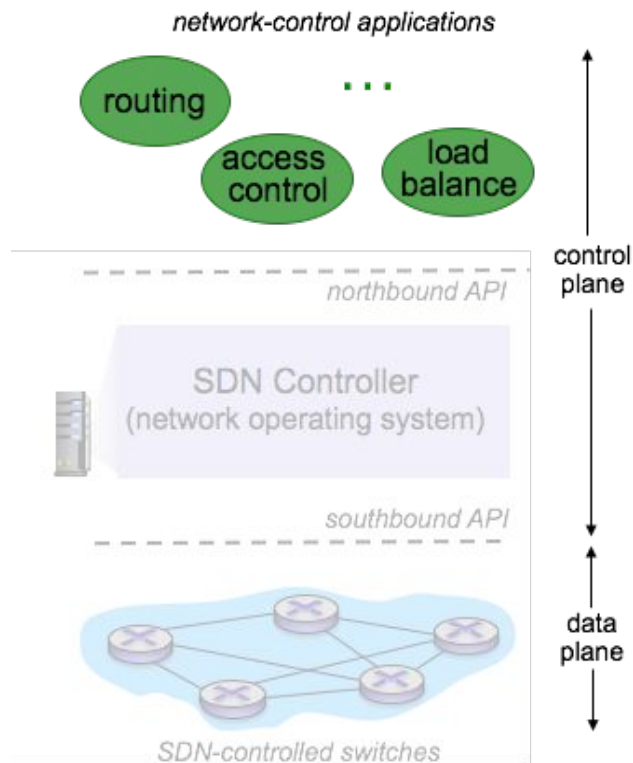
- SDN controller (network OS):
  - maintain network state information
  - interacts with network control applications “above” via northbound API
  - interacts with network switches “below” via southbound API
  - implemented as distributed system for performance, scalability, fault-tolerance, robustness






# SDN perspective: network apps

- network-control apps:
  - “brains” of control: implement control functions using lower-level services, API provided by SDN controller
  - unbundled: can be provided by 3rd party: distinct from routing vendor, or SDN controller




# Components of SDN controller

- Interface layer to network control apps: abstractions API
- Network-wide state management layer: state of networks links, switches, services: **a distributed database**
- communication layer: communicate between SDN controller and controlled switches




# OpenFlow Protocol



- operates between controller, switch
- TCP used to exchange messages
  - optional encryption
- three classes of OpenFlow messages:
  - controller-to-switch
  - asynchronous (switch to controller)
  - symmetric (misc)


# OpenFlow: controller-to-switch messages

- Key **controller-to-switch** messages
  - **features:** controller queries switch features, switch replies
  - **configure:** controller queries/sets switch configuration parameters
  - **modify-state:** add, delete, modify flow entries in the OpenFlow tables
  - **packet-out:** controller can send this packet out of specific switch port



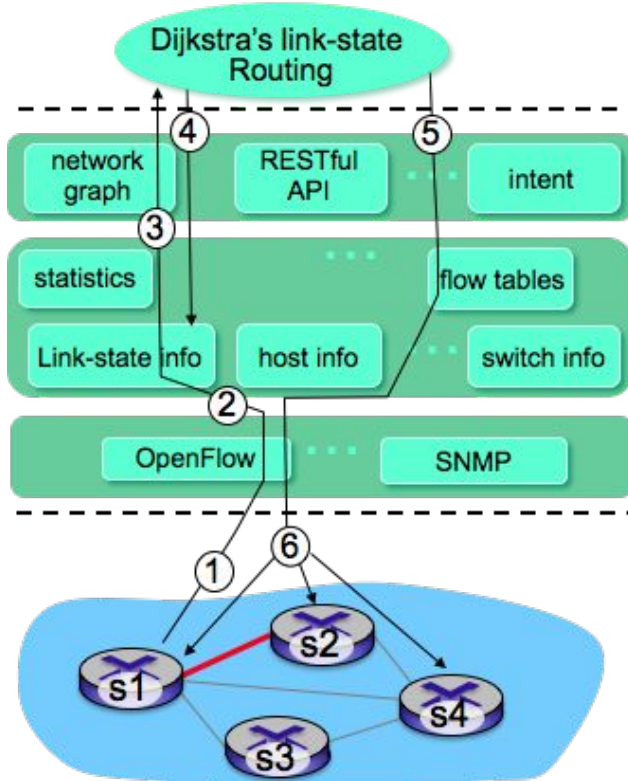
# OpenFlow: switch-to-controller messages

- Key **switch-to-controller** messages
  - **packet-in:** transfer packet (and its control) to controller. See packet-out message from controller
  - **flow-removed:** flow table entry deleted at switch
  - **port status:** inform controller of a change on a port.




Fortunately, network operators **don't "program" switches** by creating/sending OpenFlow **messages directly**. Instead use **higher-level abstraction at controller**

# SDN control/data plane interaction example



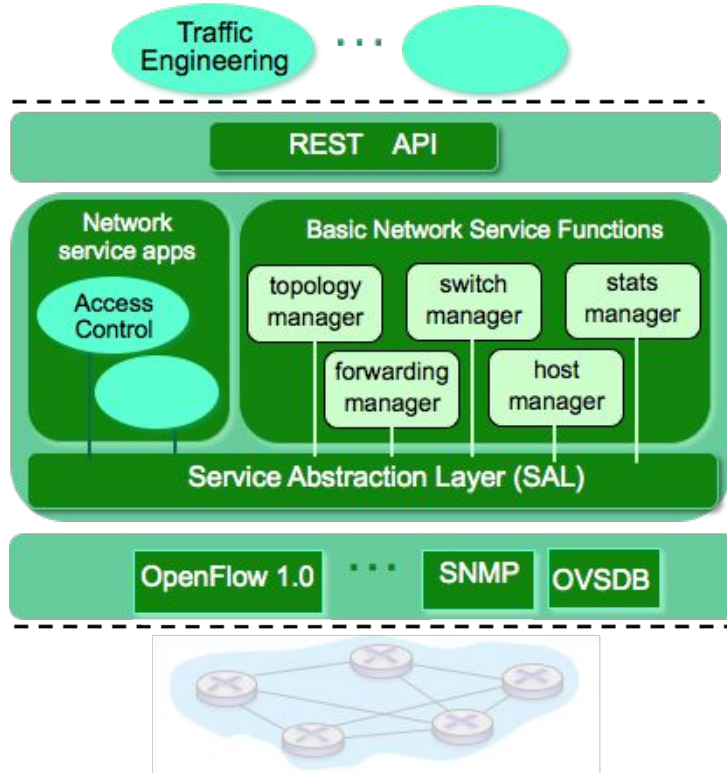
1. **S1**, experiencing link failure using **OpenFlow port status** message to notify controller
2. **SDN controller receives OpenFlow message**, updates link status info
3. **Dijkstra's routing algorithm application has previously registered to be called when ever link status changes**. It is called.
4. **Dijkstra's** routing algorithm access network graph info, link state info in controller, **computes new routes**

# SDN control/data plane interaction example



5. link state routing app interacts with flow-table-computation component in SDN controller, which computes new flow tables needed
6. Controller uses OpenFlow to install new tables in switches that need updating

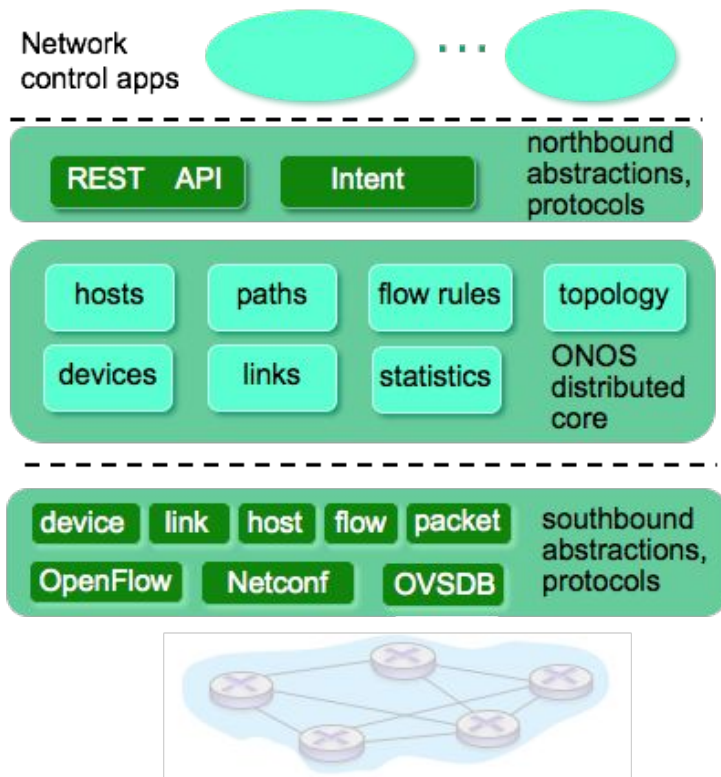
# OpenDaylight Controller



- ODL Lithium controller
- network apps may be contained within, or be external to SDN controller
- **Service Abstraction Layer:** interconnects internal, external applications and services



# ONOS Controller



- control apps separate from controller
- **intent framework:**  
**high-level specification of service:** what rather than how
- considerable **emphasis on distributed core:**  
service reliability,  
replication performance  
scaling

# SDN: selected challenges

---

- hardening the control plane: dependable, reliable, performance-scalable, secure distributed system
  - robustness to failures: leverage strong theory of reliable distributed system for control plane
  - dependability, security: “baked in” from day one?
- networks, protocols meeting mission-specific requirements
  - e.g., real-time, ultra-reliable, ultra-secure
- Internet-scaling